

WELCOME TO DAY 2

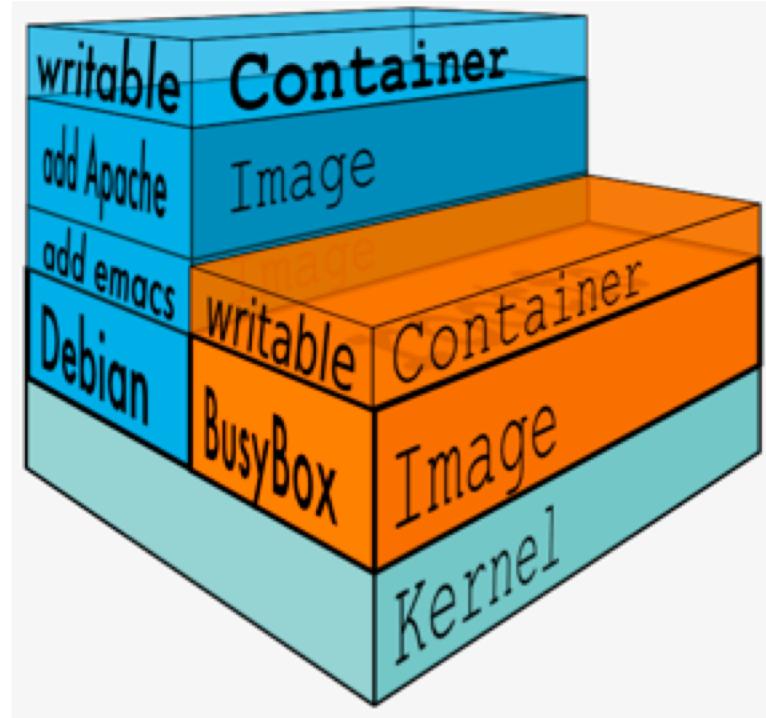
MINING THE DOCKER FILE SYSTEM

module 04

Section 4.1 | remember this!

docker's first technical drawing

- visualize the layers of a docker container image
- distinguish how a container layers on the image
- depict possible actions at each layer



Section 4.1 | remember this!

Images

Images are read-only, how do we change it?

- We create a new container from that image.
- Then we make changes to that container.
- When we are satisfied with those changes, we transform them into a new layer.
- A new image is created by stacking the new layer on top of the old image.

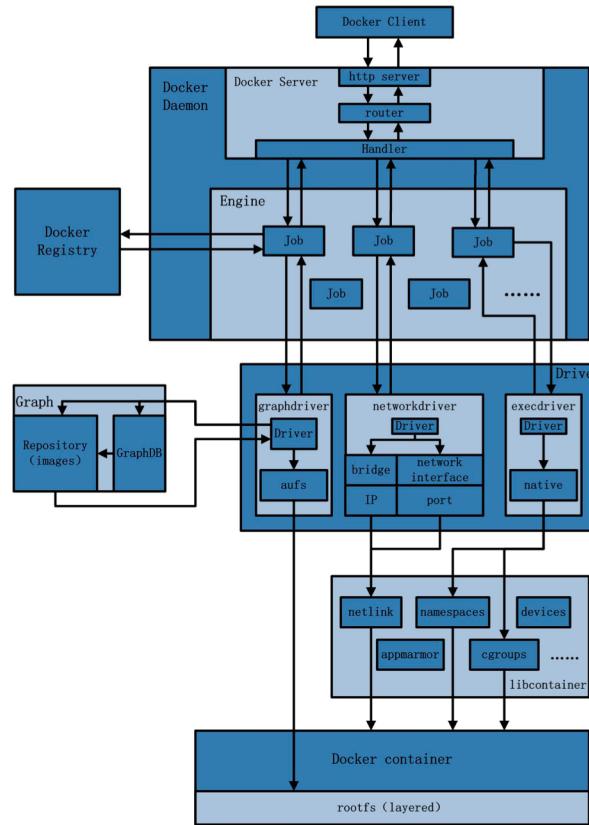


the graph driver

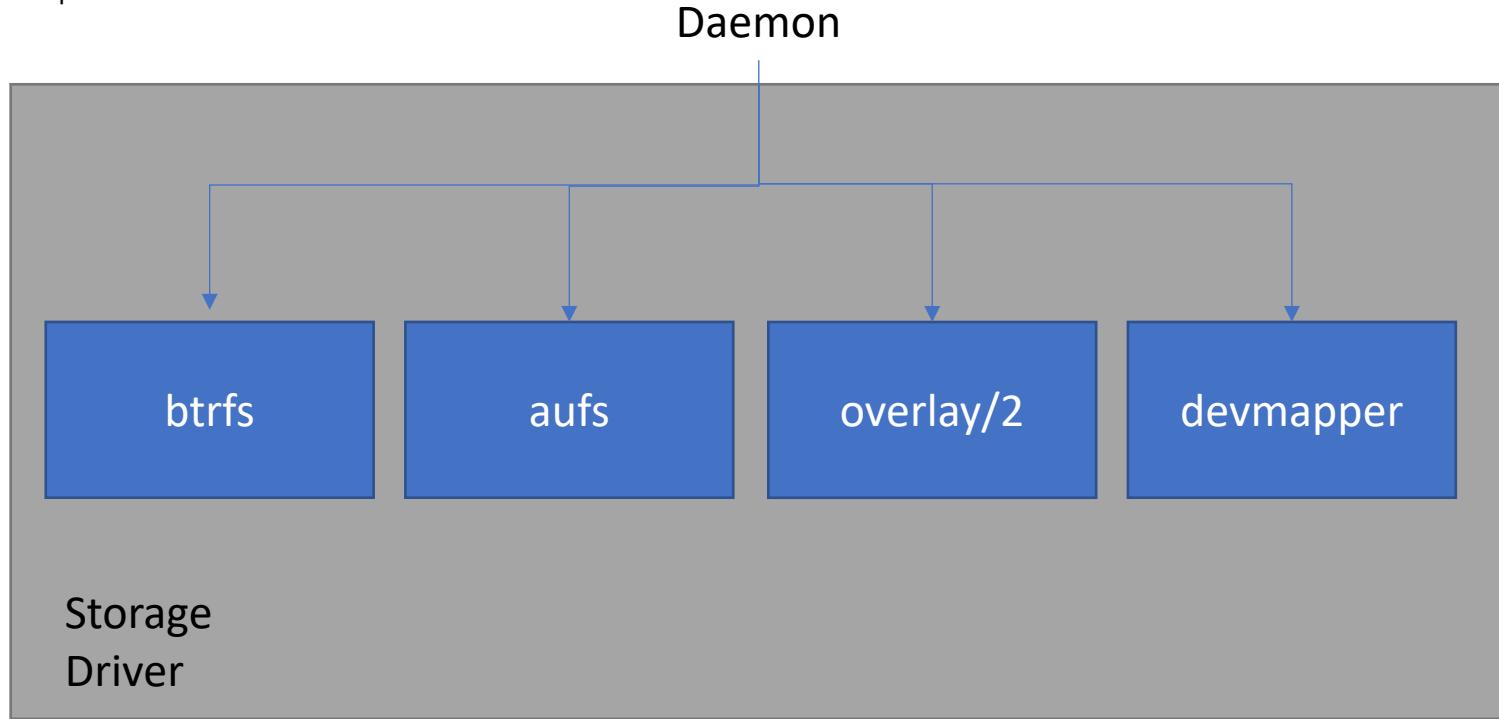
- manages how containers and container images are stored
- independent from volume mount storage
- default == /var/lib/docker/



Section 4.1 | remember this!

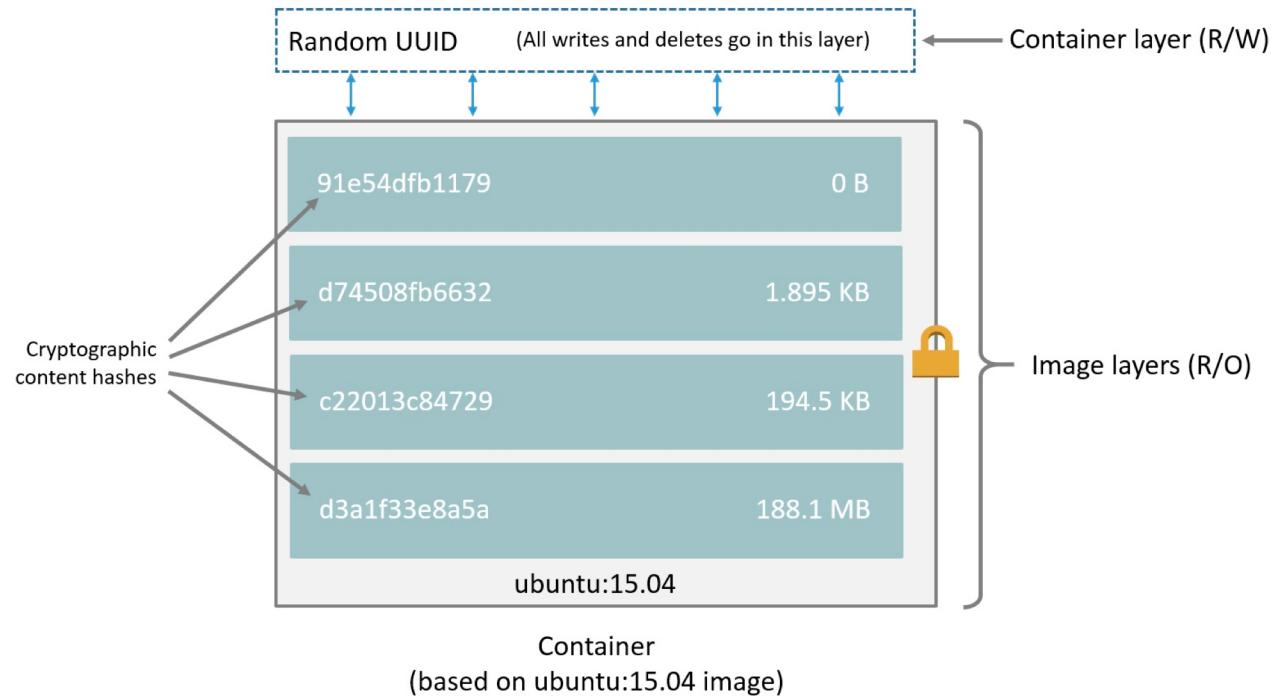


Section 4.1 | remember this!

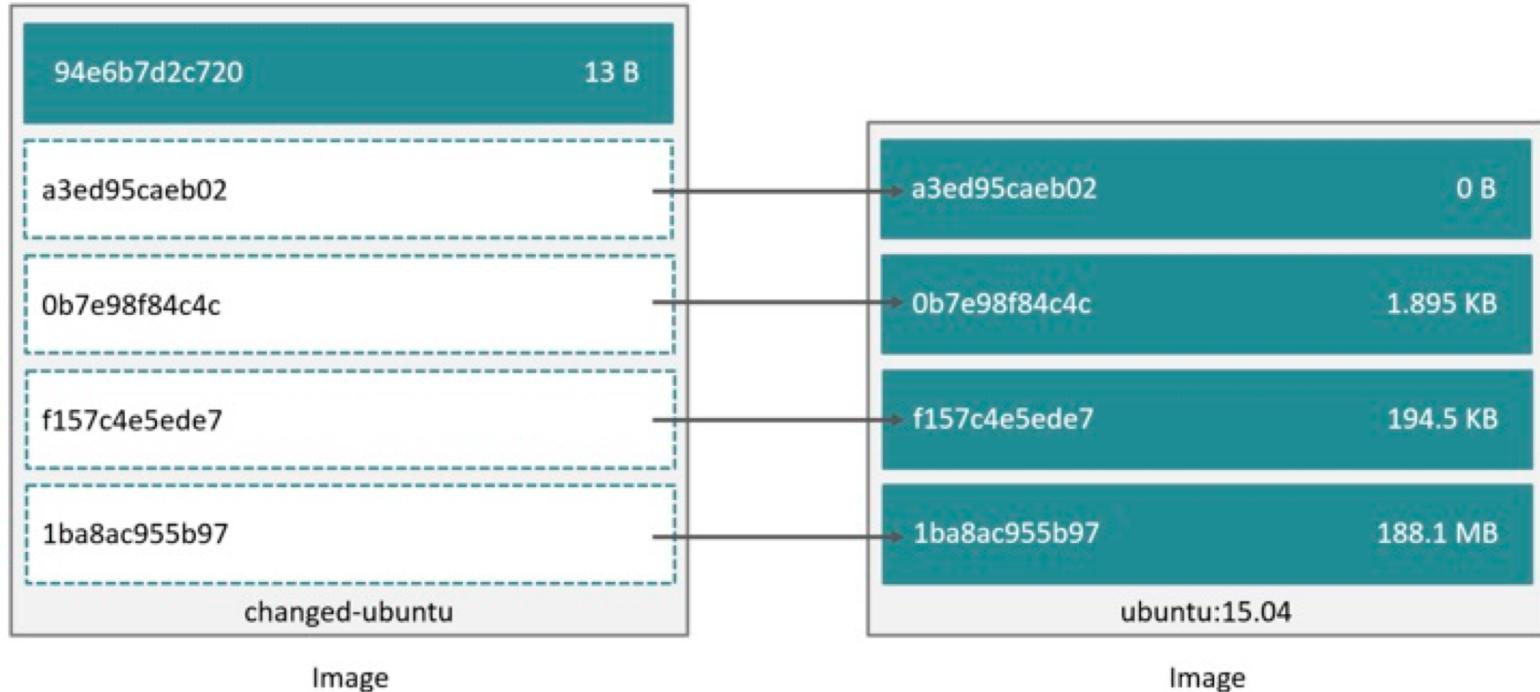


Section 4.2 | images

content addressable storage

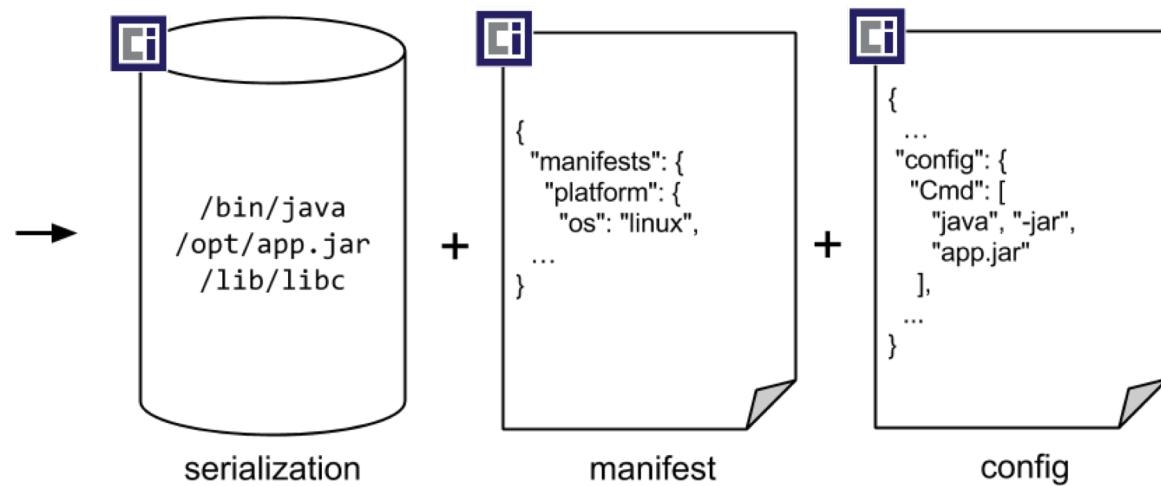


content addressable storage



oci image spec

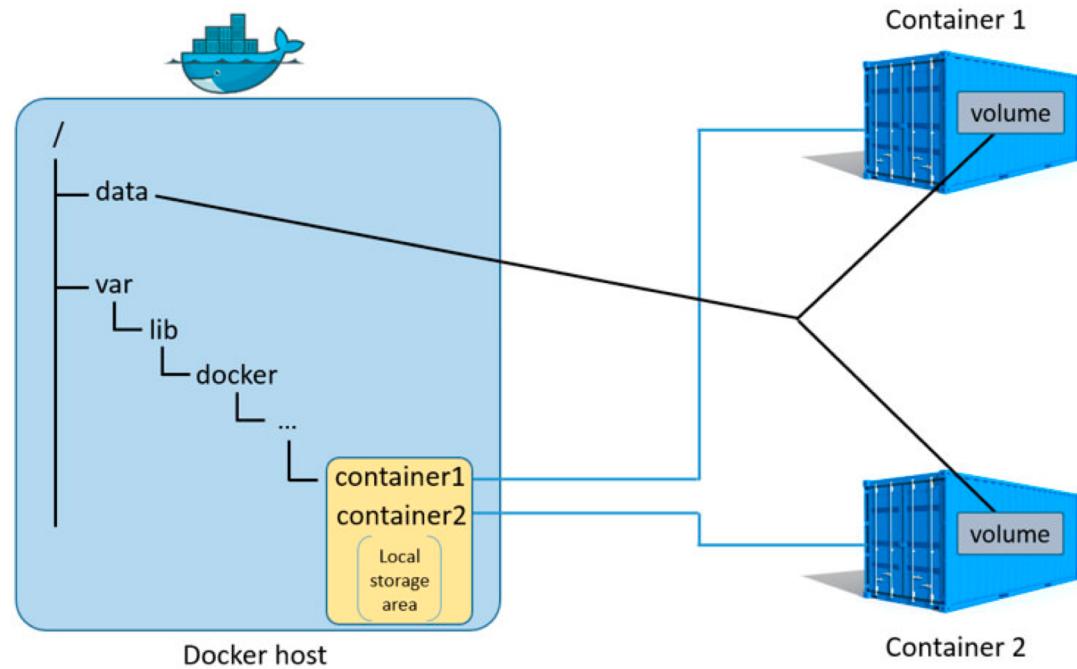
```
public class HelloWorld {  
    public static void main(String[] args) {  
        System.out.println("Hello, World");  
    }  
}
```



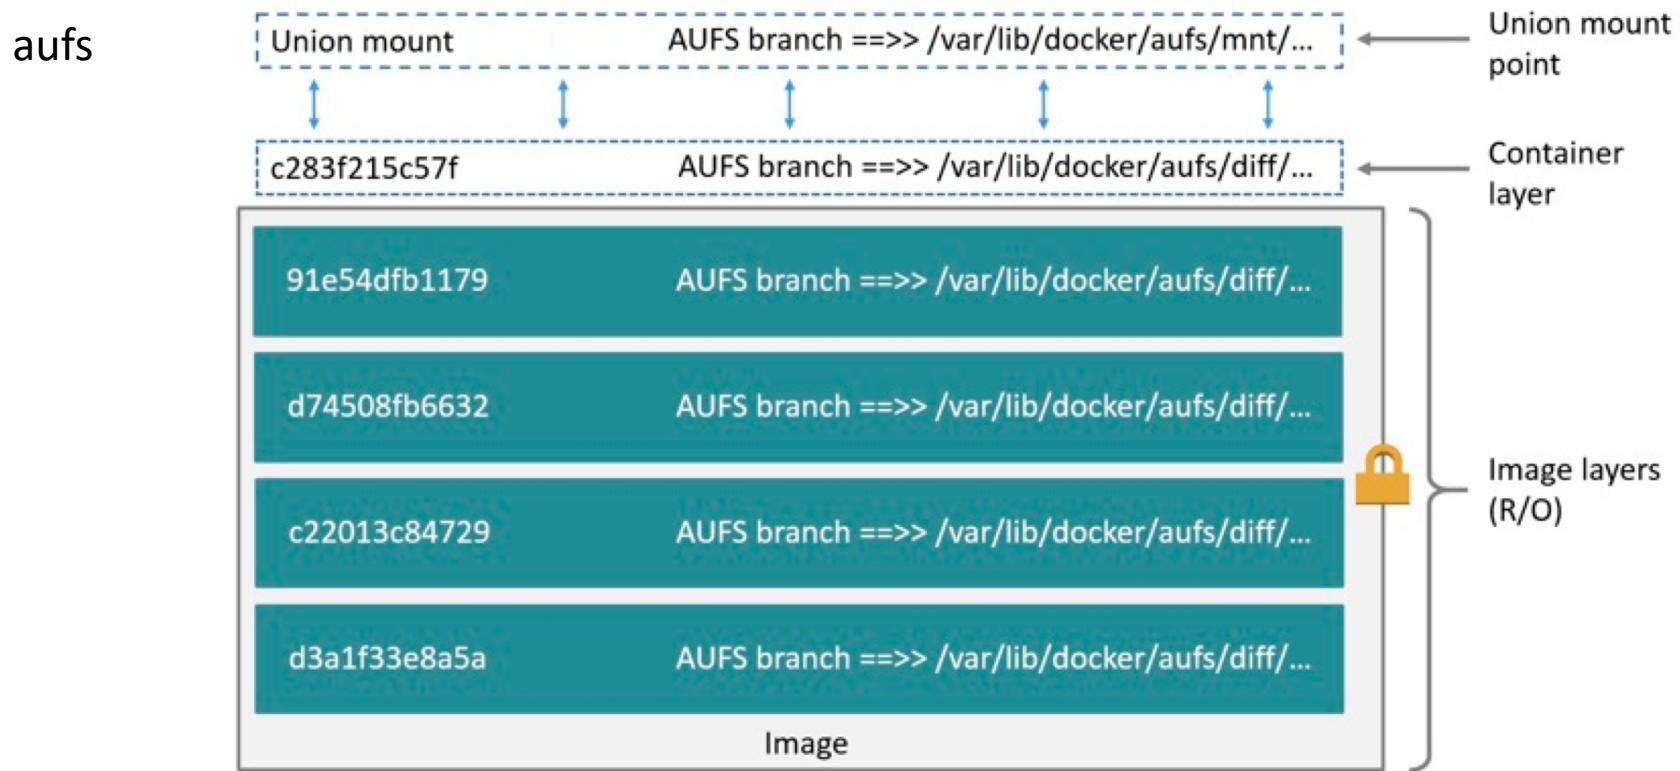
Section 4.3 | storage drivers

Section 4.3 | storage drivers

graph directory



Section 4.3 | storage drivers



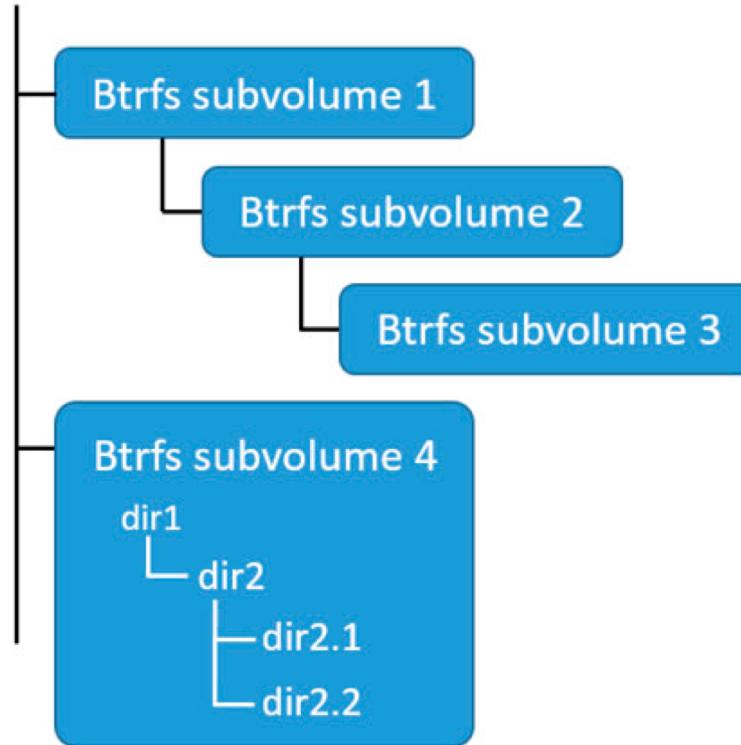
aufs

- Container layer read scenarios
 - File does not exist in container
 - File only exists in the container layer
 - Exists in both
- Write Scenarios
 - Writing to file for the first time (copy up)
 - Deleting files (white files, opaque file)
 - Renaming Directories (not supported)



Section 4.3 | storage drivers

btrfs



btrfs

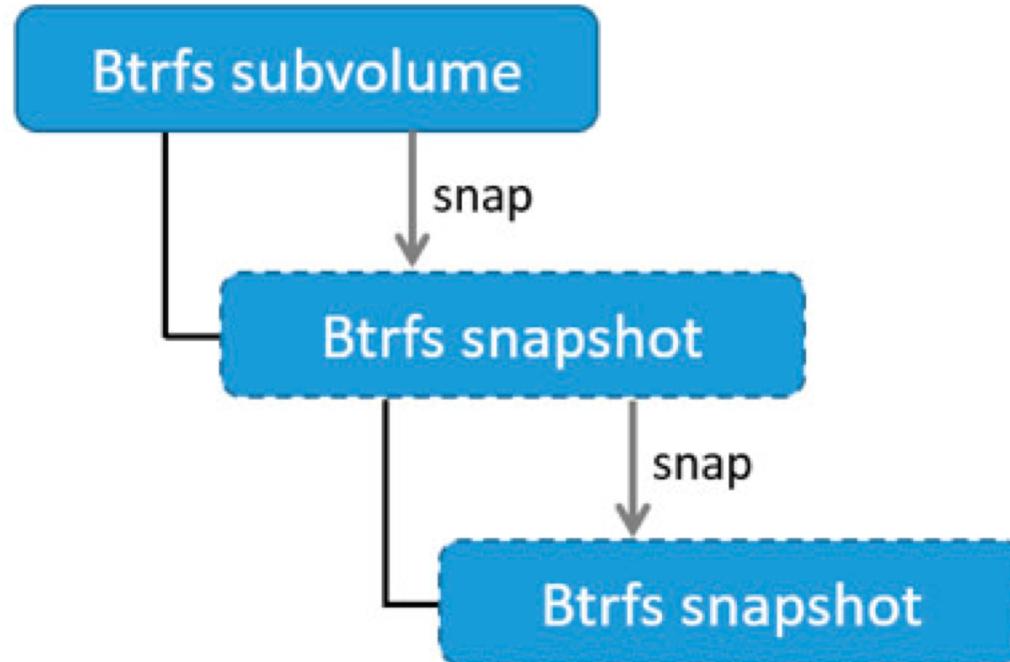
- Block level ops, thin provisioning, copy on write
- Supported on Docker CE on Ubuntu, Docker EE/CS on SLES
- Requires dedicated storage, btrfs mount at /var/lib/docker
- Requires a “btrfs” tool for creating mounts
- Configure daemon.json



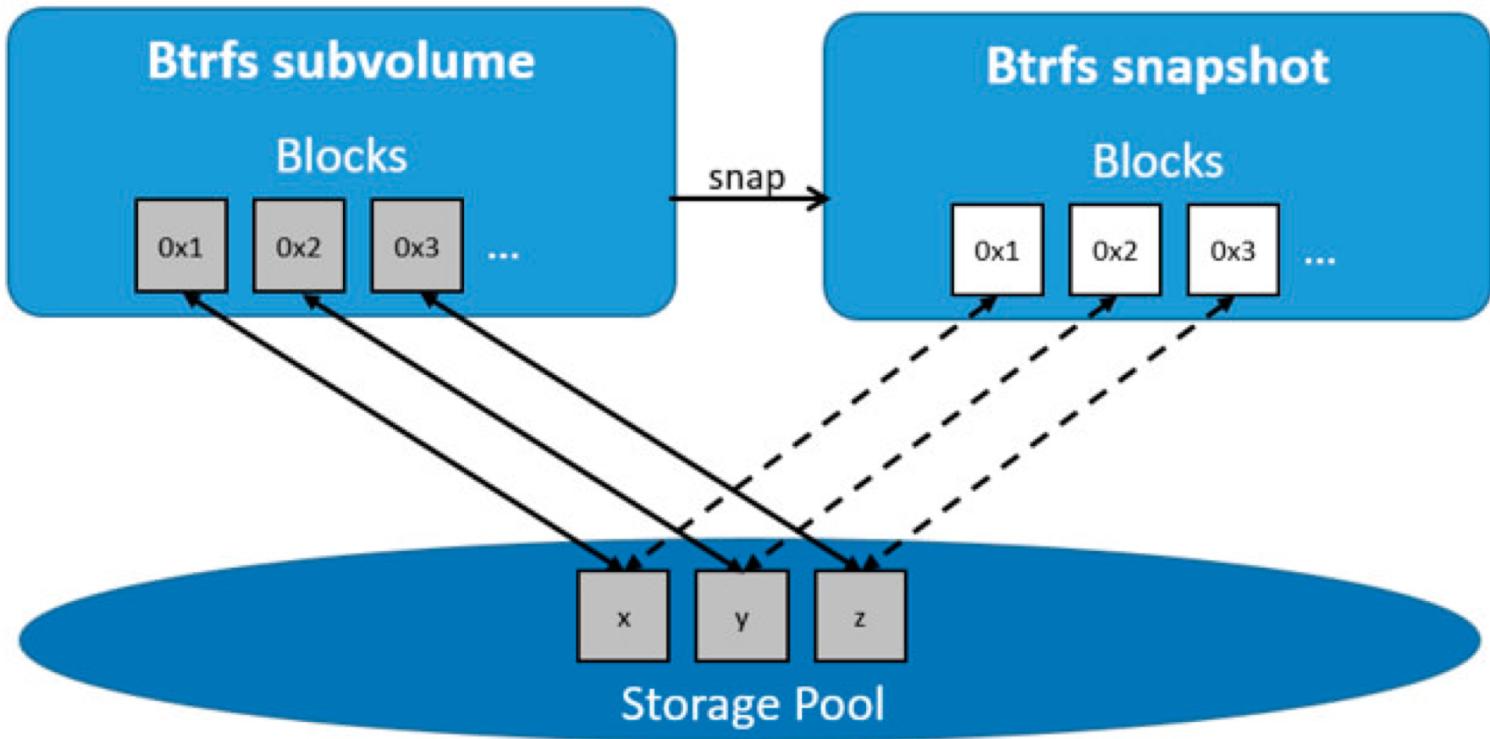
Btrfs (continued...)

- When using balanced mode with BTRFS, 1 GB chunks will be added when storage pressures occur
- The /subvolumes/ dir
 - One directory per image layer
 - Union filesystem built from a layer plus all its parent layers
 - Can be nested and snapshotted if necessary

btrfs



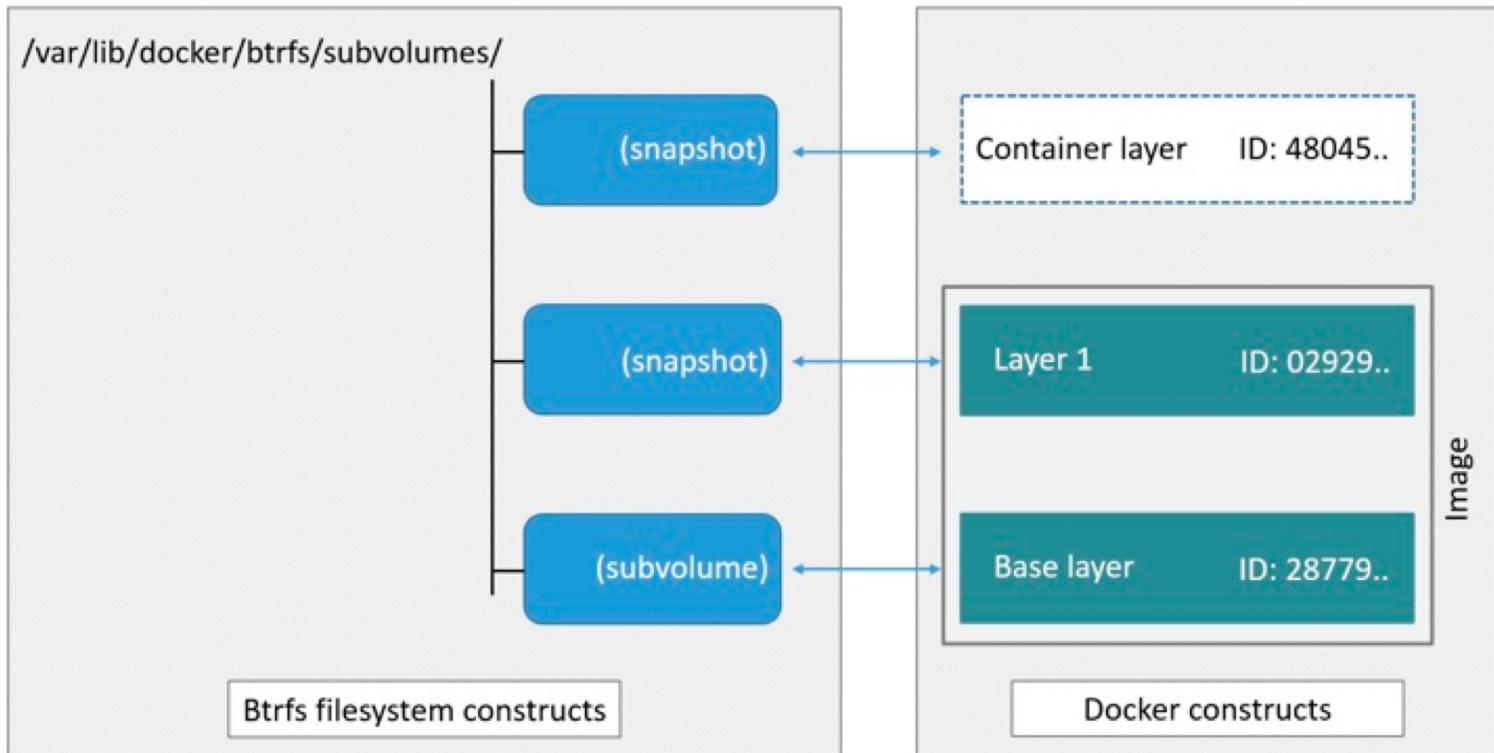
btrfs



NEBULAWORKS

Section 4.3 | storage drivers

btrfs



Section 4.1 | remember this!

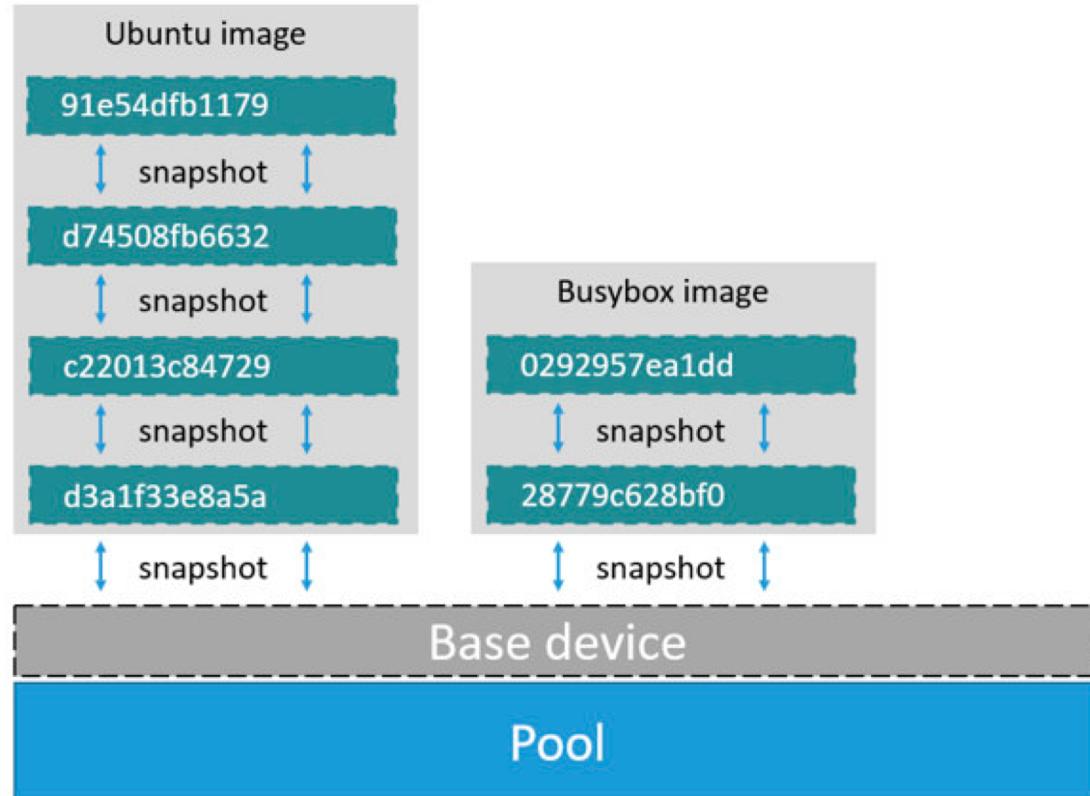
btrfs

- Reads
 - A container is a space-efficient snapshot of an image
 - Metadata blocks in the storage pool point to actual blocks
- Writes
 - Write a new file allocates new block in the containers snapshot
 - Same as writing to a subvolume
- Write to Existing file
 - Updating an existing file in a container is a COW op
 - The original data is read from the layer where the file exists, and only the modified blocks are written to top layer
 - Performing many writes on small files may slow performance



Section 4.3 | storage drivers

dev mapper



NEBULAWORKS

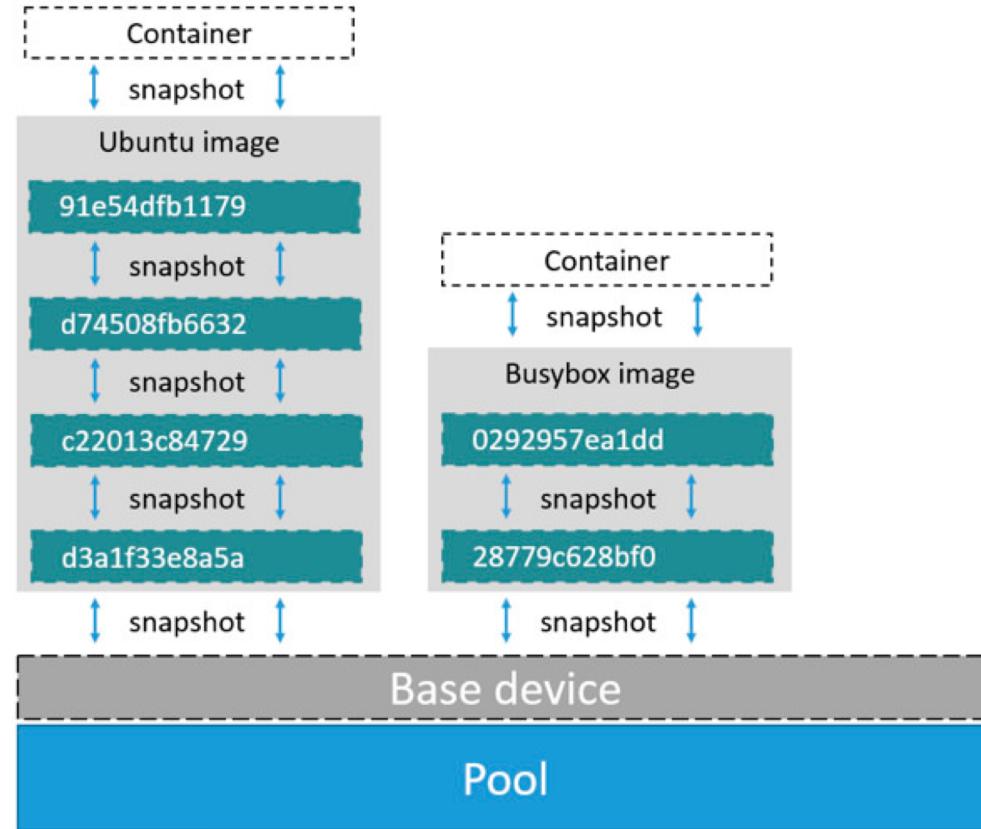
Section 4.1 | remember this!

dev mapper

- /var/lib/devicemapper/mnt/ contains metadata about the dev mapper config
- /var/lib/docker/mnt
 - Contains a mountpoint for each image and layer that exists

Section 4.3 | storage drivers

dev mapper

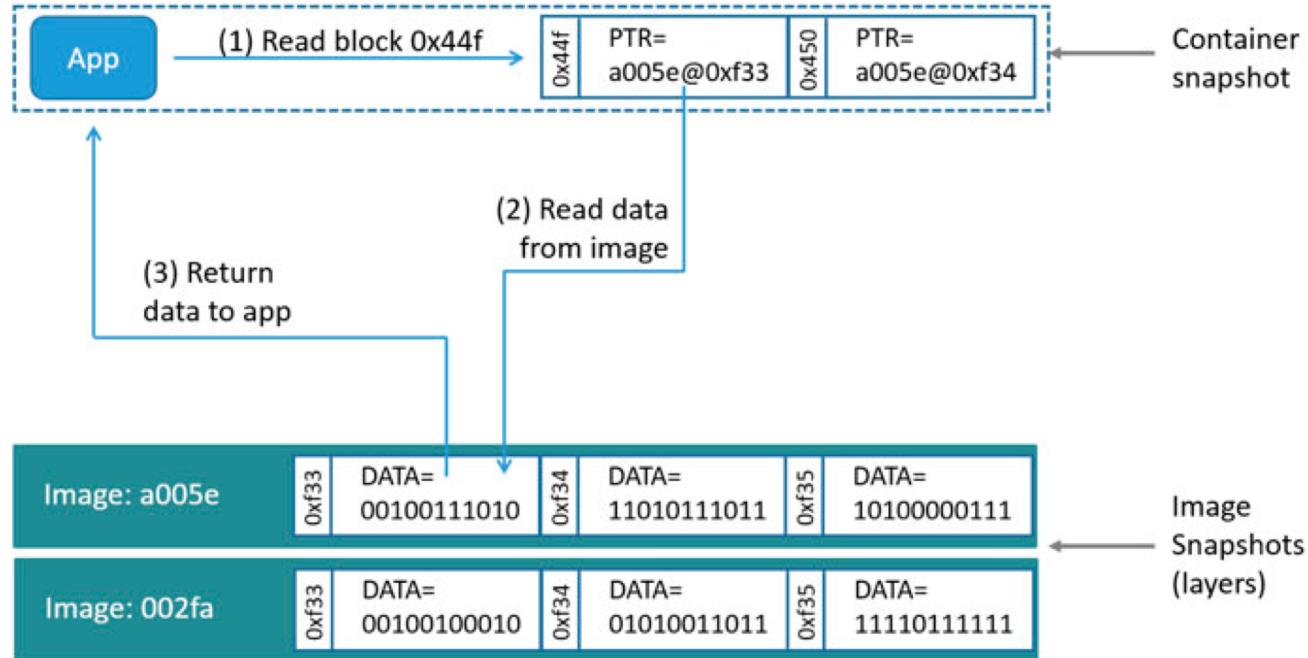


dev mapper

- Writing File Scenarios
 - Writing a new file uses allocate-on-demand op
 - Each block of the new file is added in the containers writable layer
- Update a file
 - Block of the file is read from nearest layer where it exists, only modified blocks are written to writable layer
- Delete a file
 - Devicemapper will intercepts all reads to this resource and will through directory does not exist



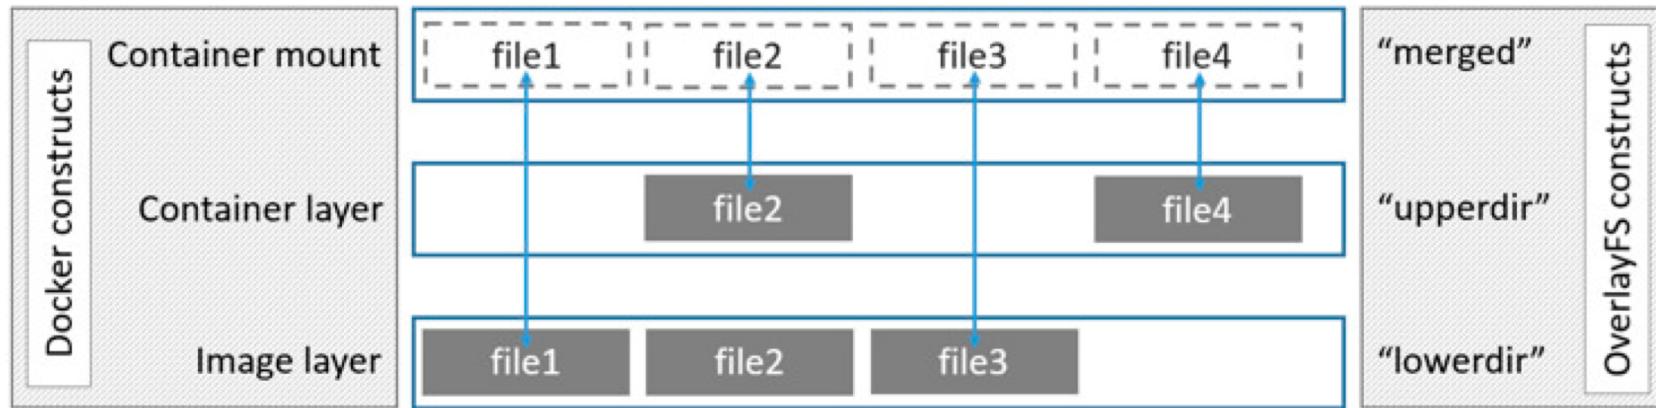
dev mapper



Overlay and overlay 2

- Recommend to use overlay2, more efficient
- Overlay 2 supported on Docker EE/CE
- Needs Kernel 4.0 or higher
- Recommended to use a separate backing filesystem form the one used by /var/lib/docker
- Overlay2 has up to 128 lower level layers, while overlay supported 2

overlays

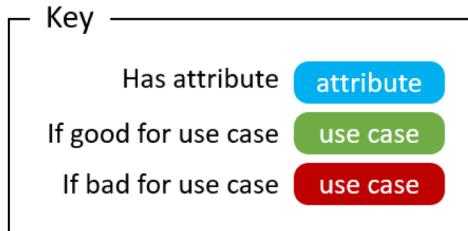


overlay2

- Use overlay2
- /var/lib/docker/overlay2/l
 - Lowercase L
 - Layer identifiers symbolic links
- Lowest layer is called a link, that contains a name of the shortened identifier
 - Contains diff directory as well
- Second lowest layer and all above it have “lower” directory
 - Lower denotes a parent, and diff contains the layer contents

Section 4.3 | storage drivers

AUFS	stable	production-ready	good memory use	smooth Docker experience	high write activity	PaaS-type work
Devicemapper (loop)	stable	in mainline kernel	smooth Docker experience	production	performance	lab testing
Devicemapper (direct-lvm)	stable	production-ready	in mainline kernel	smooth Docker experience	PaaS-type work	
Btrfs	in mainline kernel	high write activity	container churn	build pools		
Overlay	stable	good memory use	in mainline kernel	container churn	lab testing	Overlay2
ZFS native (ZoL)		PaaS-type work				
ZFS FUSE	stable	lab testing	production			



Section 4.5 | volume plugins

The problem

- We can store data inside of containers **BUT** we usually don't want to.
 - Containers are ephemeral
 - Storage drivers aren't very efficient.

How the heck do I persist data?

The problem

- We can store data inside of containers **BUT** we usually don't want to.
 - Containers are ephemeral
 - Storage drivers aren't very efficient.

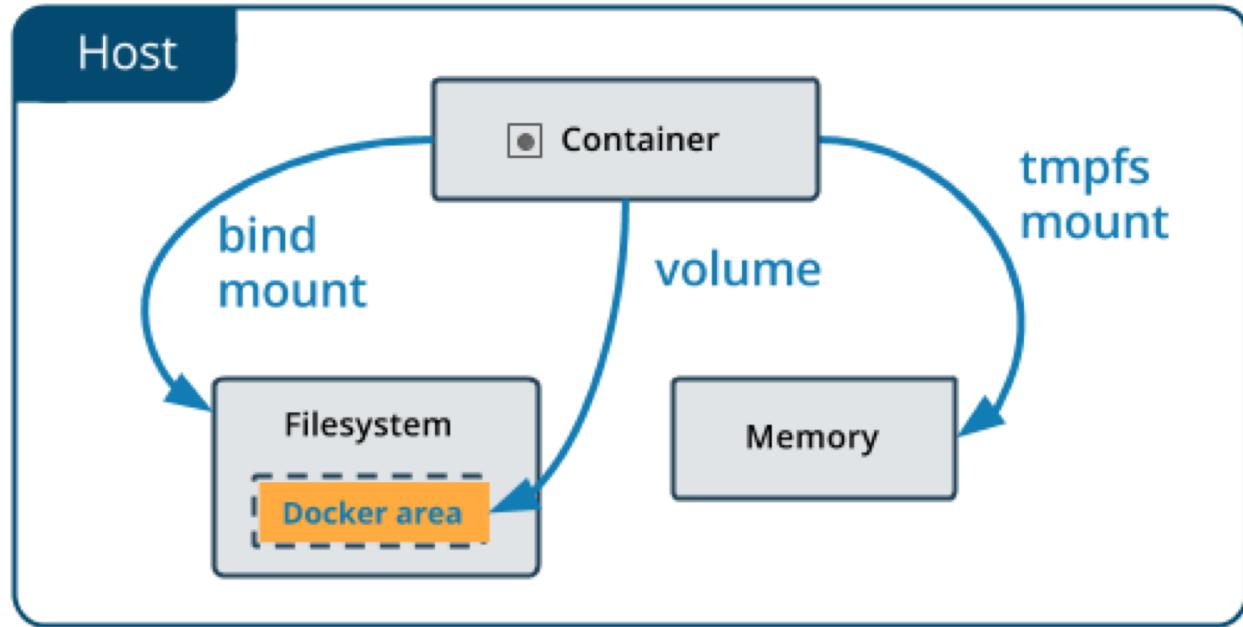
How the heck do I persist data?

What are volumes?

- Volumes are a place to store data that live outside the lifecycle of a container.
- They are mounted into a container
- First class citizens in Docker
- Work with Linux and Windows containers.
- They don't add to a containers size!
- Can be created and managed (by docker) outside the scope of any container

`/var/lib/docker/volumes/<volume_name>`

Volumes



Source : docker documentation

managing volumes with the docker volume object

manages persistent data

- local copy by default
- stored in the graph dir
- accessed by labels
- created with a plugin architecture



The bad

- Docker ships with a default volume driver that supports local, host-based volumes.
- The default volume driver may not be suitable for certain workloads (data-intensive workloads).
- It won't cut it for distributed stateful applications.

Lab 3

Volume Plugins *A closer look*

The Basics

- Docker Engine Plugins allow you to extend the capabilities of the Docker Engine.
- Extends the local volume driver
- Allow you to interface with external storage devices / platforms.

Volume Plugin Architecture

The Docker Plugin Specification

- Standardized API interface (9 endpoints)
- Plugin packaging (build as an image)
- Plugin storage (hub/store)
- Plugin Management (installation, configuration, deletion)
 - https://docs.docker.com/engine/extend/plugin_api/-plugin-lifecycle



managing volumes

docker volume command (under the hood)

- `volume create`
 - `/VolumeDriver.Get` (checks to see if it exists)
 - `/volumeDriver.Create`
- `volume inspect`
 - `VolumeDriver.Get`
- `volume ls`
 - `/VolumeDriver.List`
- `volume rm`
 - `/VolumeDriver.Remove`



Standardized API Interface

- /Plugin.Activate
- /VolumeDriver.Create
- /VolumeDriver.Remove
- /VolumeDriver.Mount
- /VolumeDriver.Path
- /VolumeDriver.Unmount
- /VolumeDriver.Get
- /VolumeDriver.List
- /VolumeDriver.Capabilities

/VolumeDriver.Create

- Request body: {"Name": "volume_name", "opts":{}}
- Response : {"Err":""}
- Notes : This should only create the volume (no mounting)

- This endpoint gives our plugin the ability to create a volume
- Docker CLI -> Docker Daemon -> Plugin API.
- The plugin must be able to respond with the required response.

Info on the plugin api :

- https://docs.docker.com/engine/extend/plugins_volume/-command-line-changes



Managing a Plugin Lifecycle

```
first class citizens : docker plugin <cmds>
plugin create
plugin inspect
plugin ls
plugin rm
```



- Established in 2014 by Goutham Rao and friends
- Container data services platform
- Software-defined storage





Providing a “Cloud Agnostic” Storage Solution

- Portworx can run on any infrastructure.
- It will aggregate the host storage resources into storage classes and allow them to be consumed as docker volumes by application containers.

Any Stateful Container

Redis, ELK, Cassandra, Jenkins & More

Any Scheduler

Docker, Kubernetes, Mesosphere

Your Data Layer

PORTWORX

Any Infrastructure

AWS, Azure, Google, On-Prem

Section 4.5 | volume plugins





Use Cases

- **Databases** : compatible with a lot of dbs
- **CI/CD** : Testing prod env & the golden image
- **Content Management** : Horizontal scale of CMS sites
- **Docker storage** : Persistent Container storage across a swarm
- **Kubernetes storage**

Deploying portworx with Docker

Prerequisites:

- You must have a storage device that you can allocate to portworx
 - Non-root storage devices.
lsblk
- You must configure Docker to allow shared mounts propagations.

```
$ docker run -it -v /mnt:/mnt:shared busybox sh -c /bin/date  
$ sudo mount -make-shared
```

- A minimum version of Docker 1.10 must be installed.
 - runC must be installed
- ```
$ sudo apt-get install runC
```



## Deploying portworx with Docker

### Install PX OCI bundle

- Portworx gives us a Docker based installation utility that helps deploy the PX OCI bundle. You can install the bundle by running the following Docker container:
- Curl for the latest version of portworx

```
$ latest_stable=$(curl -fsSL
'http://install.portworx.com:8080?type=dock&stork=false' | awk '/image: /
{print $2}')
```

- Download the OCI bits

```
$ sudo docker run --entrypoint /runc-entry-point.sh \ --rm -i --
privileged=true \ -v /opt/pwx:/opt/pwx -v /etc/pwx:/etc/pwx \
$latest_stable
```

## Deploying portworx with Docker



### Configure PX under runC

- Once the PX OCI bundle is installed, you can run the following command from the bundle to configure system to start PX runC.

```
$ sudo /opt/pwx/bin/px-runc install -c MY_CLUSTER_ID \ -k
etcd://myetc.company.com:2379 \ -s /dev/xvdb
```

## Deploying portworx with Docker



### Starting PX runC

- Reload systemd configurations and enable and start the Portworx service

```
$ sudo systemctl daemon-reload
$ sudo systemctl enable portworx
$ sudo systemctl start portworx
```

- Now that portworx is installed we can create a stateful application with Docker Swarm!

## Deploying a Jenkins Service with the portworx driver for Docker



### Create the Volume

```
$ docker volume create -d pxd --name jenkins_vol --opt size=4 \
--opt block_size=64 --opt repl=3 --opt fs=ext4 --opt shared=true
```

### Create the Service

```
$ docker service create --name jenkins \
--replicas 3 \
--publish 8082:8080 \
--publish 50000:50000 \
-e JENKINS_OPTS="--prefix=/jenkins" \
--reserve-memory 300m \
--mount "type=volume,volume-driver=pxd,source=jenkins_vol,target=/var/jenkins_home"\
```

### volume plugin – single platform

| plugin                    | description                                                                                                                                                                                                                                                                         |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Blockbridge Volume plugin | A volume plugin that provides access to an extensible set of container-based persistent storage options. It supports single and multi-host Docker environments with features that include tenant isolation, automated provisioning, encryption, secure deletion, snapshots and QoS. |
| Contiv Volume Plugin      | An open source volume plugin that provides multi-tenant, persistent, distributed storage with intent based consumption using ceph underneath.                                                                                                                                       |

## Section 4.5 | volume plugins

| plugin                                                 | description                                                                                                                                                                   |
|--------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Flocker plugin<br>(last source code commit 1 year)     | A volume plugin that provides multi-host portable volumes for Docker, allowing you to run databases and other stateful containers and move them across a cluster of machines. |
| gce-docker plugin<br>(last source code commit 2 years) | A volume plugin able to attach, format and mount Google Compute persistent-disks.                                                                                             |
| GlusterFS plugin<br>(last source code commit 2 years)  | A volume plugin that provides multi-host volumes management for Docker using GlusterFS.                                                                                       |
| IPFS Volume Plugin<br>(last source code commit 1 year) | An open source volume plugin that allows using an ipfs filesystem as a volume.                                                                                                |

## volume plugin

| plugin                                            | description                                                                                                                                                                                                                                    |
|---------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Convoy plugin<br>(Last source code update 1 year) | A volume plugin for a variety of storage back-ends including device mapper and NFS. It's a simple standalone executable written in Go and provides the framework to support vendor-specific extensions such as snapshots, backups and restore. |
| Horcrux Volume Plugin                             | A volume plugin that allows on-demand, version controlled access to your data. Horcrux is an open-source plugin, written in Go, and supports SCP, Minio and Amazon S3.                                                                         |
| Netshare plugin                                   | A volume plugin that provides volume management for NFS[34], AWS EFS and CIFS file systems.                                                                                                                                                    |
| Quobyte Volume Plugin                             | A volume plugin that connects Docker to Quobyte's data center file system, a general-purpose scalable and fault-tolerant storage platform.                                                                                                     |

### volume plugin – multi-platform

| plugin             | description                                                                                                                                                                                                                                                                                                                          |
|--------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| OpenStorage Plugin | <p>It is a clustered implementation of the Open Storage specification which relies on the docker runtime. It allows you to run stateful services in Linux Containers in a multi-host environment. It plugs into CSI and Docker volumes in order to provide storage to a container. It also plugs into Kubernetes and Mesosphere.</p> |
| REX-Ray plugin     | <p>A volume plugin that provides a vendor agnostic storage orchestration engine. The primary aim is to provide persistent storage for Mesos, Docker, and Kubernetes.</p>                                                                                                                                                             |

### volume plugin

| plugin                                                   | description                                                                                                                                                                                                                                            |
|----------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Local Persist Plugin<br>(Last source code commit 1 year) | A volume plugin that extends the default <b>localdriver</b> 's functionality by allowing you specify a mountpoint anywhere on the host, which enables the files to <i>always persist</i> , even if the volume is removed via <b>docker volume rm</b> . |

# HARDENING OF CONTAINER INFRASTRUCTURE

Module 05

A man with dark hair and glasses, wearing a light blue striped shirt, is sitting at a desk and looking down at a laptop screen. He appears to be in an office environment. The background is slightly blurred.

# Section 5.1 | Hardening the Docker host

First: Determine how "hard" you need to be, and this will help determine which guidelines you should follow.

- Is this a hobby/trial/learning machine? Lab machine? Business-critical? Military-grade?
- What policies exist: corporate/organizational policy may have specific recommendations or requirements

the CIS (Center for Internet Security) benchmark reference

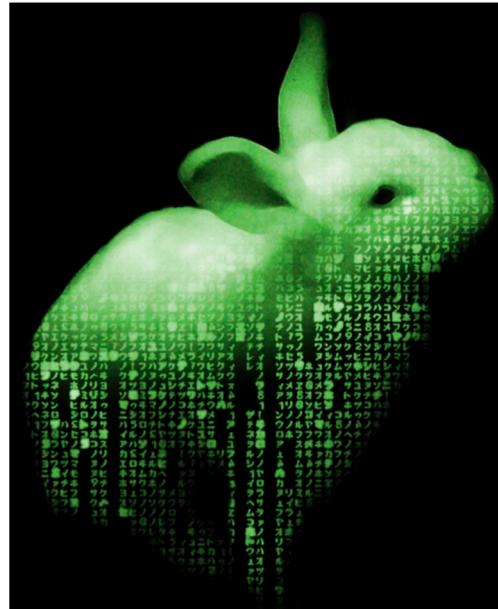
docker's security bible

- Full CIS reference doc [here.](#)



## Section 5.1 | Hardening the Docker host

Follow the white rabbit. . .



is security a function of the >

>>>host

or

>>> process



## Tip: Put Docker filesystem into its own partition

A recommended practice is to create a separate partition for `/var/lib/docker`, and mount this (in `/etc/fstab`), so that Docker doesn't wedge the entire system, should things grow out of control suddenly (or just slowly grow, unobserved, over time.)

## Building Secure Containers

- Verify authenticity of the build
- Assure public images are signed and have valid checksums
- Use HTTPS for remote file uploads
- Pin to specific files
- Always use version images
  - E.g ubuntu:14.04, vs Ubuntu

## Lightweight Images

- Favor lightweight images Alpine is a light weight distro
- Build minimum images to run apps
- Do not run SSH in a container

## Docker Security Scanning

- Provides binary level scanning of Docker image
- Presents a security profile of the image
- Push image to registry -> Docker cloud performs scanning
- Uses CVE Scanning validation service that is updated regularly
- Available in Docker Cloud, Docker EE
- Each vulnerability is associated with a severity
- Use official repos, they've been using it for a while now

Tip: Make sure you are using a modern kernel version.

- v3.10 is the bare minimum required version for Docker to function, but v4.6.1 is current, as of June 1, 2016.
- "Reasonably fresh" is generally better than stale, when it comes to security.
- (Though "too fresh" can lead to new bugs, so it's good to find the sweet spot where code is reasonably well tested, but not getting old and creaky.)

**Tip:** Remove all non-essential services from the Docker host.

- This would include other (non-Docker-related) services:
  - web services
  - database
  - user accounts & services
  - file shares
  - etc.

## Tip: Keep Docker up to date!

- Docker ecosystem is evolving VERY fast
  - months == years
  - major == minor
- Audit "docker version"
- Infrastructure as "Code" anyone

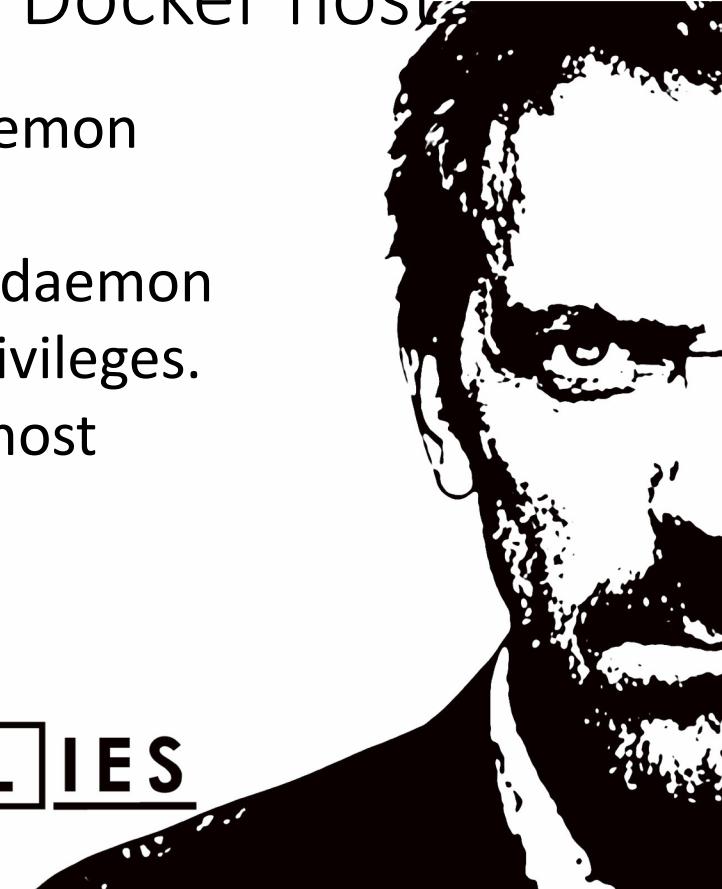


# Section 5.1 | Hardening the Docker host

Tip: Limit access to the Docker daemon

- Remember that the Docker daemon runs with elevated (root) privileges.
- Consider vulnerabilities on host

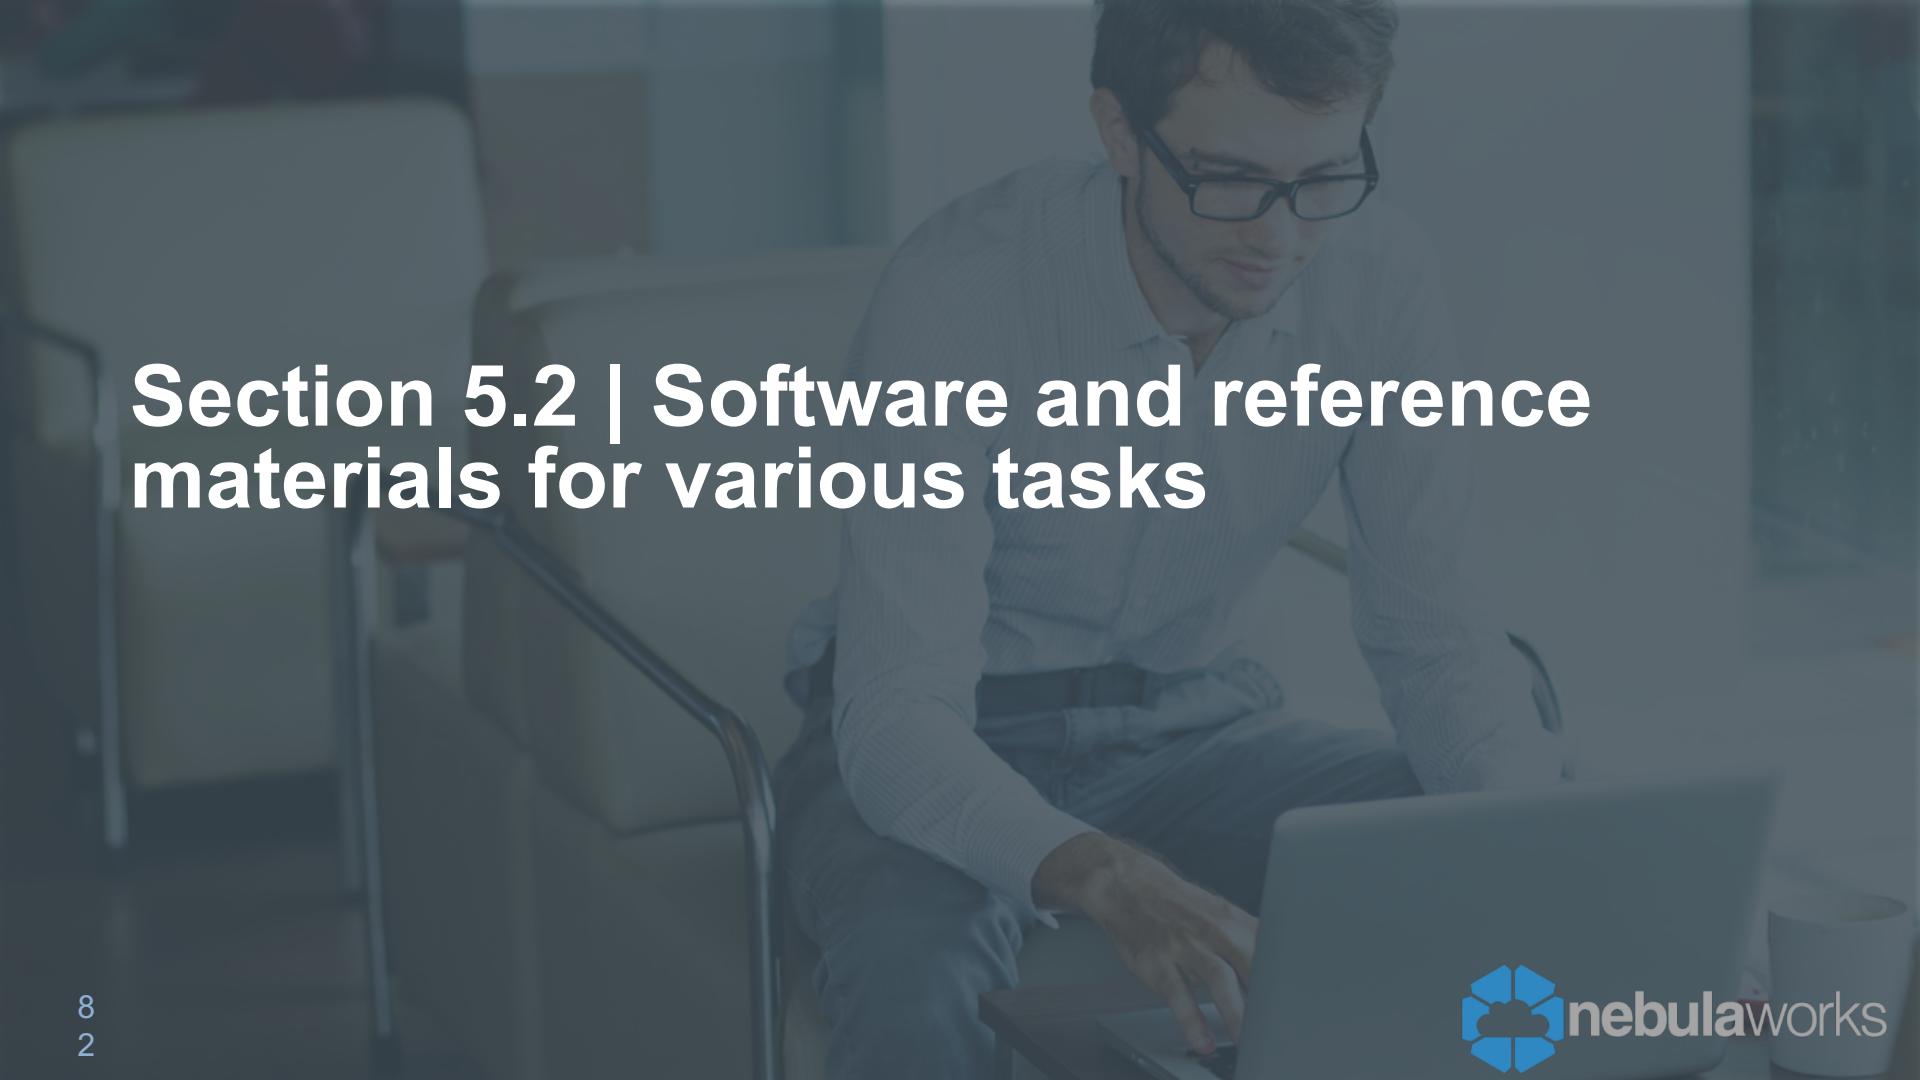
**E**VERYBODY **LIES**



Tip: Use auditd and auditctl to monitor the activity and usage of the Docker daemon, and related files.

### File to watch >>

- /var/lib/docker
- /etc/docker
- /etc/default/docker
- /etc/docker/daemon.json
- /usr/bin/docker-containerd
- /usr/bin/docker-runc
- docker.service
- docker.socket

A professional-looking man with dark hair and glasses, wearing a light blue striped button-down shirt, is seated at a desk in an office environment. He is looking intently at a laptop screen, his hands resting on the keyboard. The background is slightly blurred, showing office cubicles and a staircase. The overall atmosphere is focused and productive.

## Section 5.2 | Software and reference materials for various tasks

## Linux kernel capabilities

For fine-grained control, rather than simplistic root/non-root permissions.

- <http://man7.org/linux/man-pages/man7/capabilities.7.html>
- The "Other kernel security features" section discusses hardening options, and a few useful tool sets, profiles, etc.

## Seccomp - secure computing mode

- used to sandbox processes in the Linux kernel, and isolate them from making any system calls except for a very few approved ones.
- Leverages a whitelist that will deny system calls
- Distribution must support Seccomp
- <https://en.wikipedia.org/wiki/Seccomp>
- <https://docs.docker.com/engine/security/seccomp/>

- AppArmor - A method of restricting and controlling application capabilities on an individual basis, using profiles, and with a "learning mode" to help study usage patterns and generate profiles.
- Security profiles automatically added to containers, not daemon
- Will use default but can be overridden

```
$ docker run --rm -it --security-opt apparmor=docker-default hello-world
```

- This is often touted as a more "friendly-to-implement" alternative to SELinux.
- <https://en.wikipedia.org/wiki/AppArmor>
- <https://docs.docker.com/engine/security/apparmor/>

## SELinux - Security-Enhanced Linux

- NSA and DoD developed guidelines for hardening systems.
- Thorough, though sometimes difficult to administer.
- [https://en.wikipedia.org/wiki/Security-Enhanced\\_Linux](https://en.wikipedia.org/wiki/Security-Enhanced_Linux)

## TLS - Transport Layer Security

- Most commonly seen with `https://` URLs.
- Docker has no inherent built-in security.
- TLS is the recommended system for securing communications with encryption and authentication.
- <https://docs.docker.com/engine/security/https/>

## TLS - Transport Layer Security (cont'd)

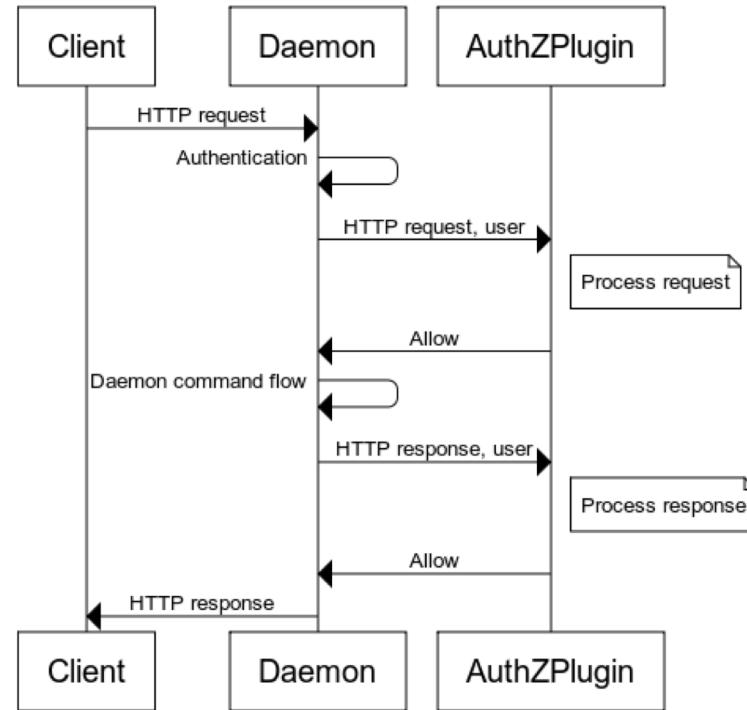
- Don't use self-signed certificates on anything other than a testing machine.
  - No one will trust it, and it will raise lots of warnings/alerts.
  - If your organization is larger than a startup, you probably have someone on staff who deals with certificates and security. Talk to an IT or Security person to track them down, and they can usually provide you with proper certificates.



### Docker authorization & authentication plugin

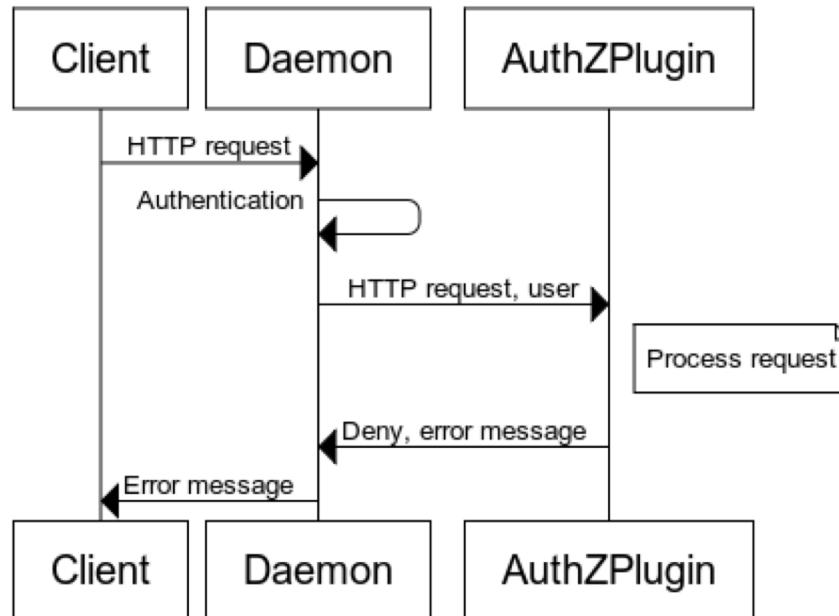
Allow >>>

#### Authorization allow scenario



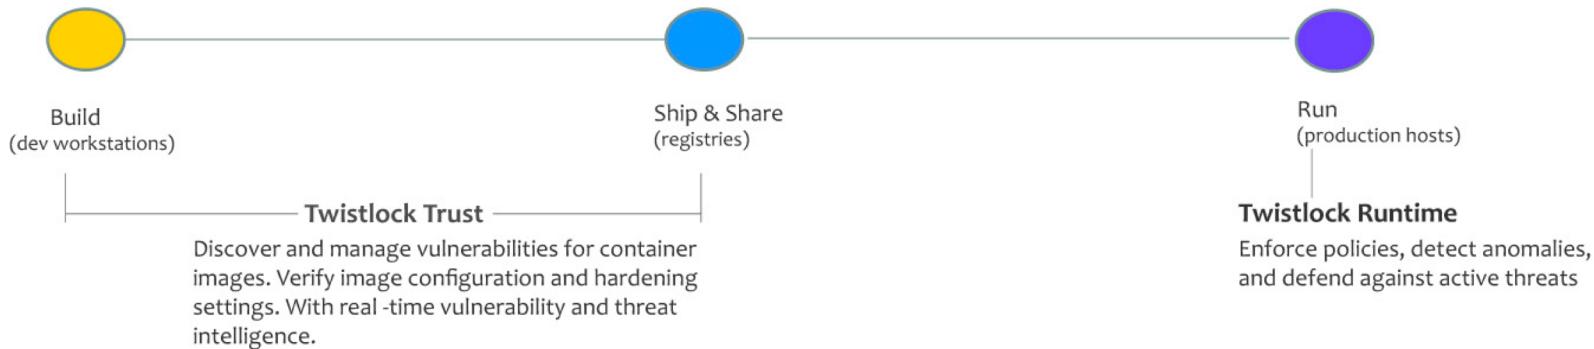
## Docker authorization & authentication plugin

### Authorization deny scenario



### Twistlock Casbin AuthZ HB Plugin

- Design: agentless, full stack, run anywhere, cradle to scale



## Vault by Hashicorp

- Provide single source of secrets for humans and machines
- Data is encrypted at rest as well as in transit
- Manage secret lifecycle
- Revoke credentials in case of attack
- Traceability through audit logs
  - <https://www.hashicorp.com>

# Section 5.3 | Docker daemon configuration

## Tuning the docker engine

- Restrict network traffic between containers
- Set logging levels
- Setting storage drivers
- Configure centralized and remote logging
- Do not use insecure registries
- Allow docker to manage/update iptables
  - (It's probably much better at it than you are. ;-)
- Use TLS authentication



## Isolating containers with a user namespace

- Prevent privilege-escalation attacks from inside a container by configuring your application to run as an unprivileged user
- If your process must run as a root user, you can re-map this user to a less-privileged user on the docker host!
- Added in v1.10 to address root privilege issues:
  - supported directly as an engine option.
  - allows for the root user in a container to be mapped to a non uid-0 user outside the container
  - all container root users are mapped to the same outside user



### How do we map user IDs from the host to the container?

- the remapping itself is handled by two files: `/etc/subgid` and `/etc/subuid`.
- Example of an entry in `/etc/subuid`:  
`user1:241082:65536`
- on the host, especially when you have not explicitly created non-root users for use in the containers.
- Example issue: RedHat has a known issue where automatic mapping of sub-UIDs and sub-GIDs doesn't happen. Manual mapping may be required for now on those systems.

### Audit the user namespace:

- ```
ps -p $(docker inspect --format='{{ .State.Pid }}' <CONTAINER ID>) -o pid,user
```
- The above command would find the PID of the container and then would list the host user associated with the container process. If the container process is running as root, then this recommendation is non-compliant.
- Alternatively, you can run the below command to ensure that the docker daemon is running with --userns-remap flag.
- ```
ps -ef | grep docker
```

### Remediation:

- Please consult Docker documentation for various ways in which this can be configured depending upon your requirements. Your steps might also vary based on platform - For example, on Red Hat, sub-UIDs and sub-GIDs mapping creation does not work automatically. You might have to create your own mapping.
- General repair steps >>>

Step 1: Ensure that the files `` `/etc/subuid` `` and `` `/etc/subgid` `` exist.

- `touch /etc/subuid /etc/subgid`

Step 2: Start the docker daemon with `` `--userns-remap` `` flag

- `docker daemon --userns-remap=default &`

# Section 5.3 | Docker daemon configuration

## Impact:

- User namespace remapping makes quite a few Docker features incompatible and also currently breaks a few functionalities. Check out the Docker documentation and referenced links for details.

## Default Value:

- By default, user namespace is not remapped.

(Above section on user namespaces from CIS Docker 1.11.0 benchmark, section 2.8)

And there's also this, which goes into more detail about kernel capabilities, and advises not relying completely on just user namespaces.

- <https://medium.com/@ewindisch/linux-user-namespaces-might-not-be-secure-enough-a-k-a-subverting-posix-capabilities-f1c4ae19cad#.cwyx9fp0l>

And this video:

Netdev 1.1 - Namespaces and CGroups, the basis of Linux containers

- <https://www.youtube.com/watch?v=zMJD8PJkoYQ> (47.5 minutes)

## Section 5.4 | Trusted content/trusted images

How do you know that an image is safe?

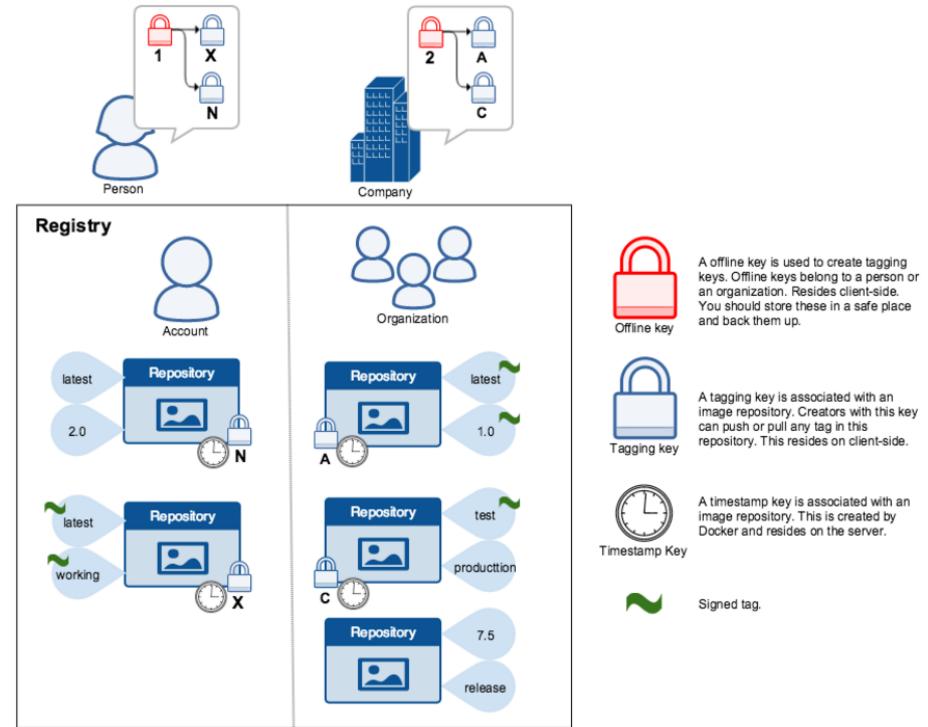
Sure, there's lots of stuff out there, but can it all be trusted?

Is it maintained with the latest patches?

(Not always, but you shouldn't blindly trust, either...)



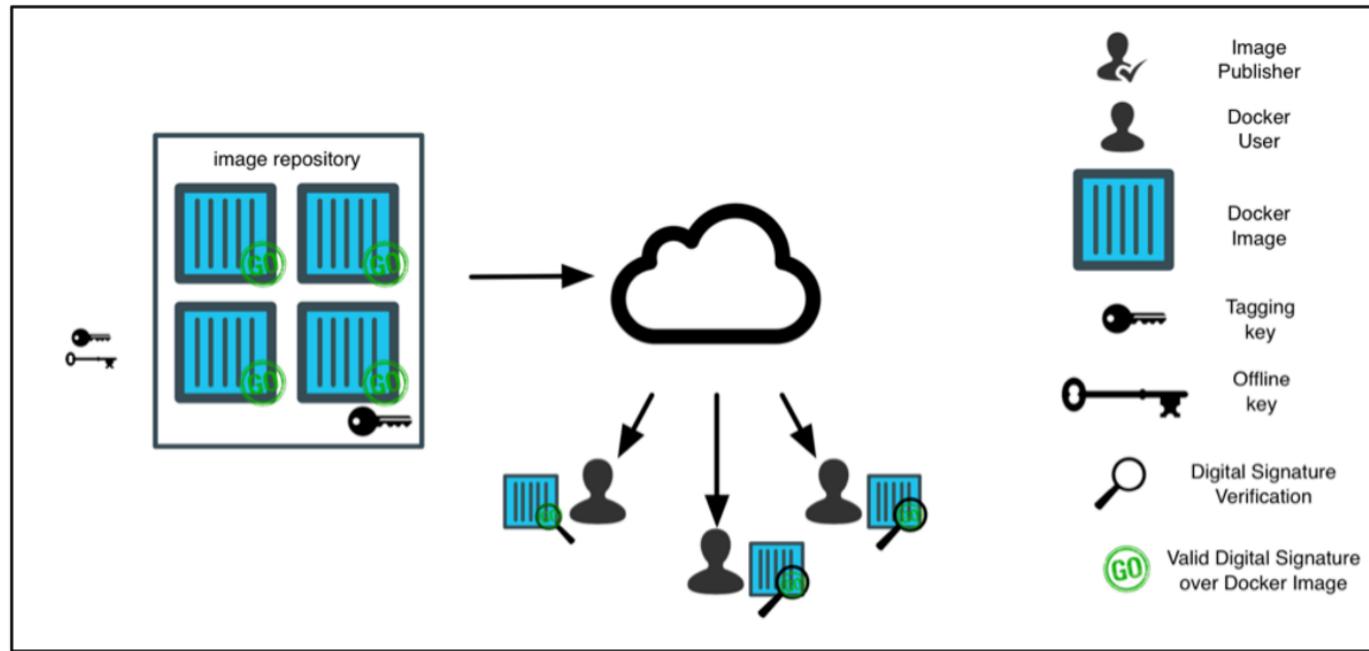
### How images get signed



## How images get signed

- Once Docker Content Trust has been enabled, only signed images can be used
- If you try to pull an image or run a container from an image that is not signed, it will fail
- The first time you pull an image, trust is established to the repository with the offline key
- All subsequent interactions with that image require a valid signature verification from that same publisher
- Once trust is established, TUF will ensure integrity and freshness on the content, via the use of a timestamp key
- Docker uses and manages the timestamp key

## How images get pulled



## Docker Content Trust (Image signing and verification)

- <https://blog.docker.com/2015/08/content-trust-docker-1-8/>

## The Update Framework (TUF)

### Docker Notary project:

- <https://github.com/docker/notary>

### Docker Security Scanning safeguards the container content lifecycle

- <https://blog.docker.com/category/security-2/>

There is also the Banyan study finding 30-40% of images on the Docker Hub having known security vulnerabilities.

- Over 30% of Official Images in Docker Hub Contain High Priority Security Vulnerabilities
- <http://www.banyanops.com/blog/analyzing-docker-hub/>



# Section 5.5 | Thinking outside the VM

# Section 5.5 | Thinking outside the VM

...to unlearn what you have learned is the most necessary kind of learning.

-- Antisthenes



# Section 5.5 | Thinking outside the VM

Don't run sshd in a container

- think in terms of loosely coupled microservices.
- See advice from [Jerome Petazzoni](#)
- Alternatives to ssh : nsenter or docker exec



# Section 5.6 | Reference materials, links, and useful stuff

## Docker security references

- <https://docs.docker.com/engine/security/security/>
- <https://blog.docker.com/2015/05/understanding-docker-security-and-best-practices/>
- [https://www.docker.com/sites/default/files/WP\\_Intro%20to%20container%20security\\_03.20.2015%20%281%29.pdf](https://www.docker.com/sites/default/files/WP_Intro%20to%20container%20security_03.20.2015%20%281%29.pdf)

## CIS Docker Benchmark documentation:

- [https://benchmarks.cisecurity.org/tools2/docker/CIS\\_Docker\\_1.11.0\\_Benchmark\\_v1.0.0.pdf](https://benchmarks.cisecurity.org/tools2/docker/CIS_Docker_1.11.0_Benchmark_v1.0.0.pdf)

## Docker Security Cheat-Sheet

- <http://container-solutions.com/docker-security-cheat-sheet/>

## Trust models in Docker. (There's a LOT on this page.)

- [https://docs.docker.com/engine/security/trust/content\\_trust/](https://docs.docker.com/engine/security/trust/content_trust/)

## Docker Security, Hardened Containers and a Layered Strategy

- <http://blog.flux7.com/docker-security-hardened-containers-and-a-layered-strategy>
- A review of some of the security features introduced in v1.10

## "On the Security of Containers" (2014):

- <https://medium.com/@ewindisch/on-the-security-of-containers-2c60ffe25a9e#.hhuvxhmb6>

## Five security concerns when using Docker

- <https://www.oreilly.com/ideas/five-security-concerns-when-using-docker>
- This is an excerpt from (and has a link to) a larger paper, which is a security-related report/chapter from one of O'Reilly's Docker book authors. (It can be downloaded from their site if you give them your contact info.)

## Docker Security Resources & Best Practices

- <https://www.twistlock.com/docker-security-resources/>

## Understanding and Hardening Linux Containers

- <https://www.nccgroup.trust/us/about-us/newsroom-and-events/blog/2016/april/understanding-and-hardening-linux-containers/>
- 122 page PDF version available

## Linux-Audit: Docker Security: Best Practices for your Vessel and Containers

- <http://linux-audit.com/docker-security-best-practices-for-your-vessel-and-containers/>
- There's quite a bit of good information here. Worth mining.

## Hardening Docker Containers: Disable SUID Programs

- <https://blog.tutum.co/2015/02/03/hardening-containers-disable-suid-programs/>

## Docker: Best Security Practices

- <https://elastic-security.com/2016/04/11/docker-best-security-practices/>

## Docker Bench for Security project

- <https://github.com/docker/docker-bench-security>

Docker crocker-blocker aims at stopping Docker shockers ([Oldy but Goody](#))

- [http://www.theregister.co.uk/2015/05/08/hardening\\_guide\\_a\\_docker\\_shocker\\_blocker/](http://www.theregister.co.uk/2015/05/08/hardening_guide_a_docker_shocker_blocker/)
- Humorous title, from The Register (UK), and an introduction to some of the CIS security benchmarks.
- This article is about a year old, and refers to an older version (1.6) of the security docs, but there are links to the new version for v1.11 listed elsewhere in these slides.:

# RE-IMAGINING INFRASTRUCTURE

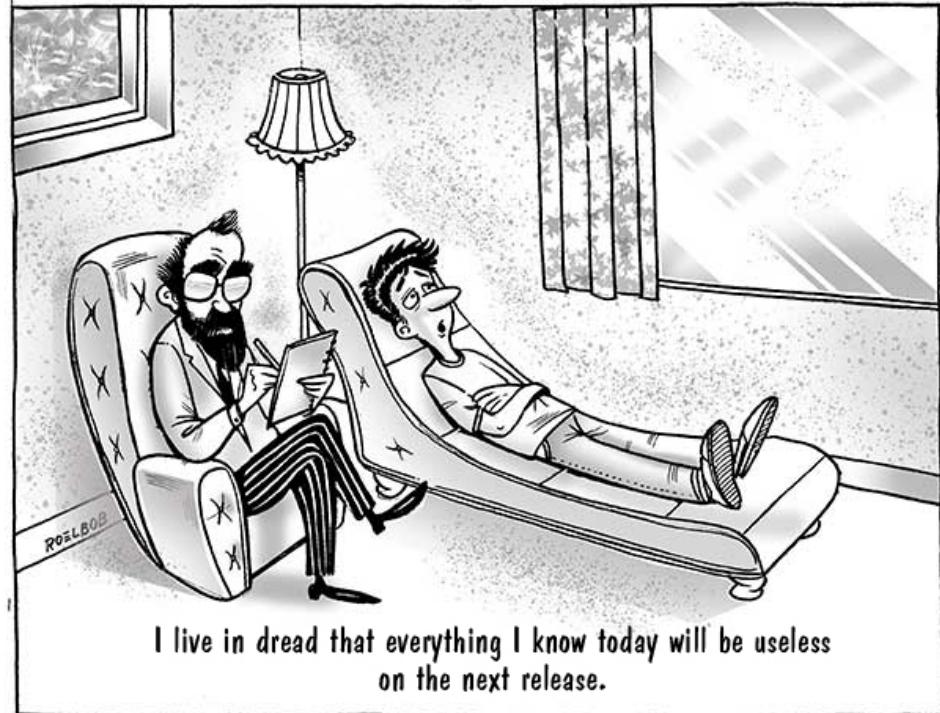
module 06

A man with dark hair and a beard, wearing black-rimmed glasses and a light blue striped button-down shirt, is seated at a desk. He is looking intently at a laptop screen, which is partially visible in the lower right corner of the frame. The background is slightly blurred, showing what appears to be an office environment.

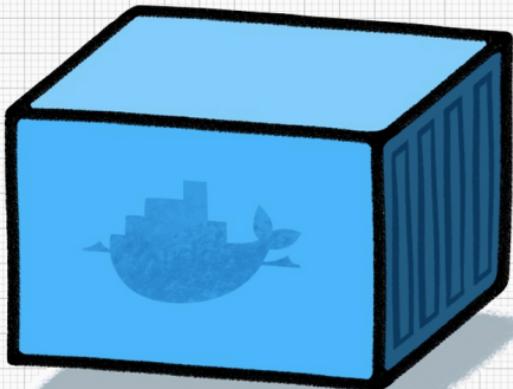
# Section 6.1 | embracing a new artifact

Docker isn't about containers, it's about artifacts.

The benefits of embracing the “container” artifact are realized when you treat them as the universal means for delivering application services.



The real value of Docker is not technology



It's getting people to agree on something

## artifact types

- development
- debug
- test
- production

## development artifact characteristics

- complete dev environment
- able to consume mounted source
- self assembling support services
- tied to build or branch



## debug artifact characteristics

- full debug tool sets
- embedded tests
- auto attaches to feedback systems
- ulimits should conform to production environment



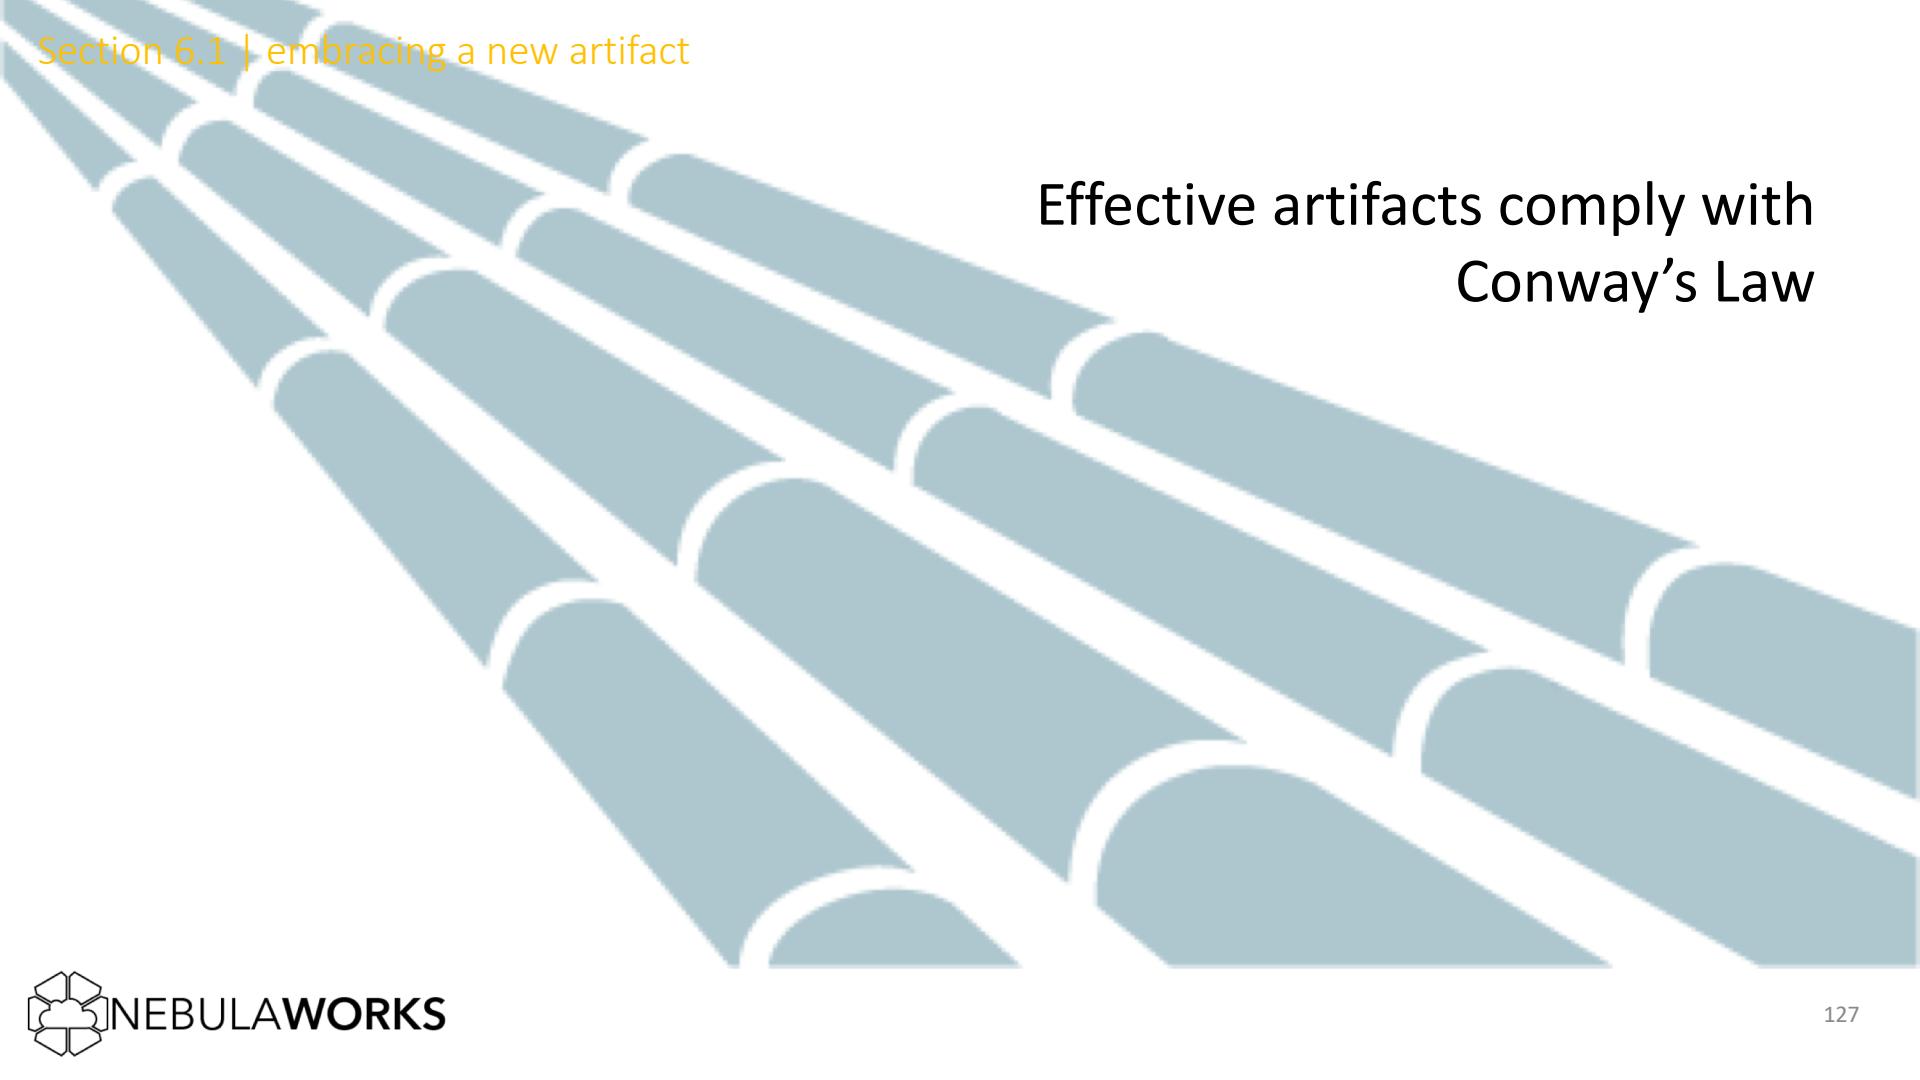
## test container characteristics

- service plumbing
- minimal debug tools
- test frameworks
- feedback hooks



## production artifact characteristics

- lean
- composable
- auto configuring
- minimal attack surface

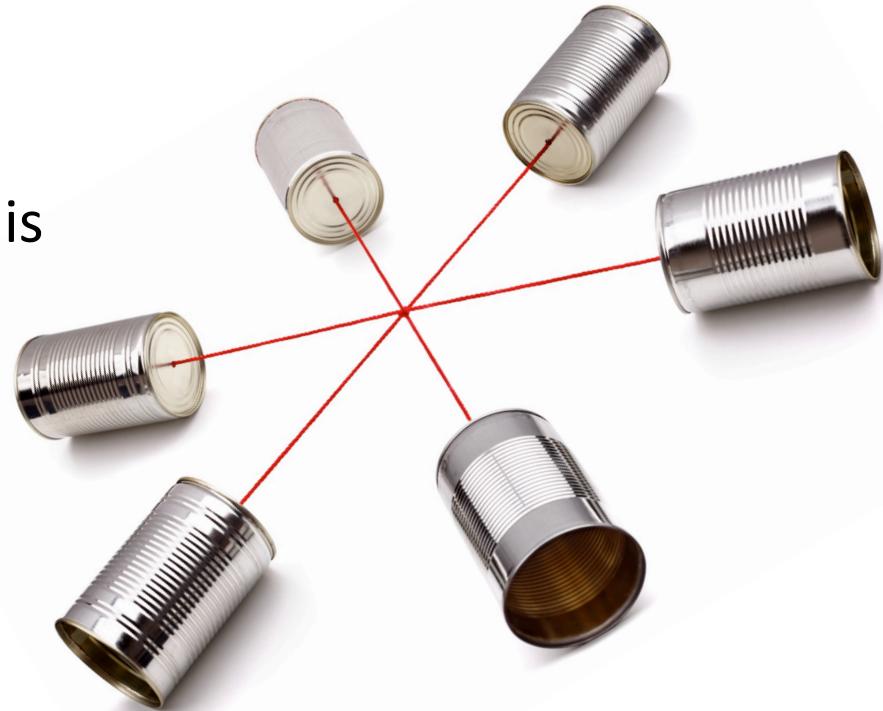


Effective artifacts comply with  
Conway's Law



Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure.

-- Melvyn Conway, 1967



# Section 6.2 | application logistics

applying logistics as a metaphor for IT servicing of application products. this encourages designing an automated JIT ("just in time") process for manufacturing and delivery applications to production end points.



artifacts need hubs



## hubs provide >>>

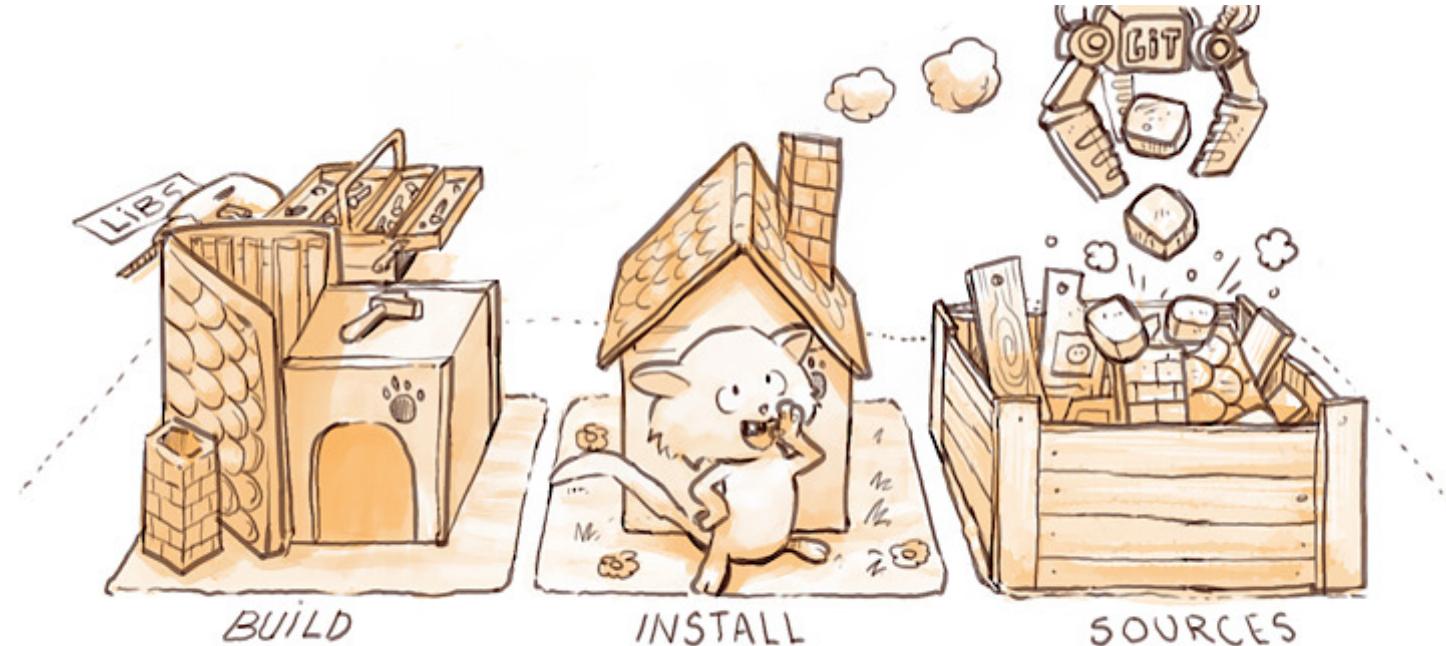
- version control
- garbage collection
- replication facilities



## hub types >>>

- source files
- container images
- automation tools

## artifacts need integrated application pipelines



cc-by deevad

## integrated application pipelines

- automated
- closed
- scalable

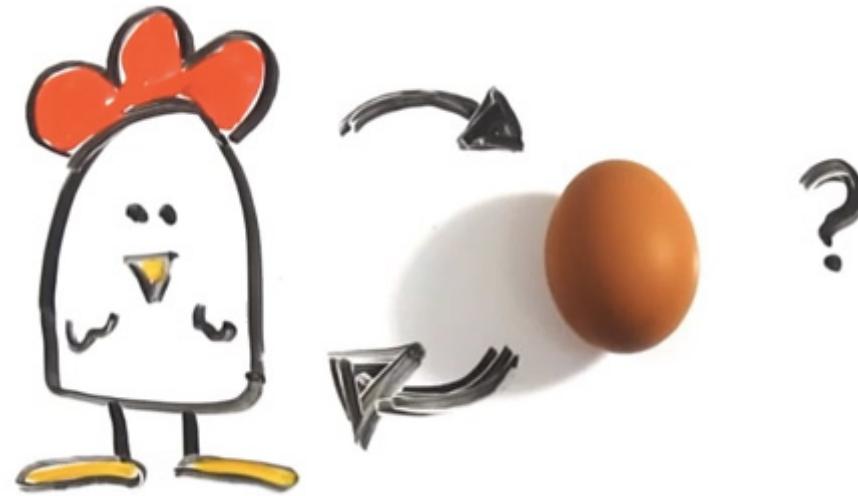


## how does agile development fit in

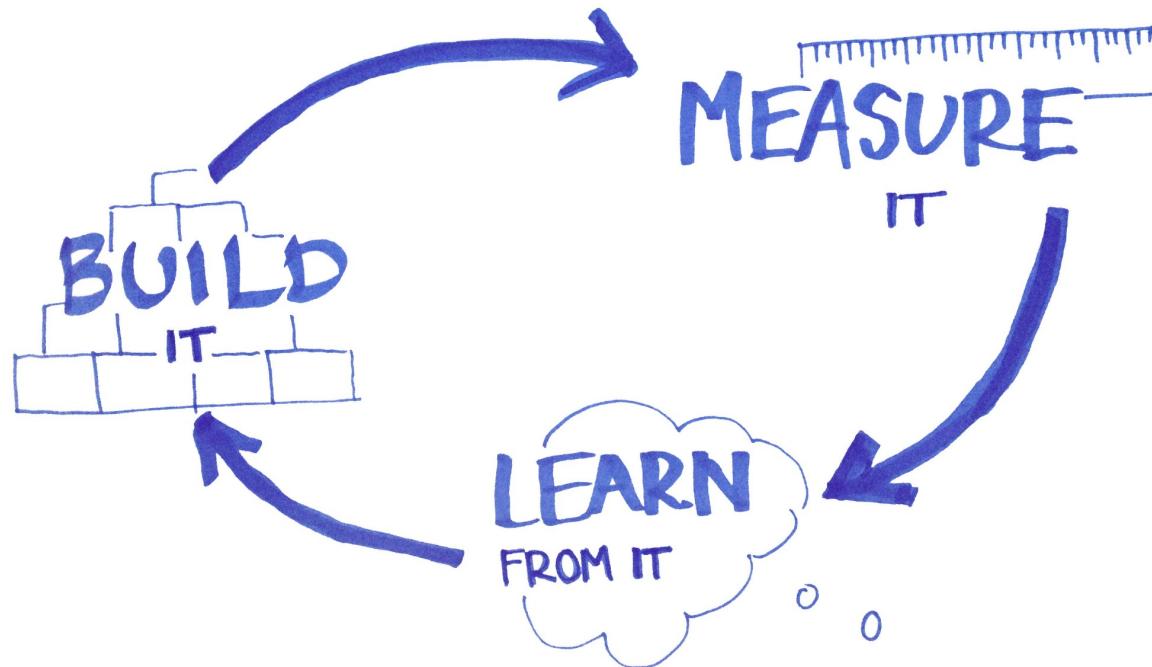


## Section 6.3 | containerization process

## Section 6.3 | containerization process



## Section 6.3 | containerization process





## restructuring services

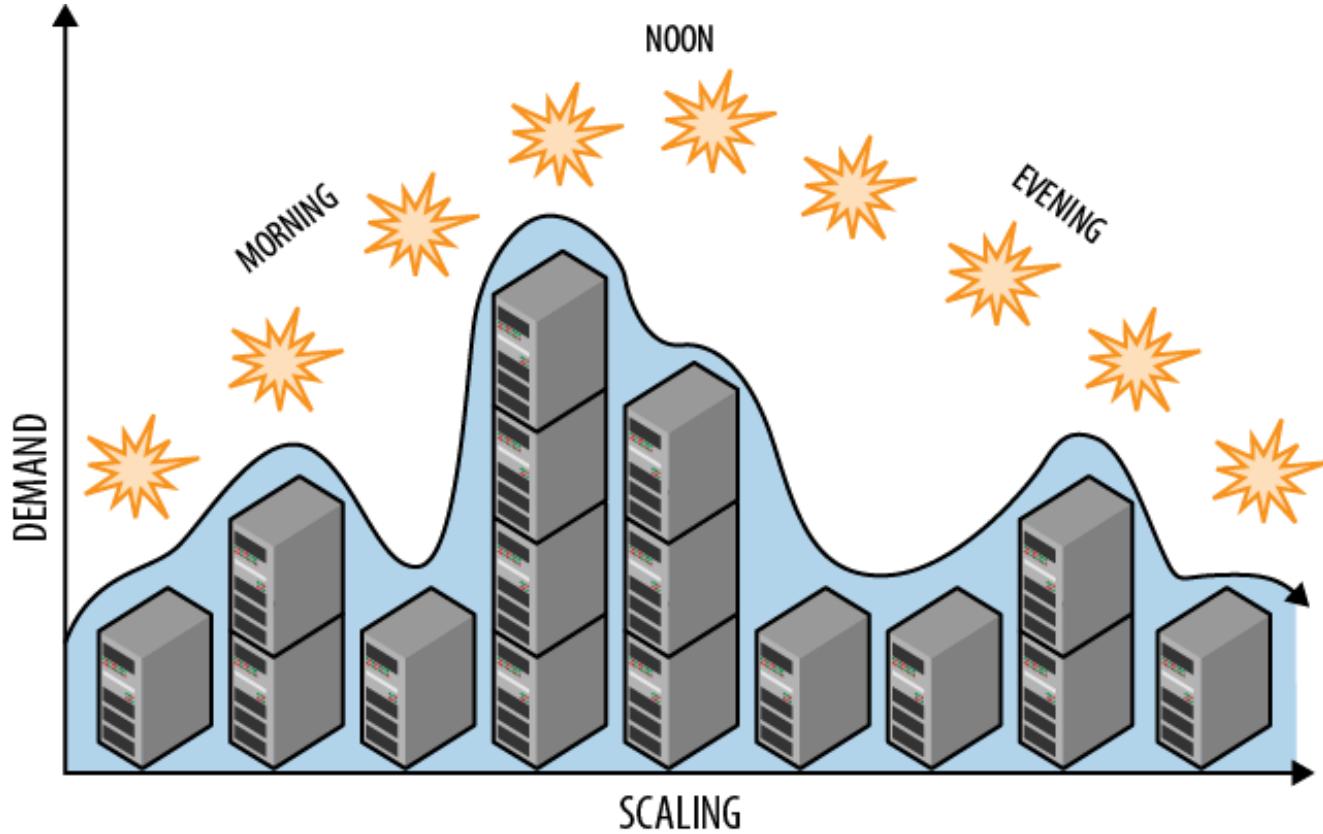
- networking
- logging
- process management
- discovery
- persistence

containers must be everywhere



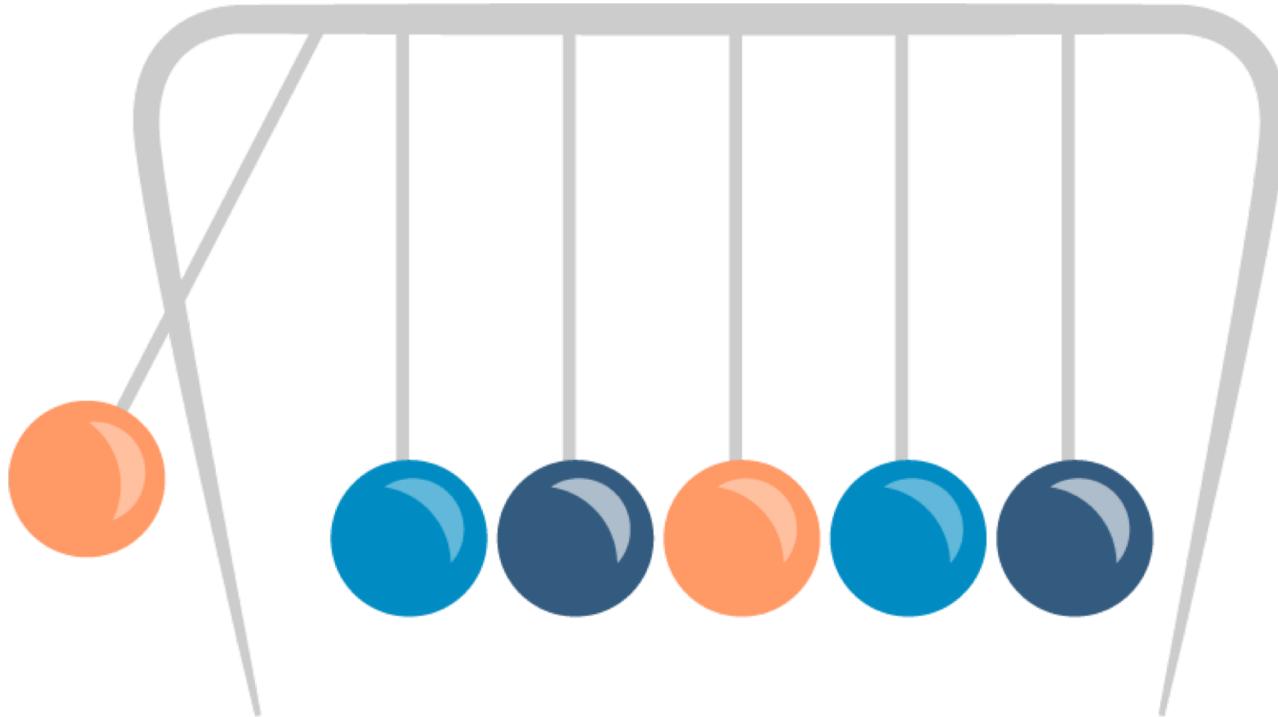
## Section 6.4 | a question of scale

## Section 6.4 | a question of scale

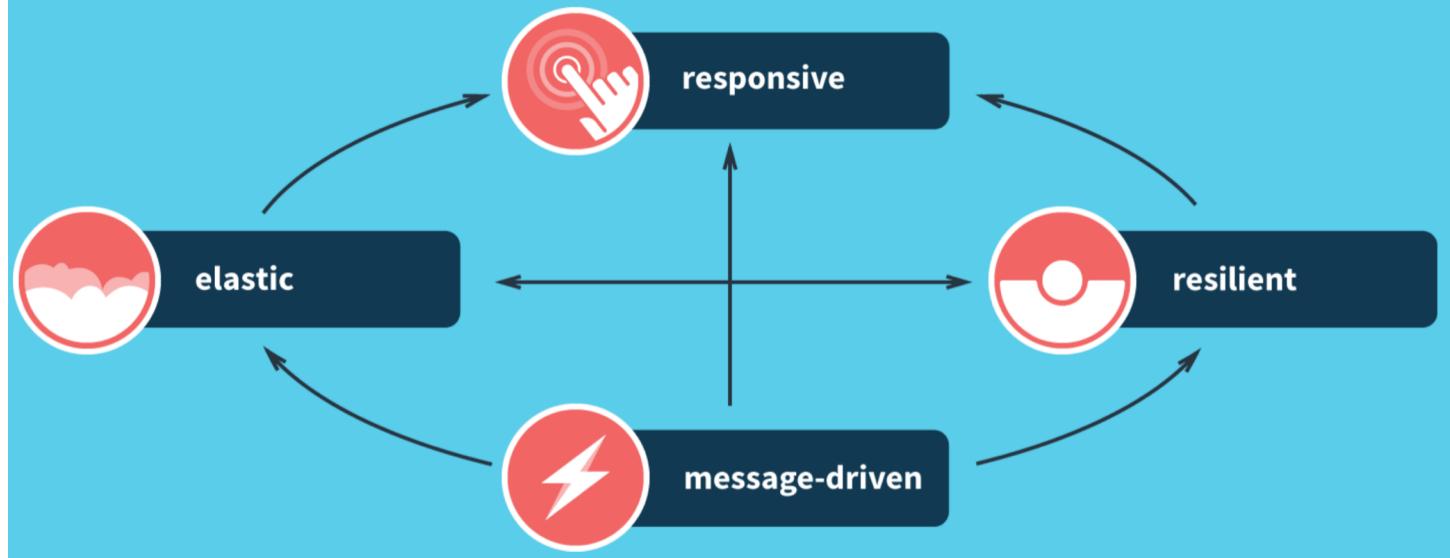


## Section 6.4 | a question of scale

REACTIVE



# Reactive applications share four traits



responsive

- timely
- consistent

resilient

- highly available
- ensured by replication

message driven

- asynchronous message-passing
- loosely coupled

elastic

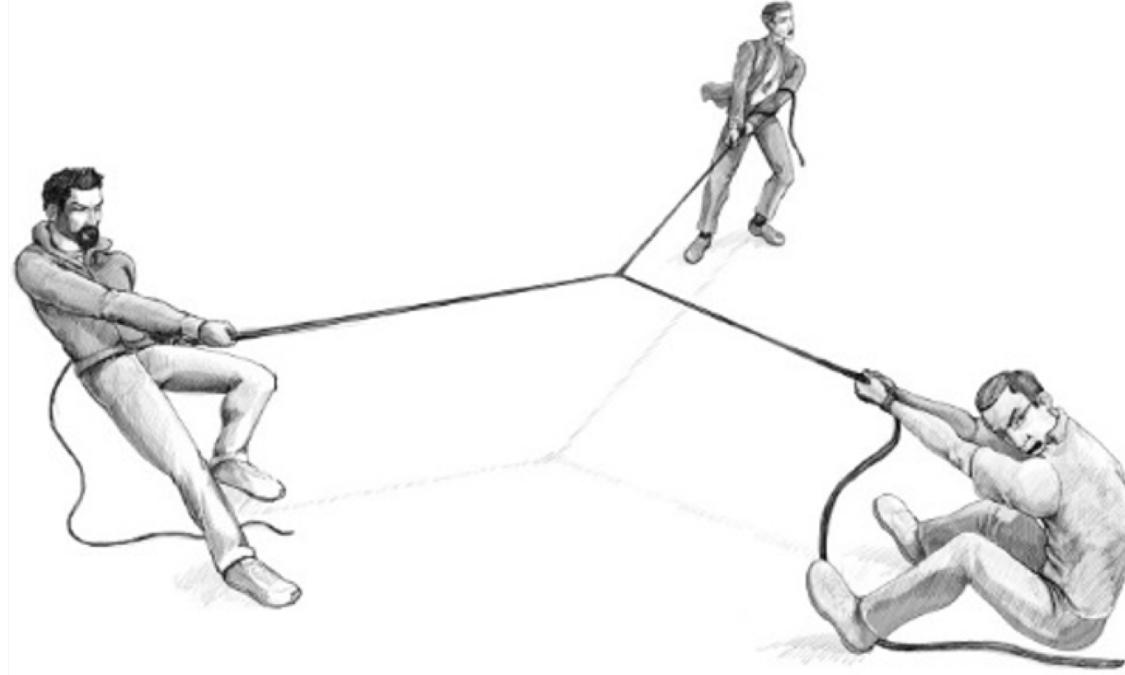
- auto-scaling

• demand driven



Consistency

CAP theorem



Availability

Partitioning



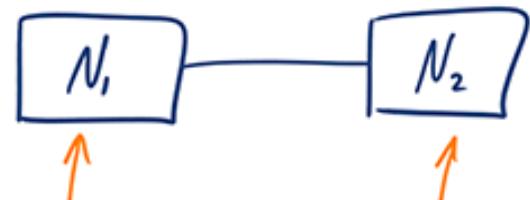
NEBULAWORKS

## CAP theorem explained

**Consistency**



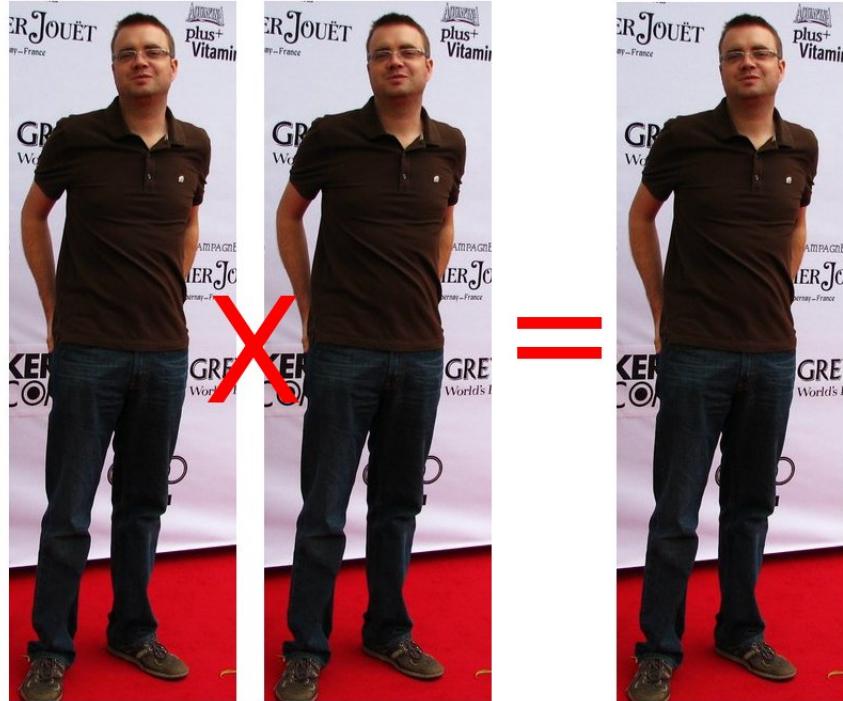
**Availability**



**Partition Tolerance**



idempotent  
&  
immutable



# Rebuild vs Upgrade



Full stack  
Versioning  
& idempotency

Patch

vs  
immutable

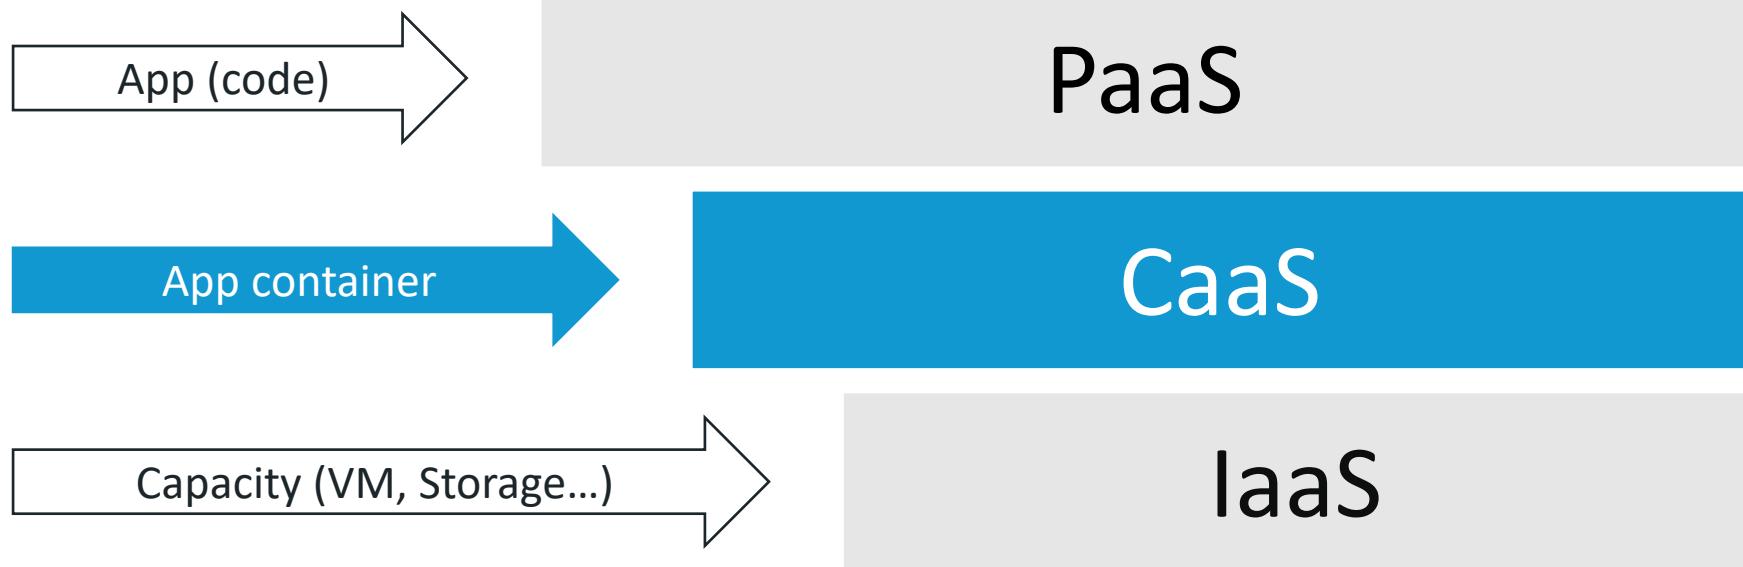


vs

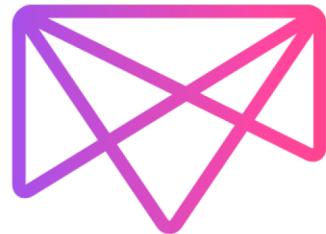
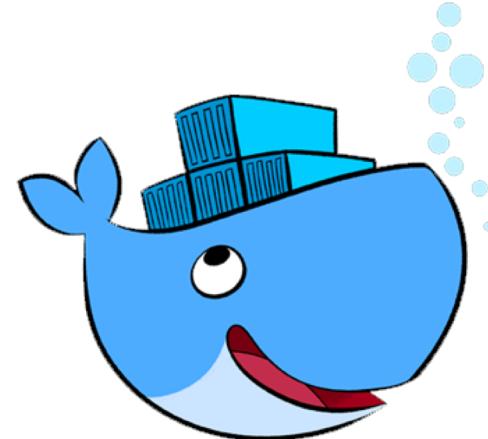
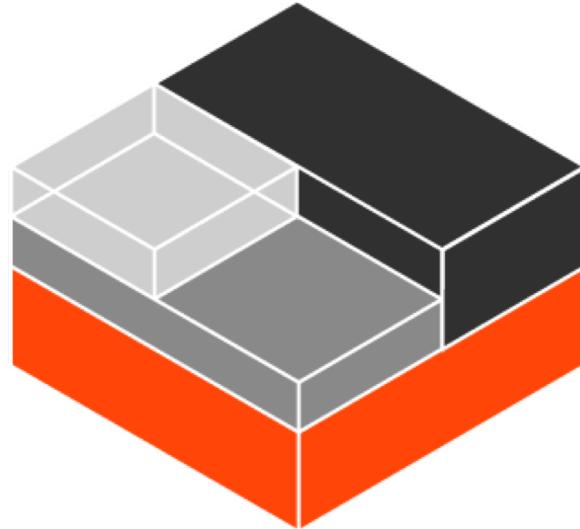
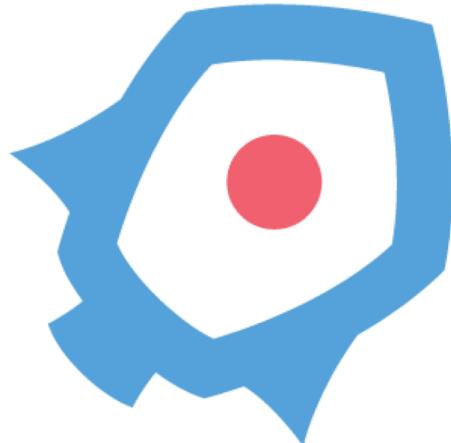


A professional-looking man with dark hair and a beard, wearing black-rimmed glasses and a light blue striped button-down shirt, is seated at a desk in an office environment. He is looking intently at a laptop screen in front of him. The background is slightly blurred, showing office cubicles and other office elements.

## Section 6.5 | platform characteristics



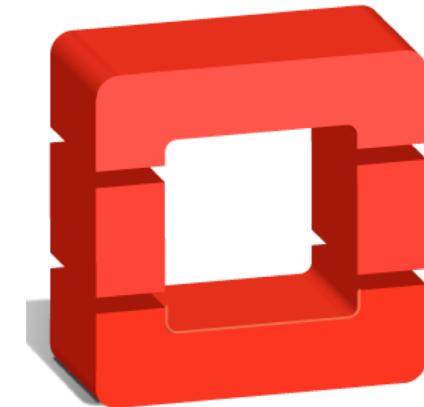
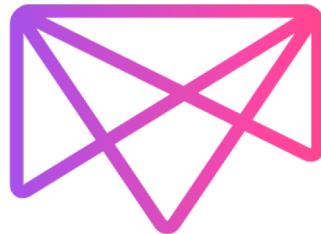
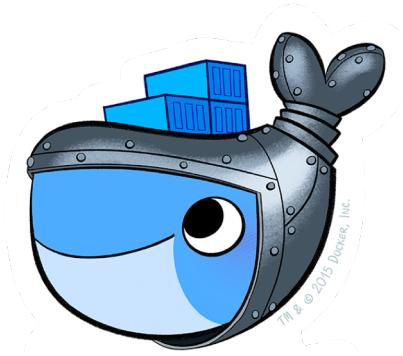
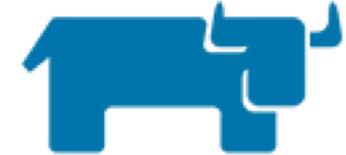
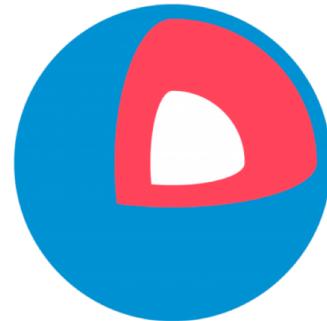
## Section 6.5 | platform characteristics



## Section 6.5 | platform characteristics

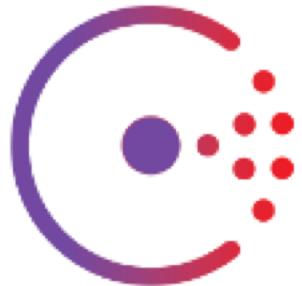


## Section 6.5 | platform characteristics



NEBULAWORKS

## Section 6.5 | platform characteristics



## A quick demo of gossip protocol in action

- Details

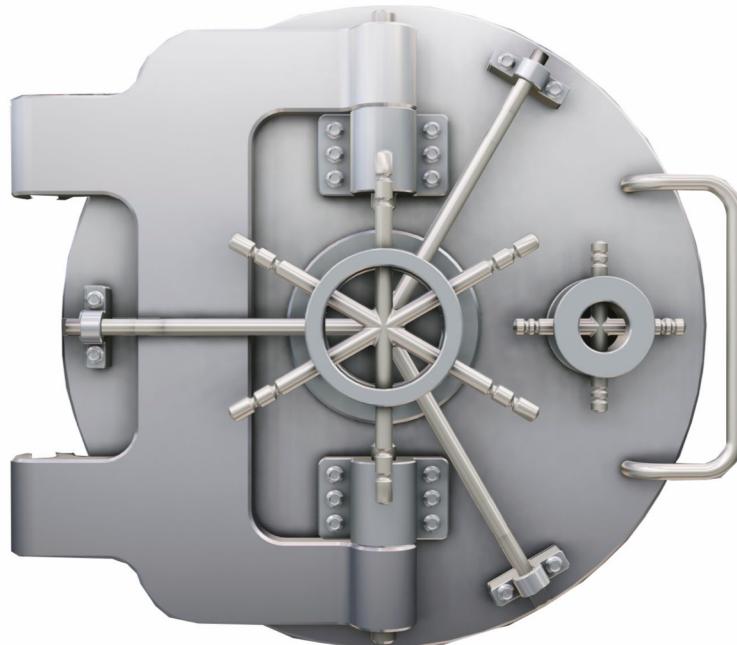
<https://www.serf.io/docs/internals/gossip.html>

- Convergence simulator:

<https://www.serf.io/docs/internals/simulator.html>

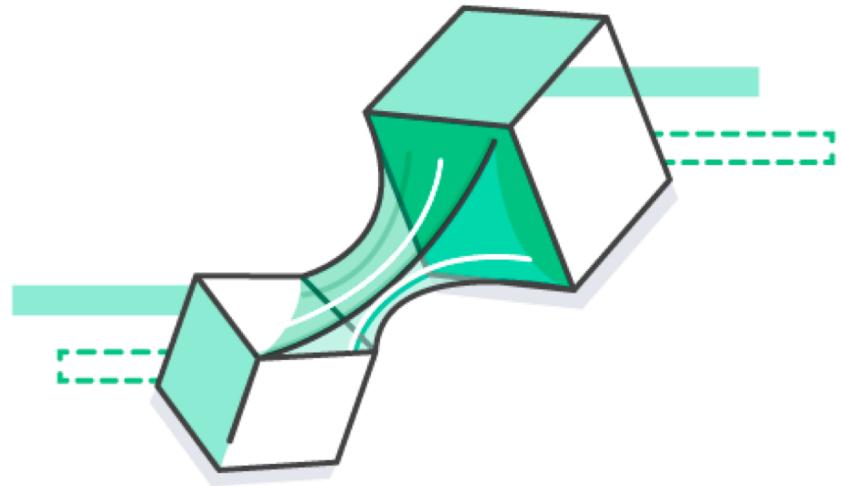
## secrets manager

- secure storage
- key rolling (lease)
- detailed audit trail
- secret revocation
- dynamic secrets

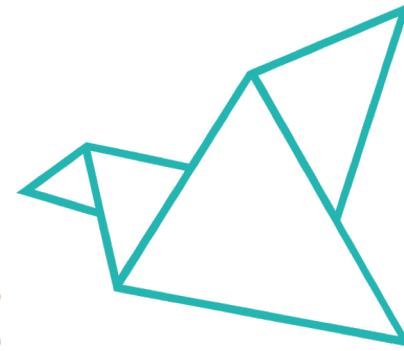


elastic storage

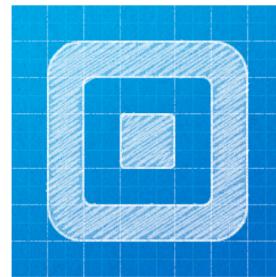
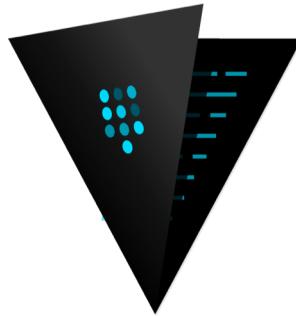
- persistent
- independent
- highly available



## Section 6.5 | platform characteristics

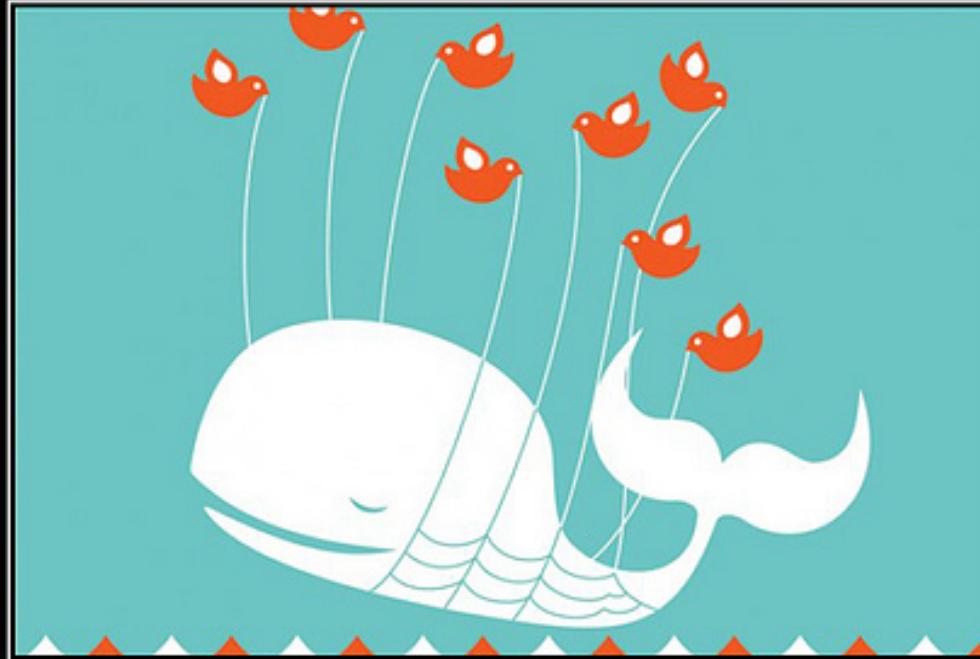


## OPTIONS



- Vault by Hashicorp
- Keywhiz by Square
- Amazon KSM (Key Storage Management)
- Competent by Lyft





# FAIL WHALE

Twitter: Failure is an option. At least once a day, or whenever you need it.

# END OF DAY 2