

# WELCOME TO DAY 3

# **OPERATING IN A CAAS**

module 07

A man with dark hair and a beard, wearing black-rimmed glasses and a light blue striped button-down shirt, is seated at a desk. He is looking intently at a laptop screen, which is partially visible in the lower right corner of the frame. The background is slightly blurred, showing what appears to be an office environment.

# Section 7.1 | app design considerations

# SOFTWARE ARCHITECTURE

## 1990's

SPAGHETTI-ORIENTED  
ARCHITECTURE  
(aka Copy & Paste)



## 2000's

LASAGNA-ORIENTED  
ARCHITECTURE  
(aka Layered Monolith)



## 2010's

RAVIOLI-ORIENTED  
ARCHITECTURE  
(aka Microservices)



## WHAT'S NEXT?

PROBABLY PIZZA-ORIENTED ARCHITECTURE

> where is the app on the spectrum?

you need borders and tools<

## the devops border

Eliminating the blame game of  
"it works on my machine."

- should be clear
- contain feedback mechanisms
- should be end-to-end



# the tools team

- script ninja
  - infrastructure plumber
  - performance magician
  - analytics & detective





<- Netflix SimianArmy logo

- <https://github.com/Netflix/SimianArmy>

## Chaos Engineering

- failure is unavoidable
- it must be tested
- it must be handled gracefully

## Section 7.2 | Docker networking



container models are different

>>>WARNING<<<  
containers are not VMs

you have a choice...  
... of network models



### Networking Fundamentals

- Container Network Model (CNM)
  - Open source
  - Docker implemented libnetwork (written in go) based on CNM
  - libnetwork (OSS) includes: service discovery, basic loading balancing for containers
- Container Network Interface (CNI)
  - Used by rkt
- Docker supports container-to-container networks, as well as connecting to existing networks as well as connecting to existing networks and VLANs

## Container Network Interface

- Came from CoreOS and used as part of Rkt
- Simple interface between container runtime and network implementation
- This is now maintained by the Cloud Native Computing Foundation (CNCF)

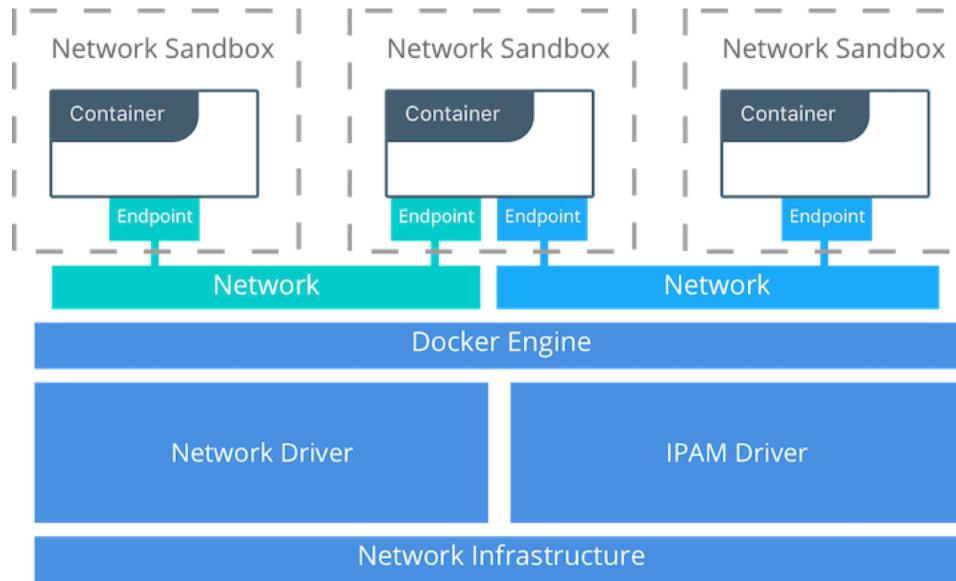
## Docker Networking Fundamentals

- CNM constructs
  - Sandbox
    - Contains the configuration of a containers network stack
    - Isolated network stack that includes
      - Routing tables, ports, interfaces, and DNS Config
  - Endpoint
    - An endpoint joins a sandbox to a Network
    - The connection to the network is abstracted
  - Network
    - Software implementation of a bridge/switch
    - Group together and isolate a collection of endpoints



## Section 7.2 | Docker networking

### Container Networking Model



Credit: Mark Church

thanks mark!

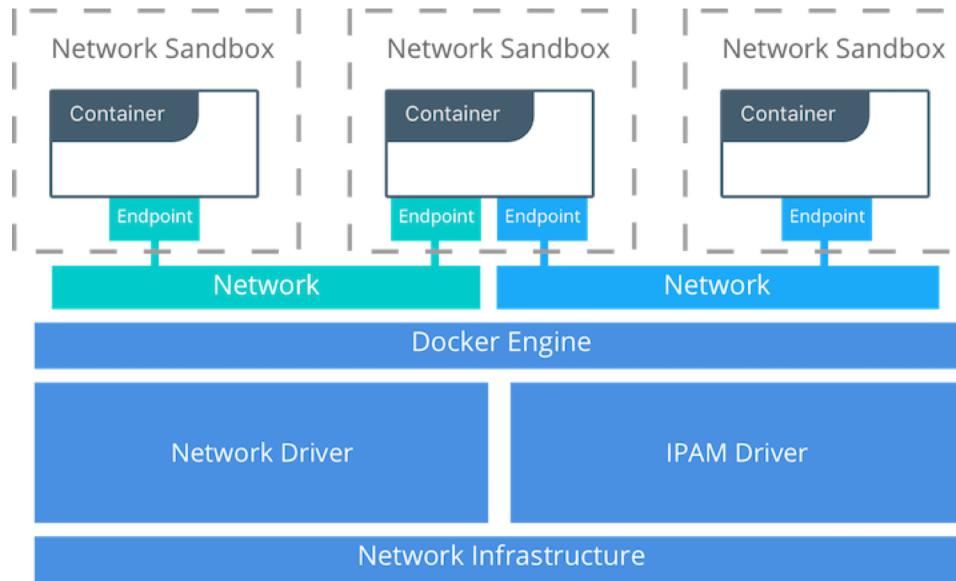
### Docker Networking Design

#### CNM Driver Interfaces

- CNM provides two interfaces for the OSS community, and 3<sup>rd</sup> party vendors to add functionality
  - Remote Drivers
    - Vendor/3<sup>rd</sup> party plugins
  - Native Drivers
    - Responsible to make networks work, and can support various drivers
      - Out of the box drivers can be used e.g bridge, overlay
- IPAM (IP Address Management Driver)
  - This driver provides default IP addresses or subnets for endpoints and networks
  - Leverage docker network, container service commands to manually assign IP addresses

## Section 7.2 | Docker networking

### Container Networking Model



Credit: Mark Church

thanks mark!

## Docker Networking

- We can use networking plugins to extend functionality
- Docker supports some common cases
  - Single host bridge networks
  - Connecting to VLANs
  - Multi-host overlays

## Docker Network Drivers

### Linux

- Bridge (layer 2 switch)
- Overlay
- Macvlan
- Access to network plugins

### Windows

- Nat
- Overlay
- Transparent
- I2bridge

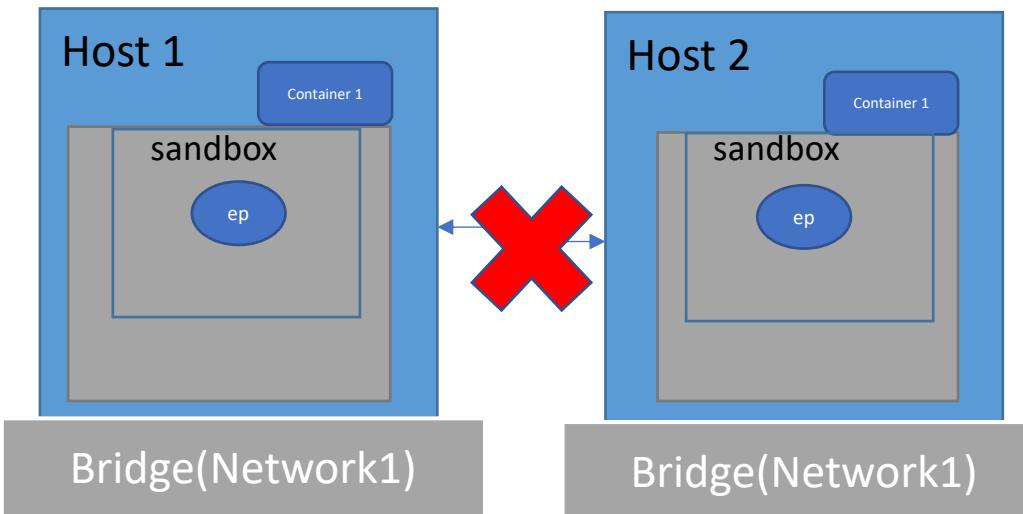
## docker network command

- docker network connect
- docker network create
- docker network disconnect
- docker network inspect
- docker network ls
- docker network prune
- docker network rm

# Lab 4

### Docker Network Drivers

- Bridge
  - By default docker uses the **docker0** bridge
  - Can create custom bridge networks



### Overlay Networks

- Bridge cannot create a distributed network to connect to multiple daemon hosts
- When a swarm is created..
  - Upon either an initialization or a join, two networks are created on the Docker host
    - **ingress** overlay network that manages data traffic related to swarm services (swarm mesh)
      - This is the default network used when a service is created
    - **docker\_gwbridge** that connects a daemon to other daemons part of the swarm

## Docker Swarm Mode

- Cluster management integrated with the Docker Engine  
    \$ docker swarm init
- Supports declarative **service** (more on this later) manifests
- Enables scaling
  - Define number of tasks you want to run
  - Feedback loops in manager monitors that desired state is maintained
- Service Discovery
  - Services are assigned DNS names and performs load balancing
- Each swarm node member enforces TLS authentication to encrypt communication between other nodes and itself
- Uses an overlay network to for distributed nodes

## Swarm Details

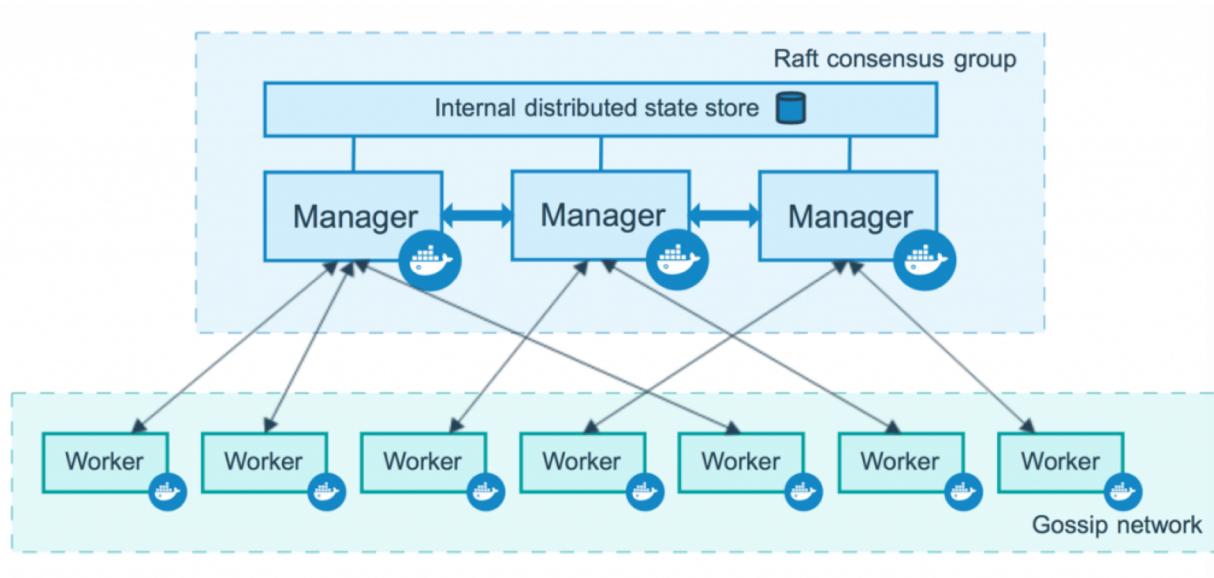
- Manager – Worker
- Docker needs a quorum (majority) for managers
  - The minimum number of hosts that a consensus group needs in order to be allowed to perform an operation
  - If you do not meet cluster quorum requirements workloads should not work
  - $(N/2) + 1$
  - Only use odd numbers for managers

| Managers | Quorum | FT |
|----------|--------|----|
| 1        | 1      | 0  |
| 2        | 2      | 0  |
| 3        | 2      | 1  |
| 4        | 3      | 1  |
| 5        | 3      | 2  |
| 6        | 4      | 2  |



## Section 7.2 | Docker networking

### Docker Swarm Mode



Source : Docker Documentation

## Docker Swarm- Raft Protocol

- Consensus algorithm
- Leader (manager) election
- Log replication
- Low learning curve (vs something like Paxos)
- You should be able to easily learn the protocol, if you wanted to (more simple than Paxos)
- Etcd, the key value store is backed by Raft
- Raft is the algorithm that manages the global cluster state

## Creating a Swarm

`docker swarm init`

- Switches the current node into swarm mode,
- Creates a swarm named **default**
- Configures the current node to be the leader manager in the swarm.
- Starts the internal distributed data store for Engine participating in the swarm.
- Generates self-signed root CA for internal swarm communication (default).
- Generates join-tokens for worker and manager nodes!

## Joining nodes to a swarm

- Joining as a worker

```
$ docker swarm join --token <worker-token>
```

- Joining as a manager

```
$ docker swarm join --token <manager-token>
```

- Retrieving join-tokens on your lead manager node

```
$ docker swarm join-token worker (manager)
```

## Locking Up a Swarm Cluster

- Despite swarms default out-of-the box security settings, a cluster can still be compromised if an old manager rejoins the cluster.
- Old managers joining a cluster will automatically decrypt and gain access to the raft log time-series db. They could then restore old backups and wipe the current swarm configuration (which would make the cluster unusable).
- In order to prevent scenarios like one described above, docker swarm gives you the option to **lock** your swarm cluster
- Locking a swarm cluster forces managers that been restarted to present a cluster-unlock key before they are allowed back into the cluster.

## Locking Up a Swarm Cluster

- You can apply a lock to your swarm at any time.
- When it is first created

```
docker swarm init --autolock
```

- Any point after creation you can also update your swarm

```
docker swarm update --autolock=true
```

- If you don't want the autolock feature remove it by updating your cluster.

```
docker swarm update --autolock=false
```

## Locking Up a Swarm Cluster

- It's important that you keep the cluster key somewhere safe. If you do lose your key, and you need to reconnect a manager to the cluster, you can view the key from an existing manager in your cluster. However, if the key has been rotated since your manager left the swarm, or your swarm cluster has fallen out of quorum then you're out of luck.
- Displaying the unlock-key (from a manager that is currently a part of the swarm cluster)

```
$ docker swarm unlock-key
```

- You also have the option to rotate your unlock keys (you should)

```
$ docker swarm unlock-key --rotate
```

## docker service (utilizing your swarm cluster)

- Goes hand-in-hand with Docker swarm
- It is a higher-level construct that wraps some advanced features around containers.
  - The ability to scale
  - Rolling updates
  - Roll backs
- Uses a declarative model
  - YOU define the desired state of your service and swarm will maintain this desired state.

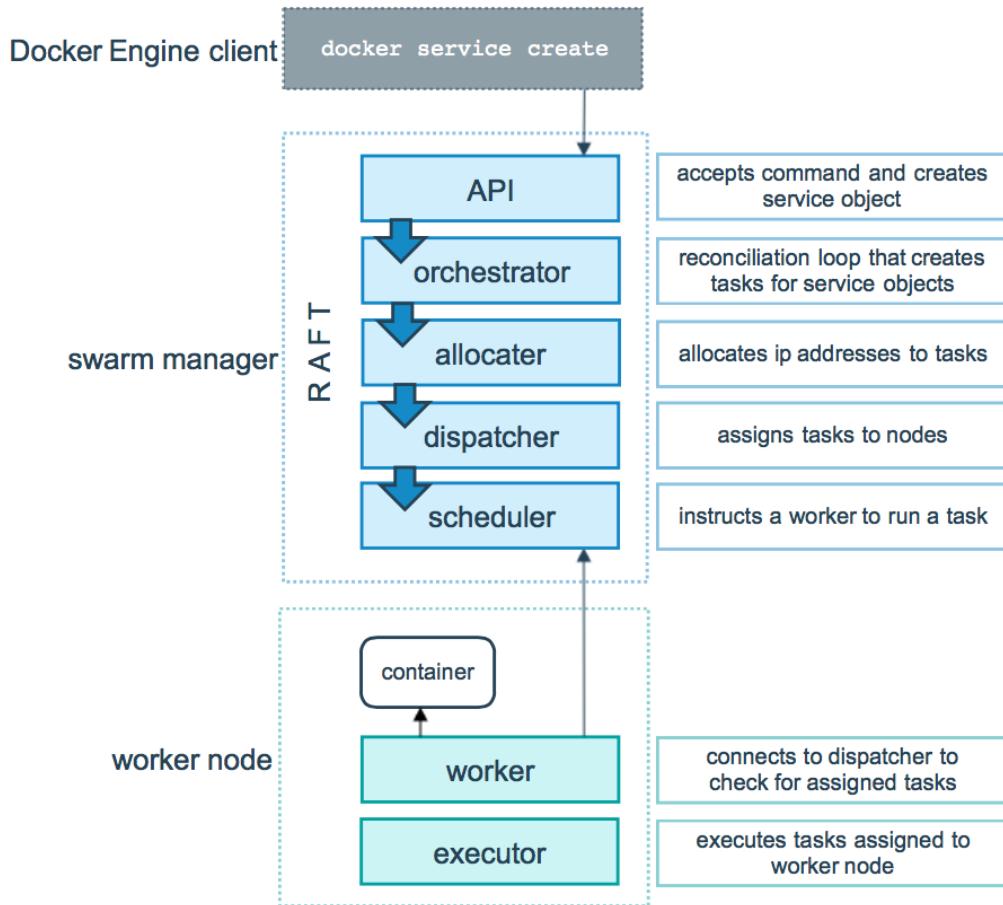
## Tasks

- The atomic unit of scheduling within a swarm.
- You can think of them as replicas or instances of your application.
- When you specify a desired state for your service (in your `docker service create` command), the orchestrator will realize this state by scheduling tasks.
- One-directional mechanism
  - Tasks progress through a series of states : assigned, prepared, running, etc.
  - If the tasks fails, the scheduler will remove it and start a new one!



## Section 7.2 | Docker networking

How the swarm mode accepts service create requests and schedules tasks to worker nodes.



## Types of Service Deployments

### Replicated

- For a replicated service, you specify the number of tasks you want to run across your cluster. This is the default type of service deployment.

### Global

- A global service deployment will run one task on every node in the cluster.

You can configure the service mode by using the `--mode` flag when creating a service

```
docker service create --name my-service --replicas 4 --mode  
replicated (global) <image_name>
```

### Scaling your Service

- There are two cli commands that will let you scale a service, with one having slightly more functionality than the other.
- The `docker service update` command allows you to scale one service at a time.

```
$ docker service update replicas=4 <service_name>
```

- The `docker service scale` command allows you to scale multiple services at the same time.

```
$ docker service scale <service1-ID>=<#-of-tasks> <service2-id>=<#-of-tasks>
```

## Updating your Service

- There are various ways you can update your service. One of which, is applying a rolling update to your service. This can be used if you want your service to use a new-and-improved image, and you want to update it gracefully.
- You configure your rolling update policy at service deployment time.

```
$ docker service create \  
  --replicas 10 \  
  --name redis \  
  --update-delay 10s \  
  --update-parallelism 2 \  
  redis:3.05
```

- Our rolling update policy is modified by the update-policy-related flags we pass in!

## Updating your Service

- The `update-delay` flag configures the time delay between updates to a service task or set of tasks.
- The `update-parallelism` flag configures the max number of service tasks that can be updated at the same time
- Now we can update our service to use the latest redis image.

```
$ docker service update --image redis:3.07 redis
```

## Networking – Docker Swarm and ingress networks

- Docker lets you configure the **ingress** with docker version 17.05 and higher
- If the subnet conflicts with an existing network, you can recreate the **ingress** network

```
$ docker network rm ingress
```

WARNING! Before removing the routing-mesh network, make sure all nodes in your swarm run the same docker engine version. Otherwise, removal may not be effective and functionality of newly created ingress networks will be impaired.

Are you sure want to continue? [y/N]

- Then create a new network
- Can only have 1 **ingress** network per host

## Networking – Customize the docker\_gwbridge

- This virtual bridge is what connects the ingress and overlay networks to a Docker daemon's physical network
- Docker\_gwbridge is a kernel feature on the Docker host (it is not a Docker device)

## Docker Swarm Mode

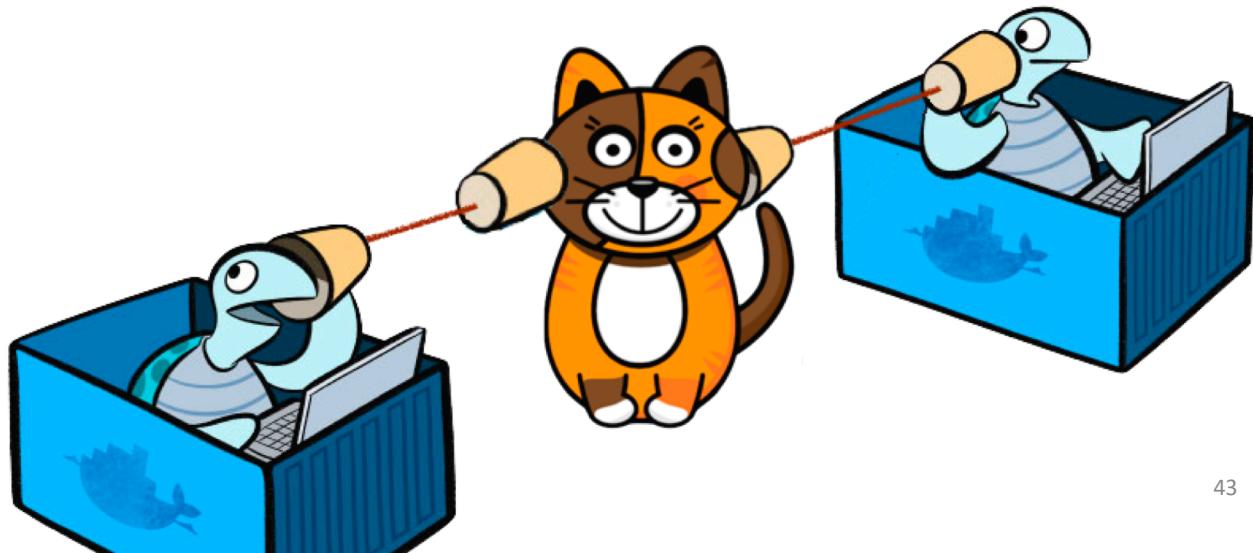
- Based on a project called SwarmKit  
<https://github.com/docker/swarmkit/>
- Introduced in Docker version 1.12
- <https://github.com/docker/swarmkit/>

## Docker Networking Design

- Networking model based on CNM
- Provides portability different infrastructures

## manual/direct connect

- macvlan
- ipvlan
- host network
- other



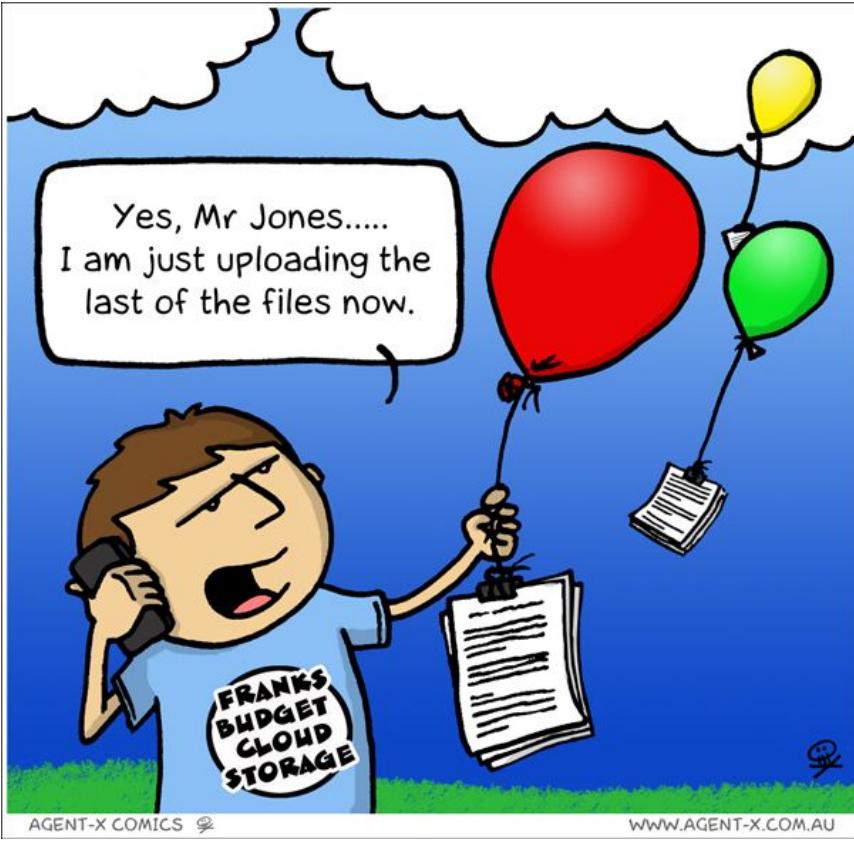
## service network sensitivity

- management network
- orchestration
- overlay sdn
  - interconnects
  - load balancing
- infrastructure services
  - storage
  - databases
  - message queues



## Section 7.3 | signal vs noise

## Section 7.3 | signal vs noise



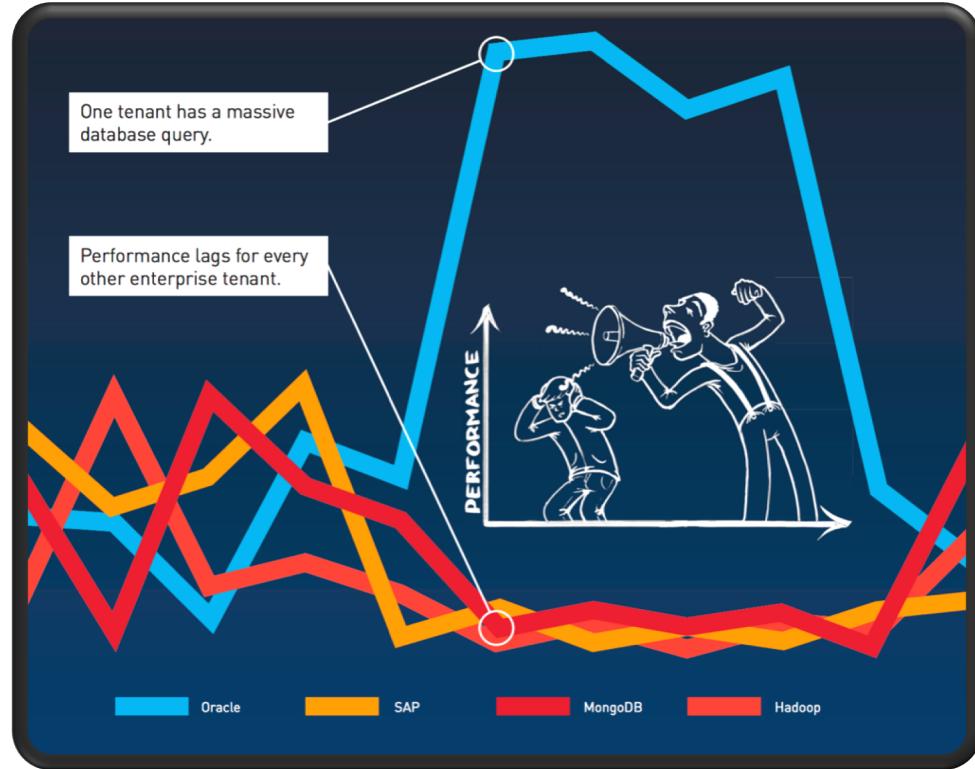
application in the “cloud” are subject to uncontrolled noise...

...a CaaS can amplify that noise significantly

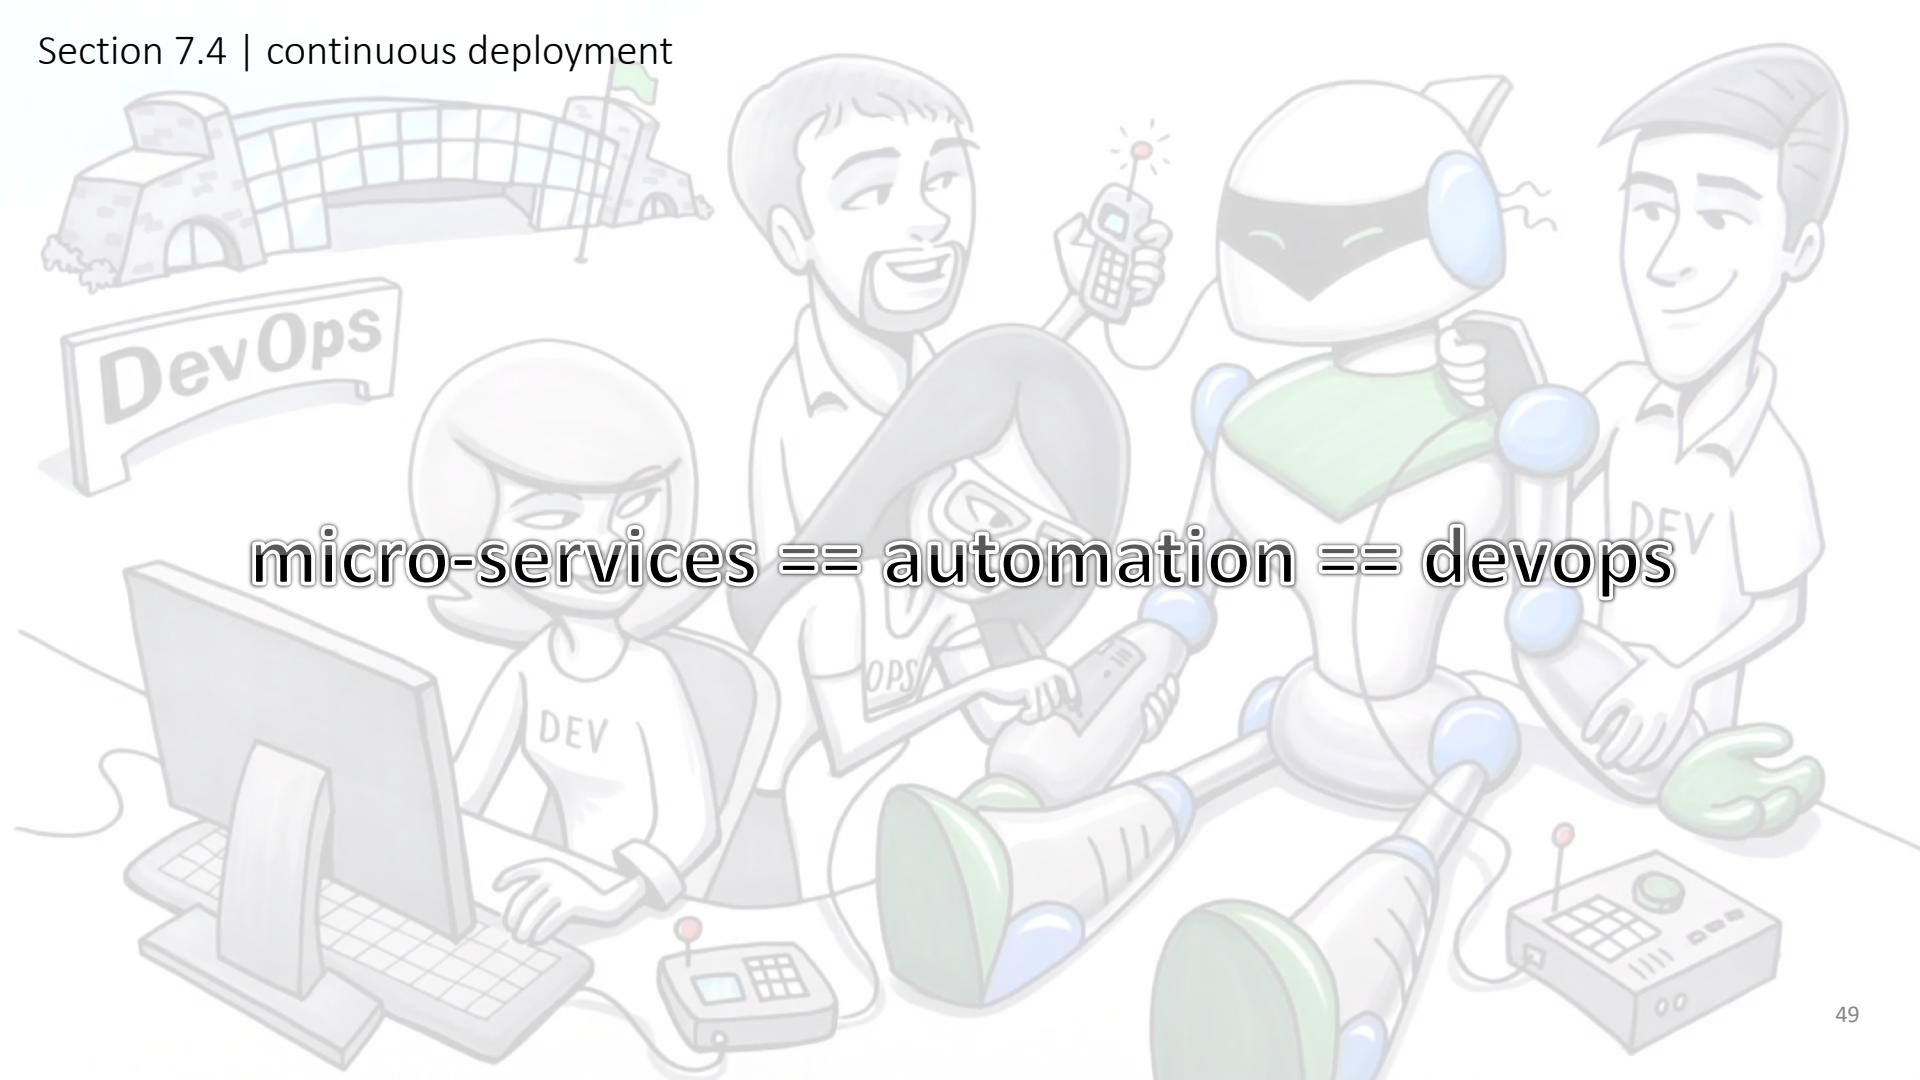


### (source) noisy neighbors

- limited visibility
- limited control
- difficult to measure



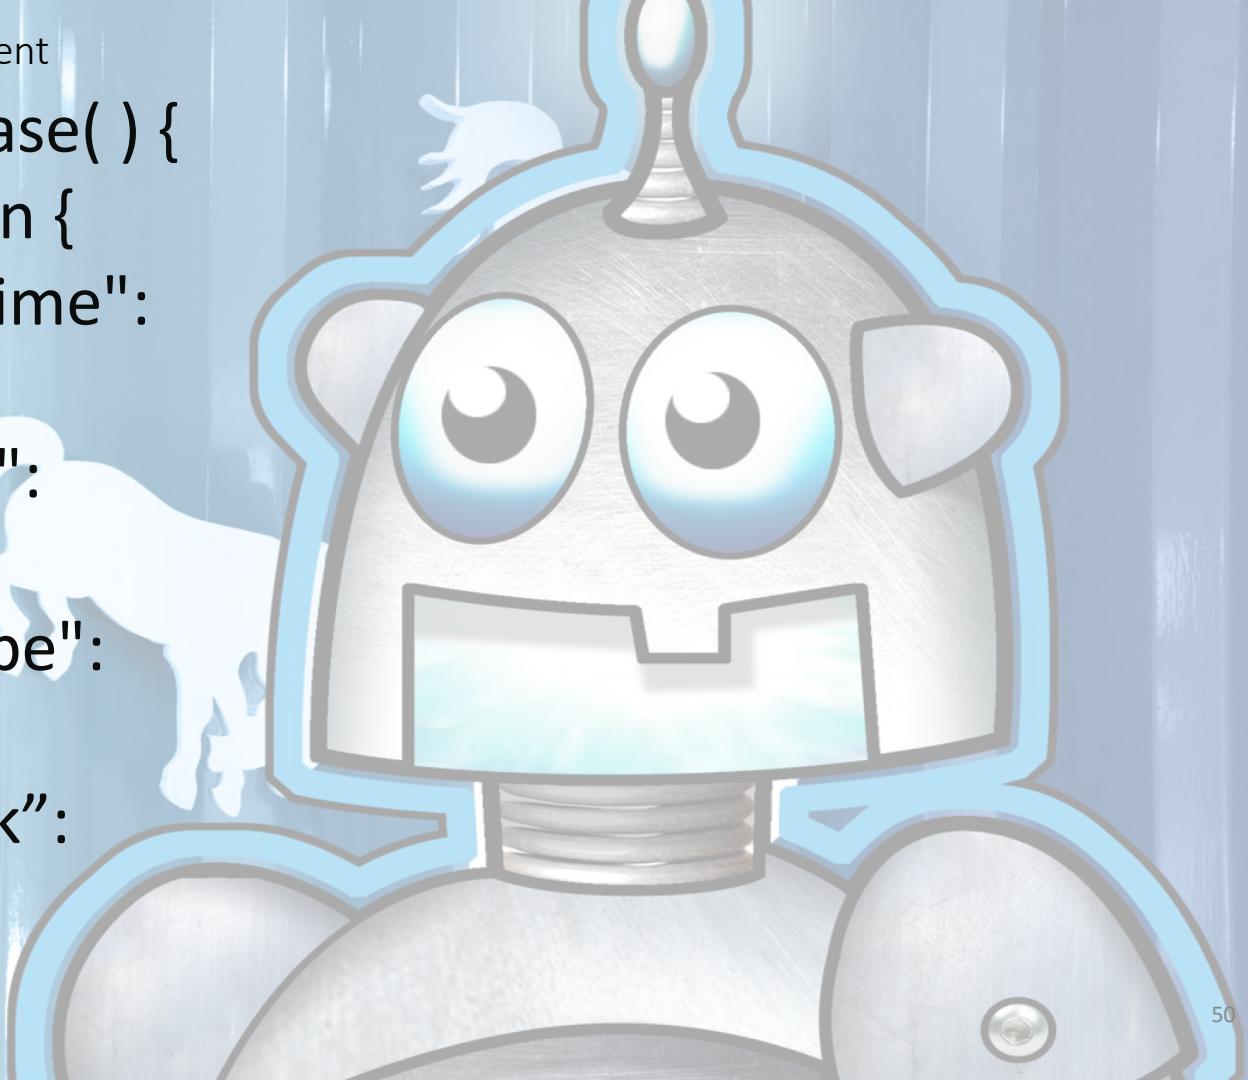
# Section 7.4 | continuous deployment



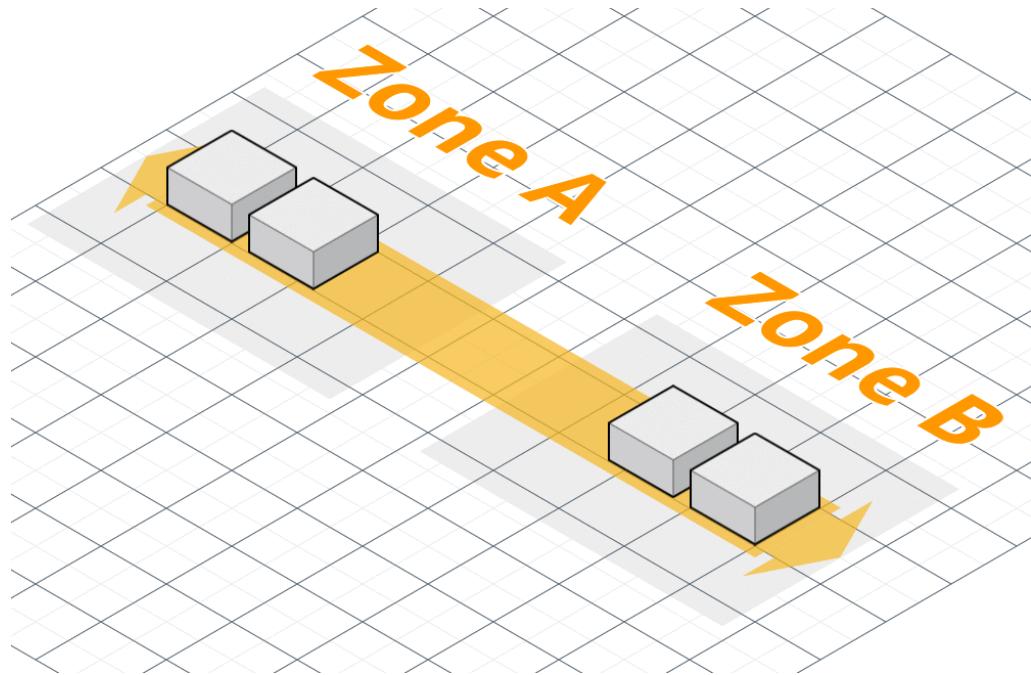
A cartoon illustration depicting a team of developers and a robot working together. In the foreground, a developer wearing a white shirt with 'DEV' on it sits at a desk, focused on a computer monitor. To their right, another developer in a similar shirt holds a smartphone. A large, friendly-looking white robot with blue glowing elements on its arms and legs is seated next to them, also interacting with the computer. In the background, a third developer stands by, holding a tablet. The scene is set against a backdrop featuring a bridge-like structure and a sign that reads 'DevOps'. The overall theme is the integration of development and operations through automation.

**micro-services == automation == devops**

```
func rolling_release( ) {  
    switch do_action {  
        case "no downtime":  
  
        case "scheduler":  
  
        case "deploy-type":  
  
        default "rollback":  
    }  
}
```



func rolling\_release( )

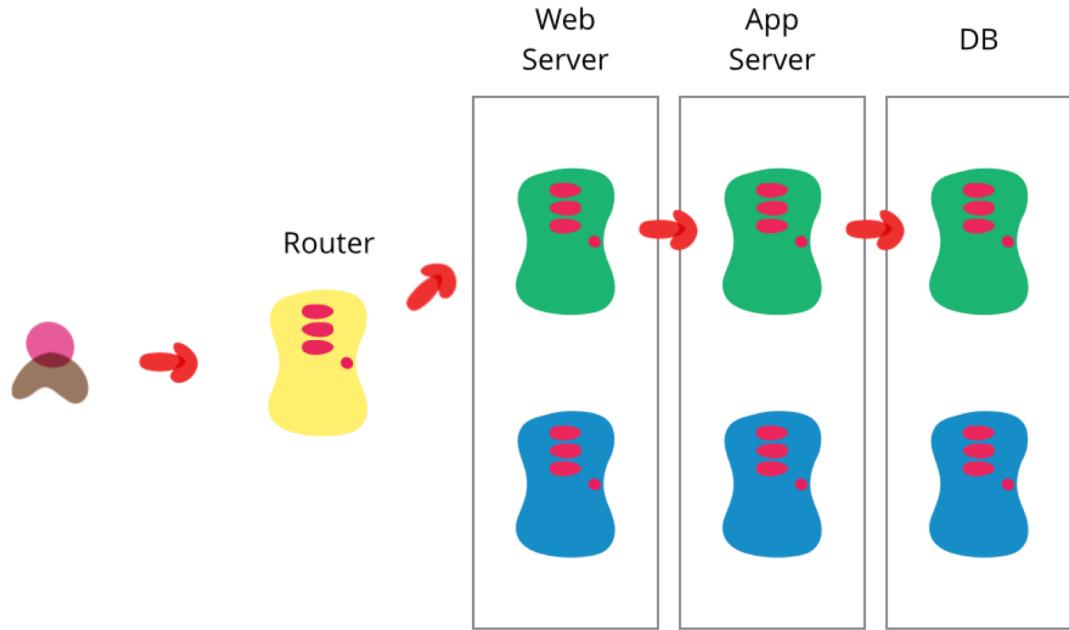




version release windows

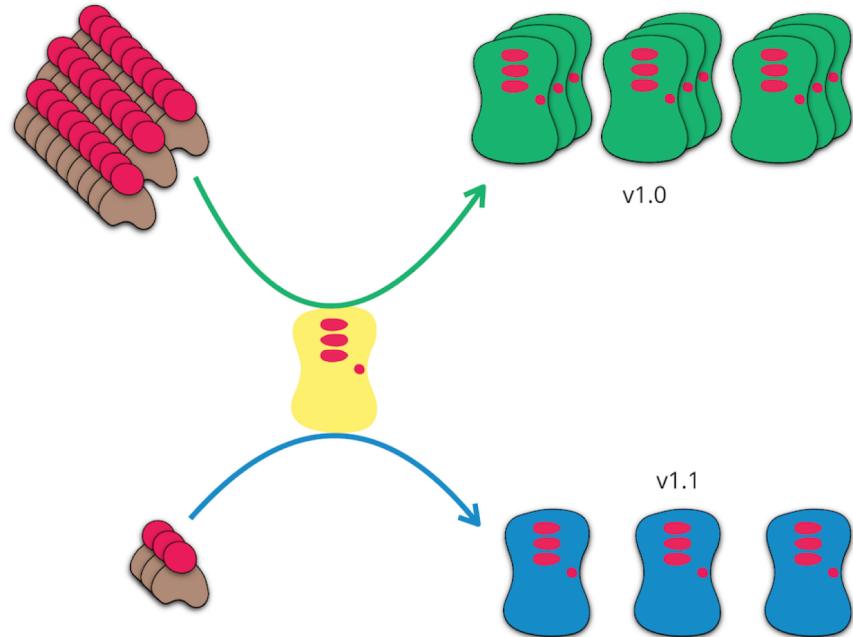
## blue green

- avoid in place upgrades
- Favor immutable infrastructure
- easy rollback plan
- Switch router to the idle environment once it has been tested and deployed



## canary

- great for experimentation of new features
- blue/green +
- great pre-test of rolling updates



# Section 7.5 | testing method vs madness



## systematic vs. specific risk

- systems must be tested as a whole
- plan in failure
- unknown unknowns

## feedback systems

- bottlenecks vs. failures
- events vs. metrics
- don't forget inspection



## Survey and Feedback

It is critical to our success as trainers to hear from you and get feedback on how we are doing. (It is also required by your employers. This is how we get graded on our work.)

- We welcome constructive criticism and feedback, and pay very close attention to your comments and suggestions. (In fact, these help shape and improve our courses, and comments from previous students have been factored in to improve the course you are taking today.)
- We're setting aside 15 minutes of class time for you to complete this before lunch, though it should actually take much less.
- Being satisfied, yet silent, actually hurts us. Please take a few minutes now to complete the survey.
- Thank you!



# Section 7.6 | lab

# Section 7.6 | lab

## lab objectives

- Deploy a CaaS in a Public Cloud
- Launch an containerized app
- Checkout a container registry



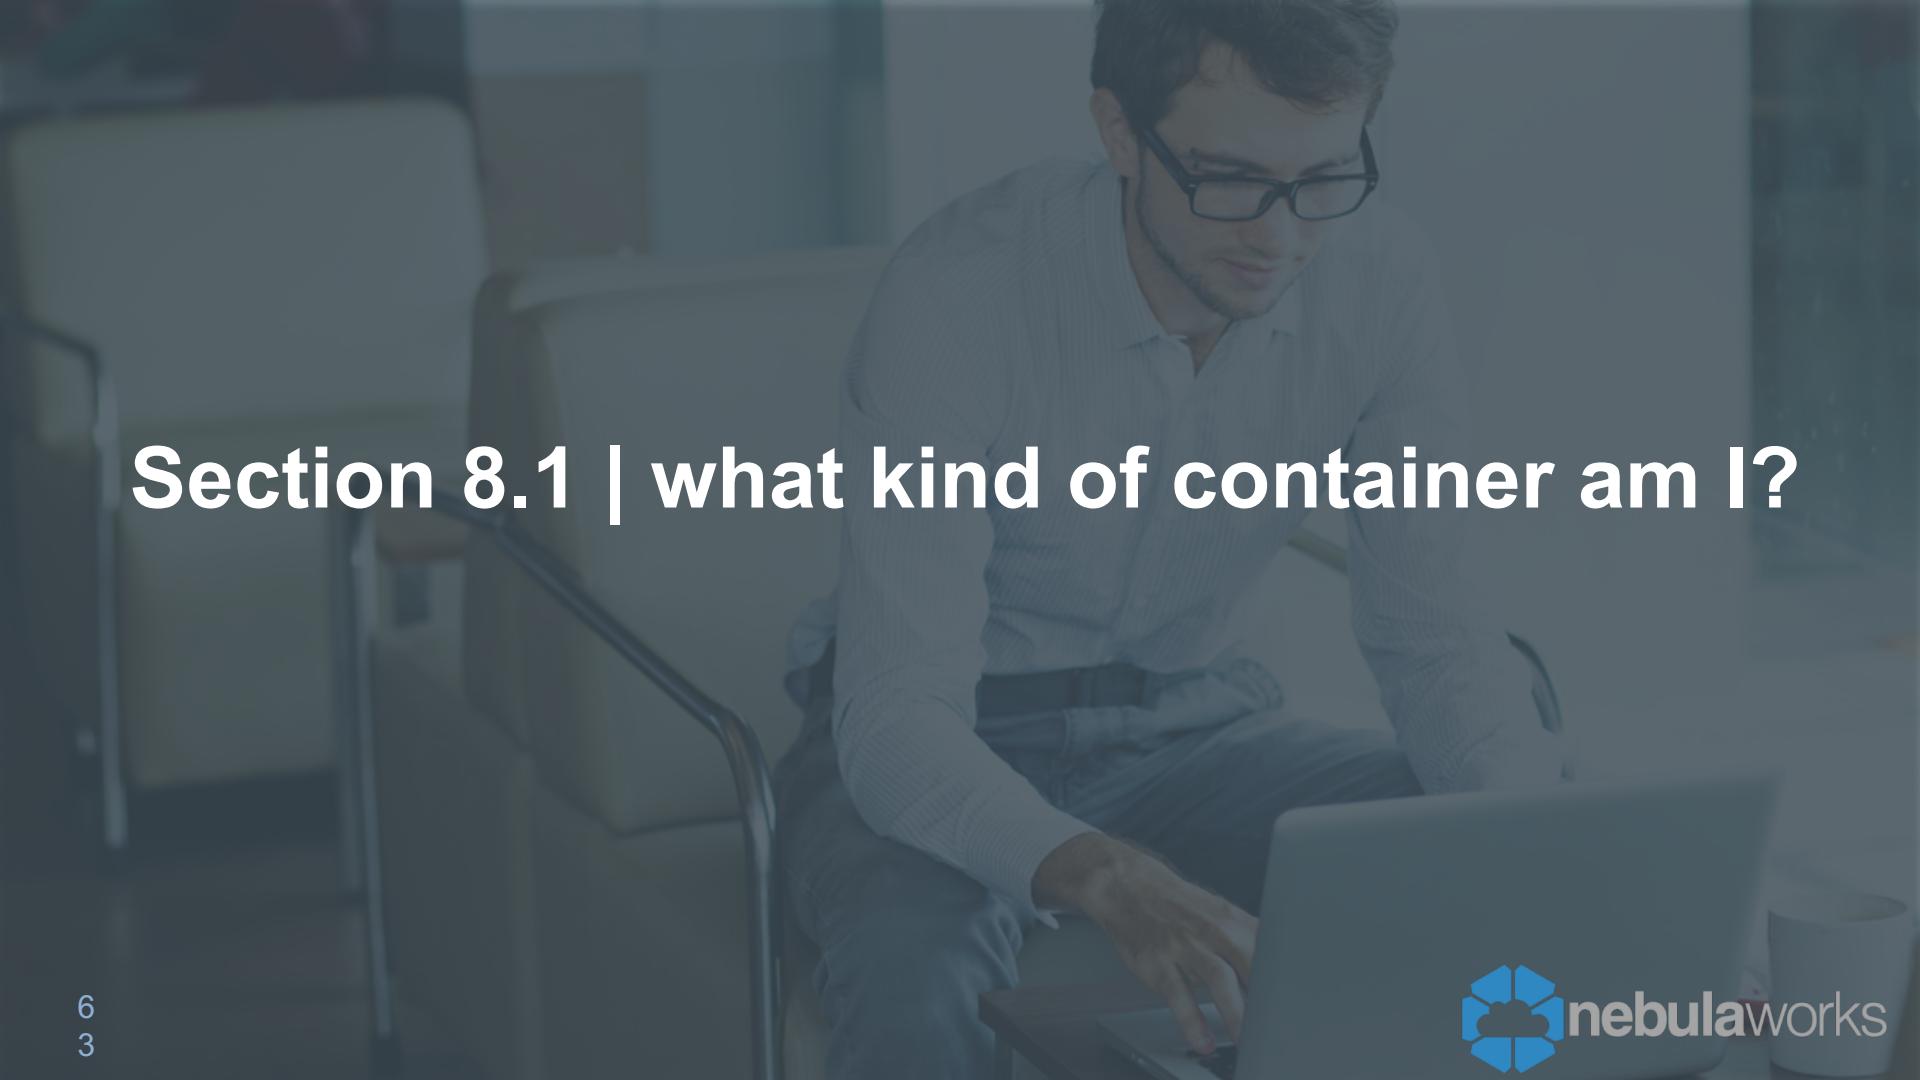
# Section 7.6 | lab

## lab 7 objectives

- Create a Kubernetes cluster
- Launch an containerized app

# **EMBRACING A PLATFORM AND AN OPINION**

module 08

A man with dark hair and glasses, wearing a light blue striped shirt, is seated at a desk in an office setting. He is looking intently at a laptop screen. The background is slightly blurred, showing office equipment and a window.

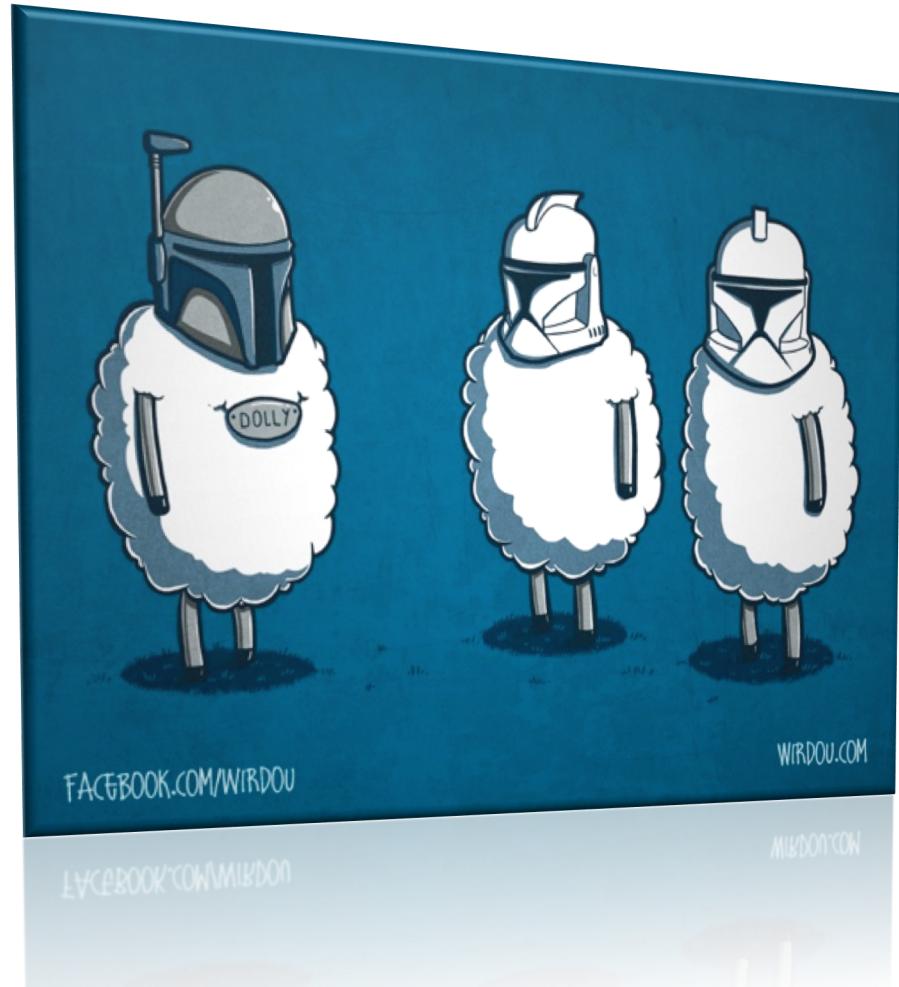
# Section 8.1 | what kind of container am I?

ultimately a container is merely a process or set of processes running within a shared kernel. the design approach used to build a given container will be one of the most deterministic factors in defining the cost against the system.

## Section 8.1 | what kind of container am I?

### full system container

- no CoW fs
- easier to turn into pets
- system vs app oriented

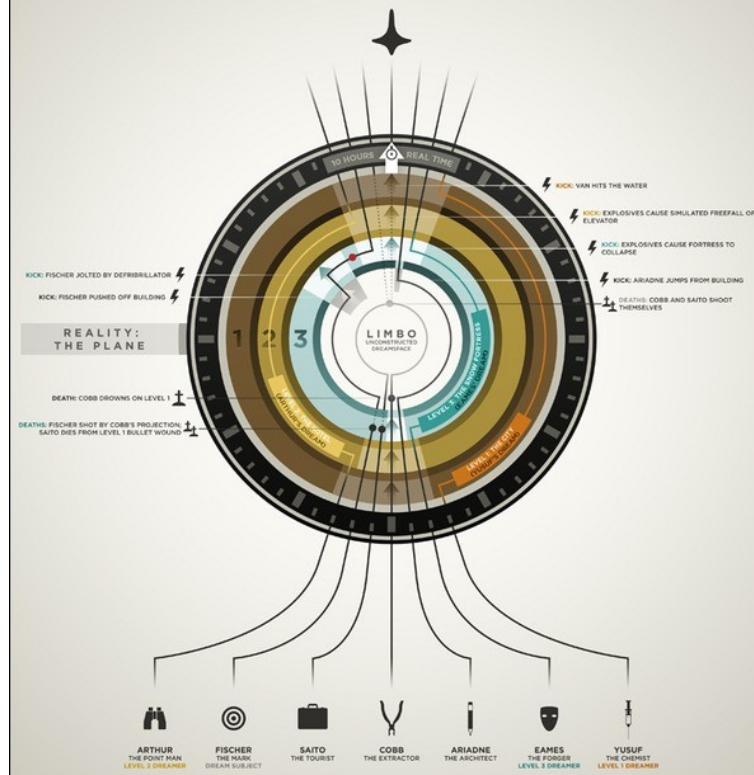


## Section 8.1 | what kind of container am I?

### virtual process container

- app vs. system oriented
- generally uses CoW fs
- closer to cattle

# INCEPTION



NEBULAWORKS

## Section 8.1 | what kind of container am I?

app container

- super lean
- micro-service friendly
- luvs pods!



NEBULAWORKS

## Section 8.1 | what kind of container am I?



### docker container

- is there any other kind
- dev friendly
- highest compatibility

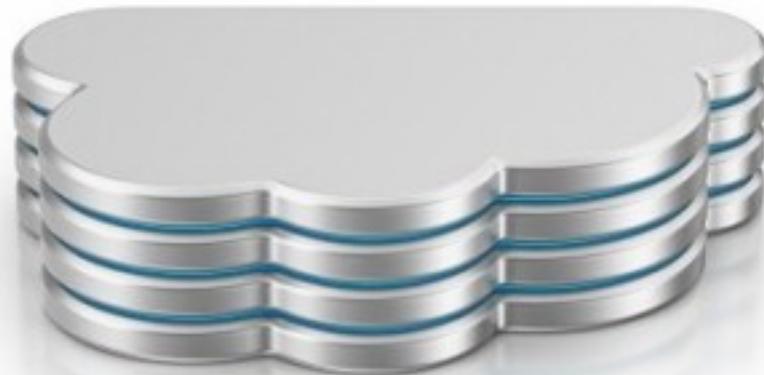


## Section 8.2 | bare-metal vs virtual machine

since containerization represents another kind of virtualization relative to machine virtualization then the choice of the underlying infrastructure used in building the host for the container engine will also impose costs. these can generally be categorized as systematic or environmental constraints.

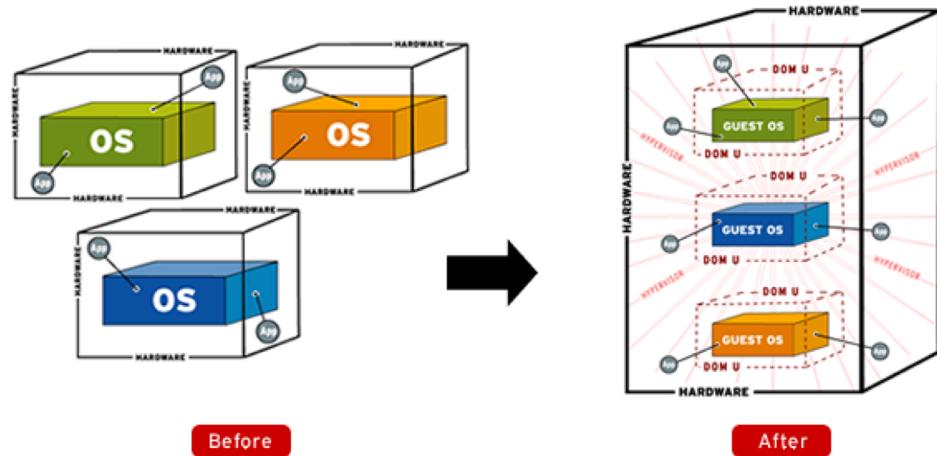
## bare-metal

- strengths
  - performance
  - simplicity
  - utilization
- weaknesses
  - expense
  - multi-tenancy
  - management



## virtual machine

- strengths
  - multi-tenancy
  - management
  - expense
- weaknesses
  - performance
  - complexity



# Section 8.3 | to persist or not to persist?

most large scale production environments deploying containers at scale treat the artifact largely as a short lived one. as such the question of persistence must be addressed to determine the limits that surround density and possibly scalability.

## Section 8.3 | to persist or not to persist?



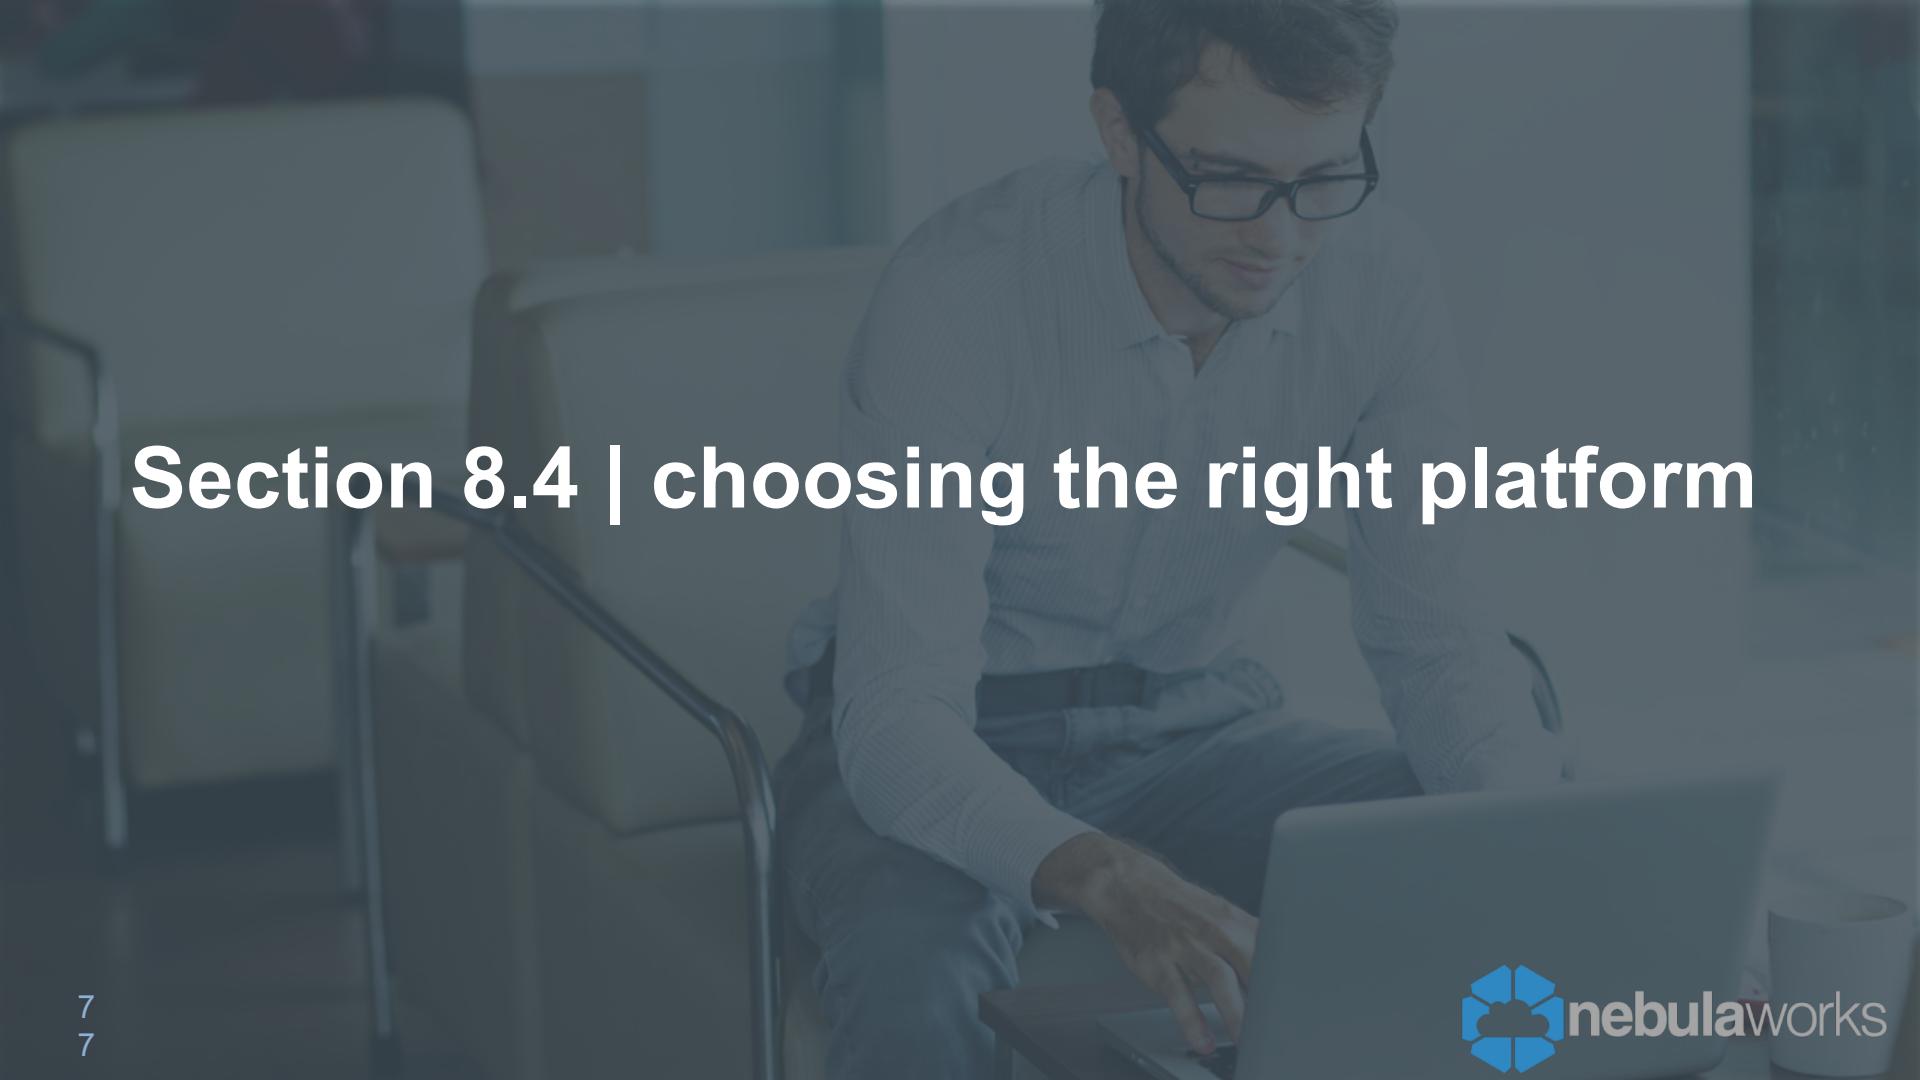
- as a network service
- persistent
  - defined as an infrastructure service
  - delivered by container?

## Section 8.3 | to persist or not to persist?

as a part of the artifact

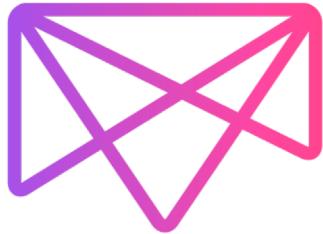
- to be or not to be ephemeral?
- how big is that artifact?
- sensitivity



A professional-looking man with dark hair and a beard, wearing black-rimmed glasses and a light blue striped button-down shirt, is seated at a desk in an office environment. He is looking intently at a laptop screen in front of him. The background is slightly blurred, showing office furniture and windows.

## Section 8.4 | choosing the right platform

## mesosphere



- Oldest Platform (at least at the core)
- Reaches beyond Containers
- Multiple schedulers available
- Container management options
- Highly scalable

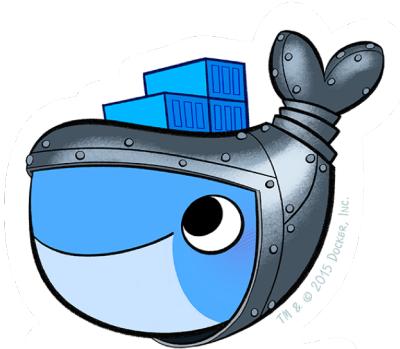
## kubernetes



- Backed by Google
- Container centric
- Complete SaaS Platform
- Container management options
- Adoption by PaaS platforms

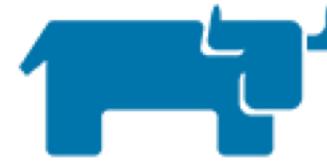
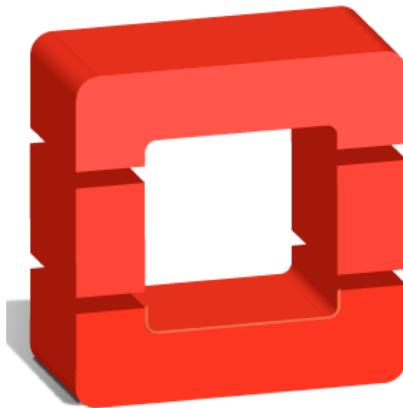
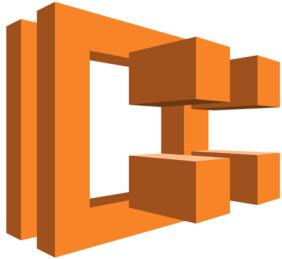
# Section 8.4 | lab

## Docker Enterprise Edition



- Docker native platform
- Looks end to end at the software cycle
- Relatively easy to use for novices
- Most widely tested Container Engine
- Available on most public clouds

others



# TROUBLESHOOTING CONTAINERS

module 09

# Section 9.1 | re-factoring a monolith



many enterprises want to use containers like “super” VMs ...  
so they start the containerization process as another form of p2v migration

base images are not AMIs (Amazon Machine Images)

- fewer default binaries
- don't have service managers
- shares a kernel

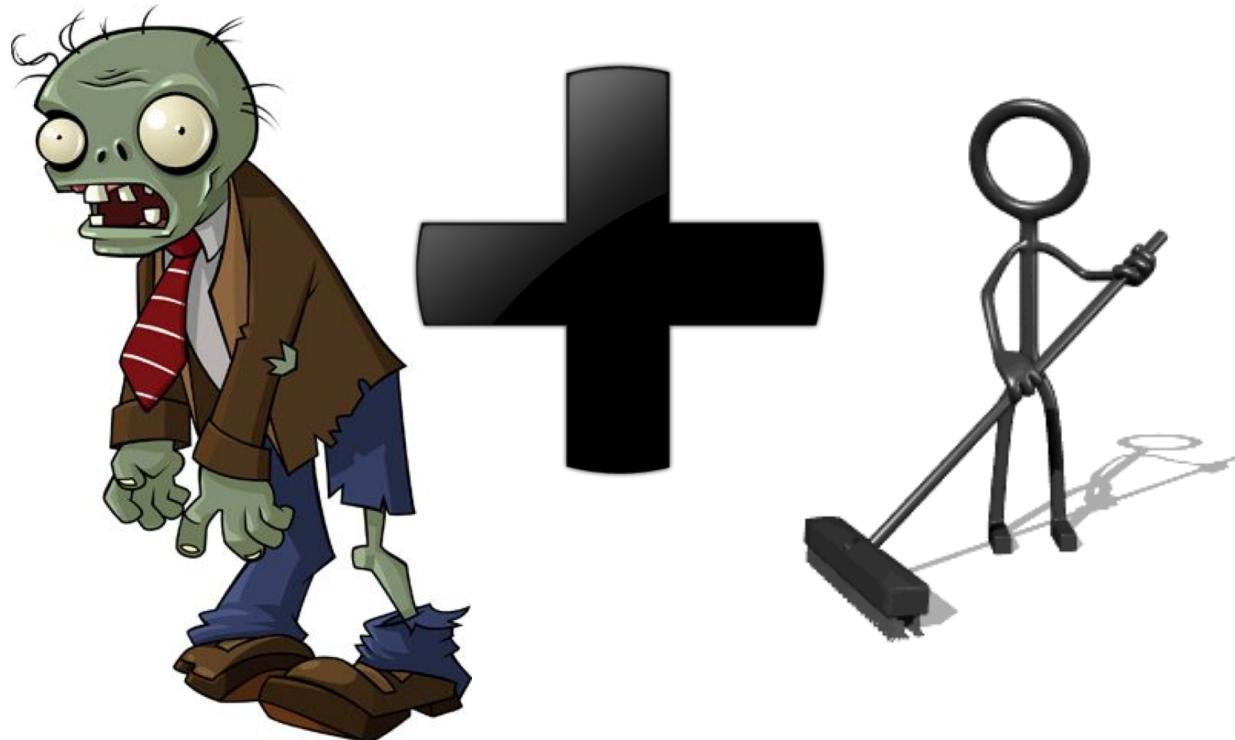




the plumbing is different

- did someone say networks
- Sometimes limited abilities to tweak things

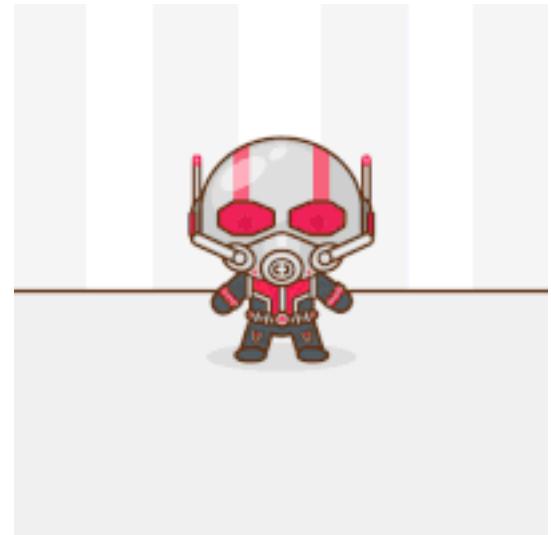
who's in charge of running the show and clean up >>>



A professional-looking man with dark hair and a beard, wearing black-rimmed glasses and a light blue striped button-down shirt, is seated at a desk in an office environment. He is leaning forward, intently focused on the screen of a laptop computer. His hands are visible on the keyboard. The background is slightly blurred, showing office furniture like a chair and a trash can.

## Section 9.2 | micro-service ready?

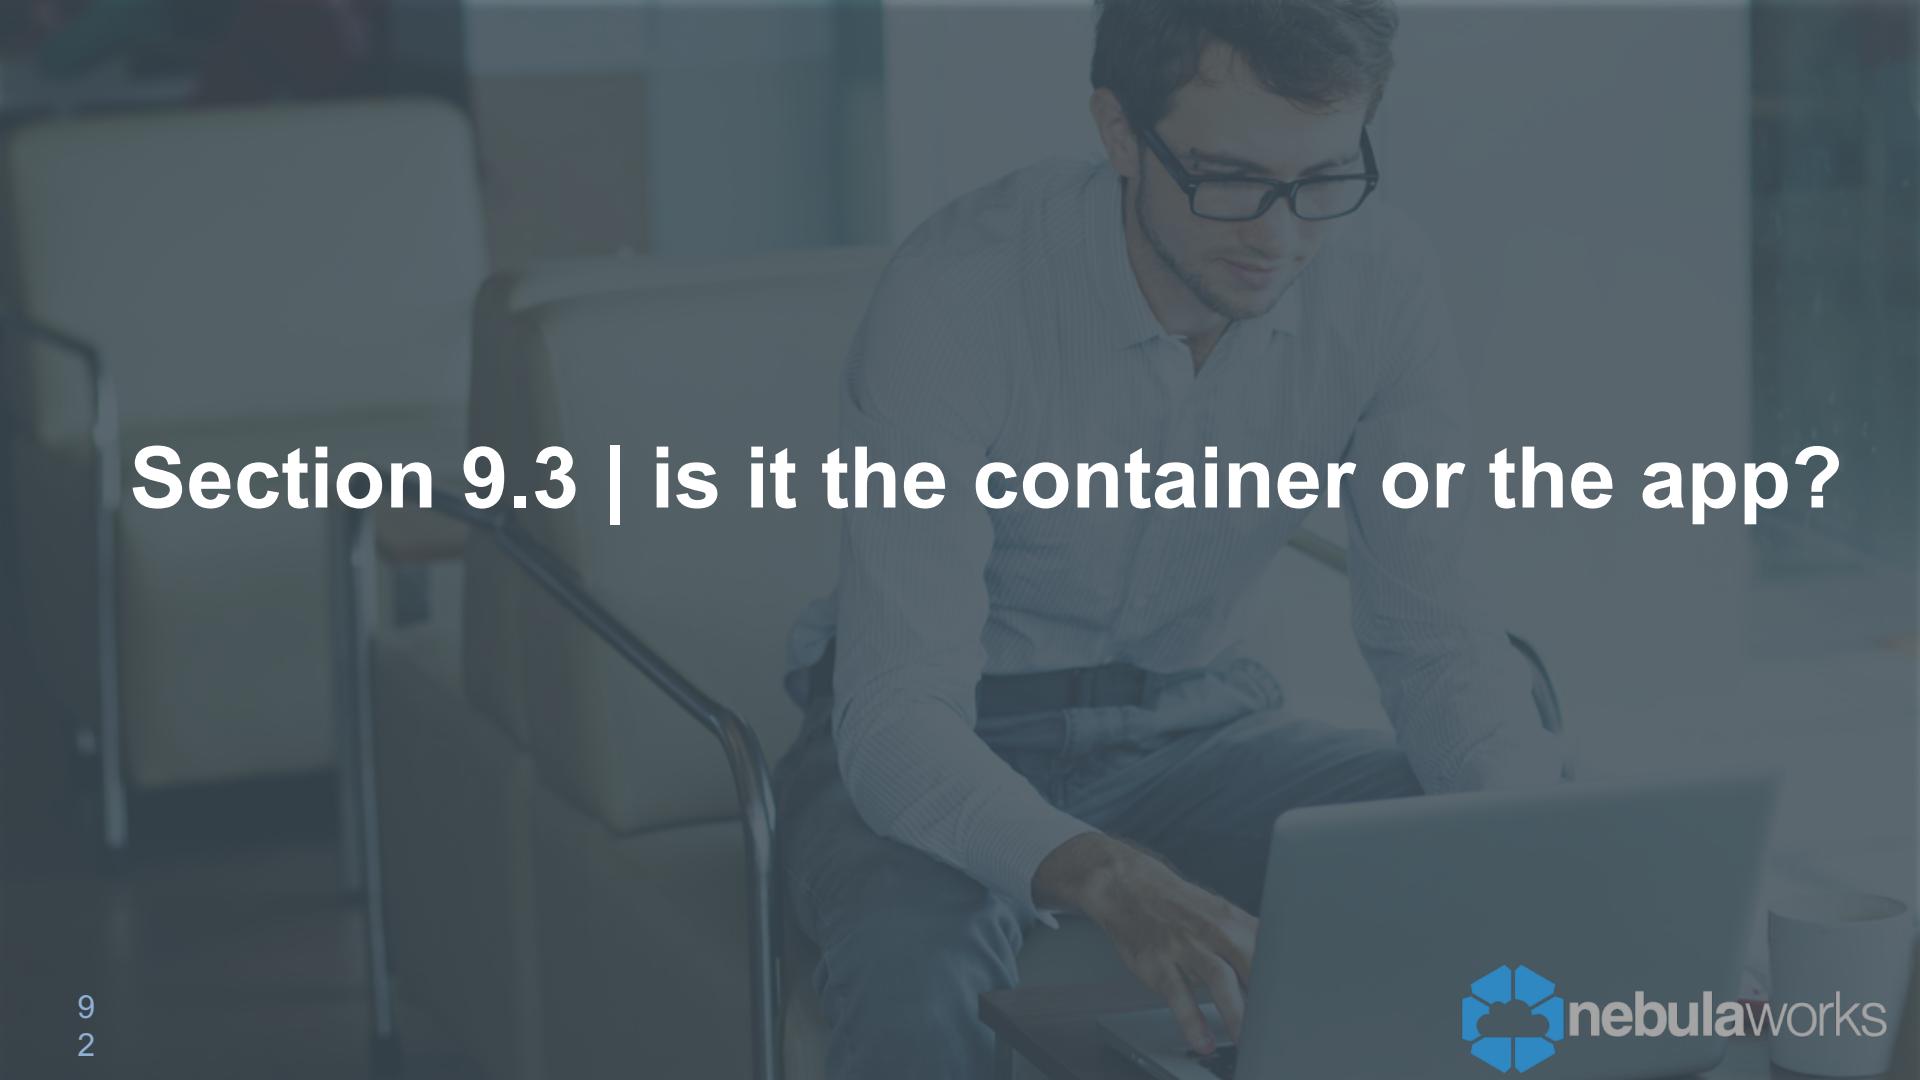
countless developers have now declared that they are "going" micro-service.



are tools available to make good determinations

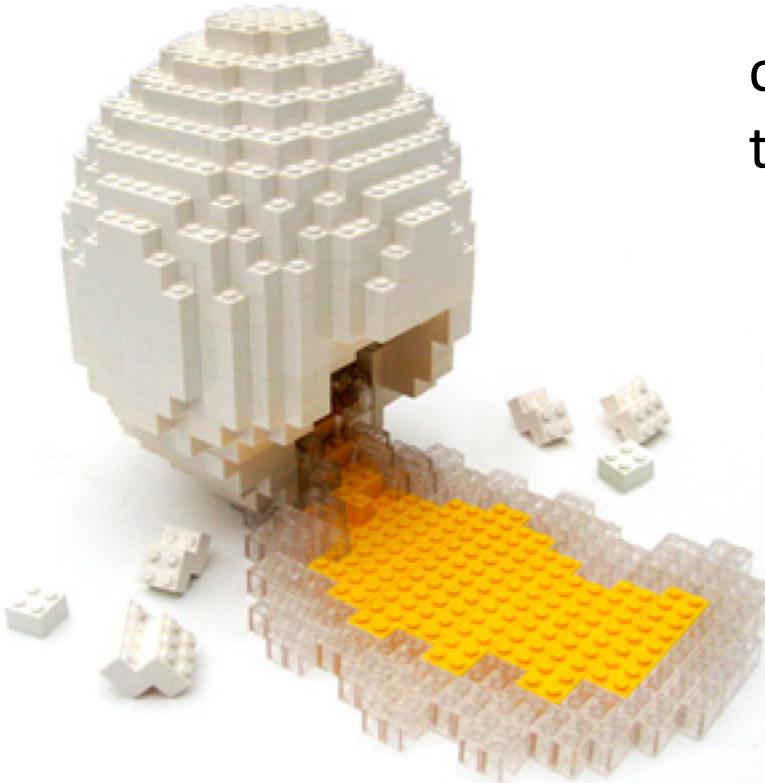
- meta data services?
- responsive to environment
- health checks?



A man with dark hair and glasses, wearing a light blue striped shirt, is seated at a desk in an office setting. He is looking intently at a laptop screen, which is partially visible in the foreground. The background shows a blurred office environment with other desks and chairs.

# Section 9.3 | is it the container or the app?

## Section 9.3 | is it the container or the app?



containers are the new lego in  
the box...

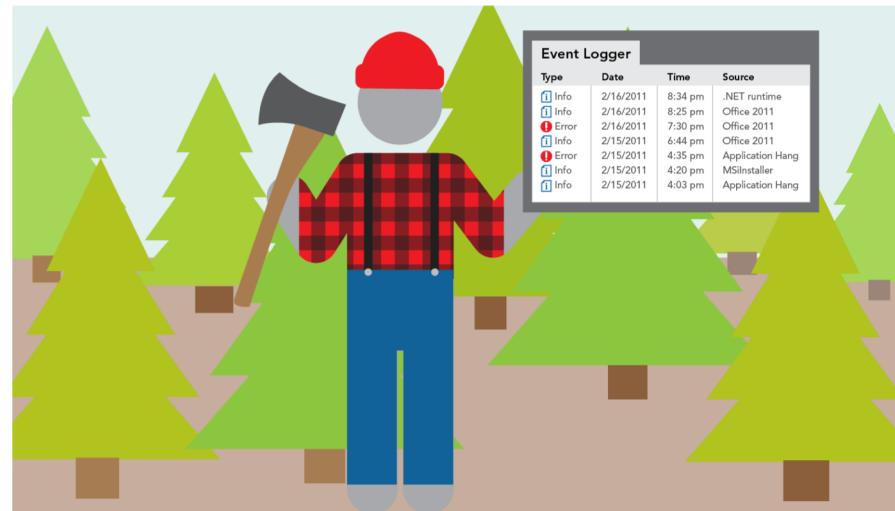
...making it the source of all new  
errors



which logs confirm your view of the problem

three-types

- daemon logs
- container logs
- application logs



how do you test conclusions

- debug/test container
- finding the root causes
- repeatable processes



# Section 9.4 | designed to be fireflies

“beautiful things never last that's why fireflies flash” – *Ron Pope*

The average lifespan of a container in the cloud is  $\frac{1}{4}$  that of an instance.



generally scaling is designed in



## problems of garbage and zombies



# **MONITORING CONTAINER SERVICES**

Module 10

# Section 10.1 | BYOLS

most existing infrastructures have logging systems in place and so there is a desire to bring containers into the fold. So when you Bring Your Own Logging System the two things to consider is how to plug containers into them and since often containers represent discrete processes what kind of metrics or events need collecting

- \* docker engine plugin architecture for logging systems
- \* alternatives to direct engine integration
- \* what are the right optics

## docker engine plugin architecture for logging systems

### key facts >>>

- supported at the daemon level
- docker logs only works with json-file and journald
- native options: json-file, syslog, journald, etwlogs (Event Tracing in Windows)
- external options: gelf, fluentd, awslogs, splunk, gcplogs

## alternatives to direct engine integration

- system level logs
  - look for systems that monitor process events and capture machine data
- mounted logs
  - containers have logging services
  - send output to volume mount
- third party platforms



what are the right optics

- types of machine data
  - metrics – performance, reliability, load
  - events – application level, system level, infrastructure level
  - stack traces – application level call traces

## Section 10.2 | metrics vs events

if the container is the new artifact and continuous deployment the new pipeline it seems obvious what's important now to watch and collect. ultimately it is the notion of predicting and managing the health of an application that will drive the selection criteria



# Section 10.3 | system lvl vs container lvl

when systems are condensed to container management system where do we plugin to get the information we need? collection services must find a ways of not only gathering the information but organizing it to close the right feedback systems



## Section 10.4 | ecosystem players

so who's providing the monitoring systems need to tame the CaaS?

- Datadog
- Sumologic
- New Relic
- CoScale
- etc



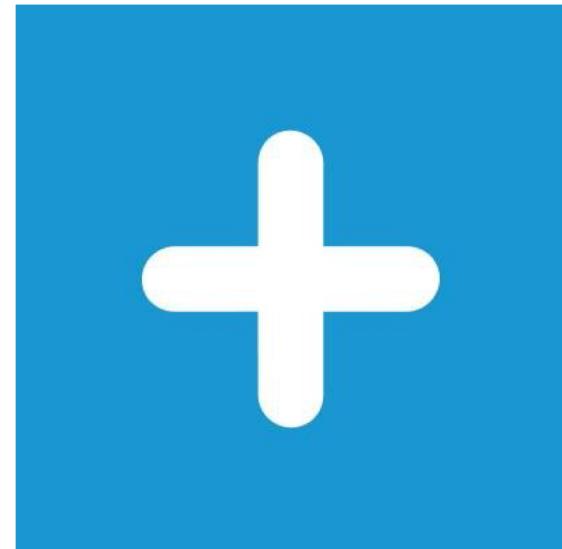


## datadog key features & facts

- founded in 2010
- SaaS based model
- custom interactive dashboards
- provides API for integrations
- provides event and metrics correlation
- <https://www.datadoghq.com/>

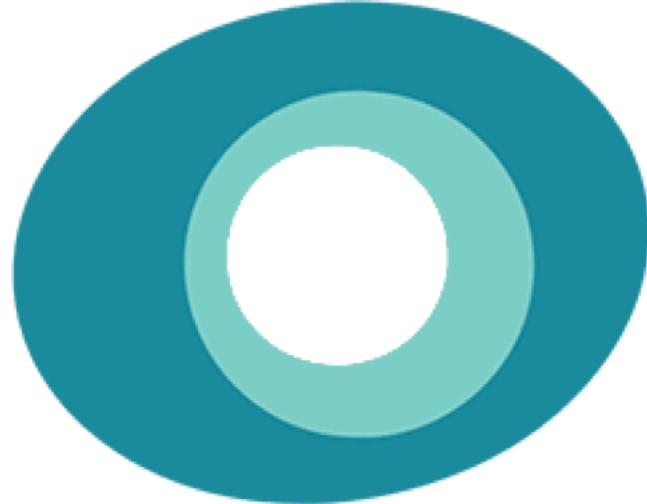
## sumologic key features & facts

- founded in 2010
- SaaS based model
- real-time log analytics
- uses machine learning
- agent based
- <https://www.sumologic.com/>



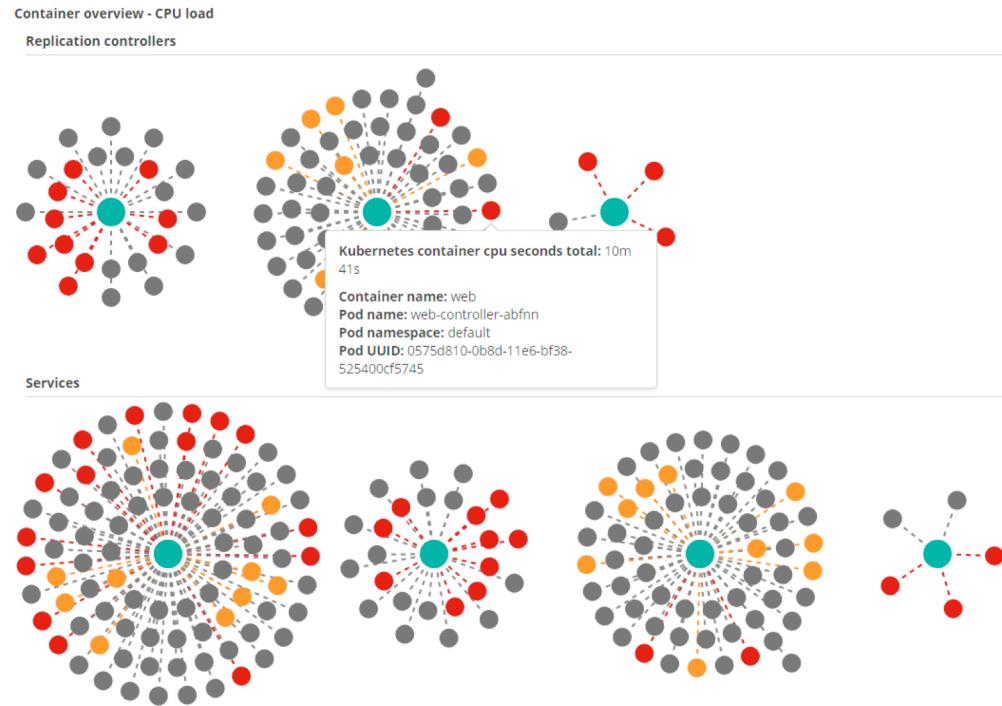
## new relic key features & facts

- founded in 2008
- SaaS based model
- application monitoring
  - end-to-end transaction tracing
  - code-level visibility
- synthetic monitoring
- iot/mobile monitoring
- <https://newrelic.com/>



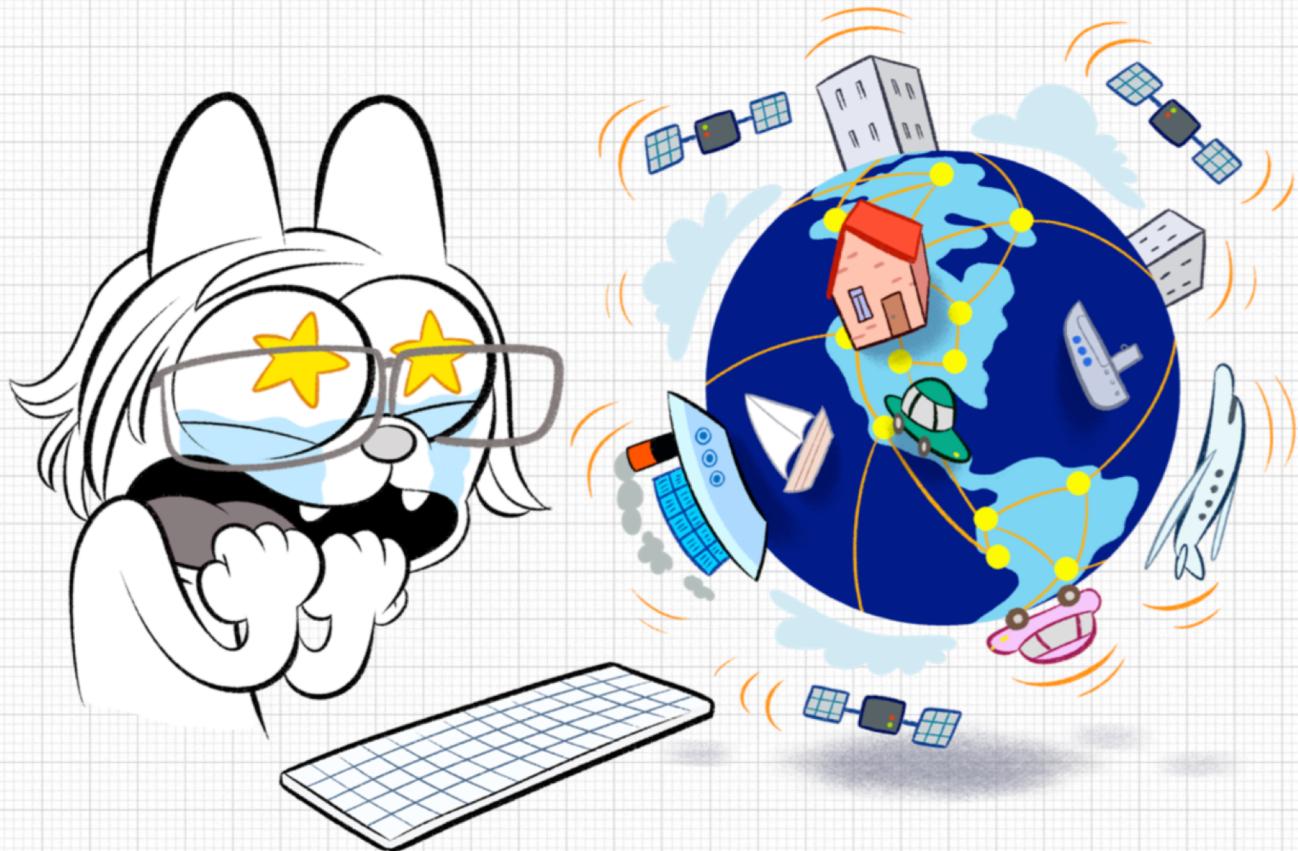
### CoScale key features & facts

- Founded ~2011
- Built on a widely-integrated platform
- Application monitoring
  - Optimized for containers and microservices
  - Container-level visibility
- Lightweight
- Automated anomaly detection
- <https://coscale.com/>



# Section 10.5 | DIY???

for the adventurous there is always the DIY approach using available OSS components to build the tailored solution to help not only monitor the obvious but one that can provide the ability to find subtle bottlenecks, work with failure and feedback systems, and forecast demand.





**THANK You!**

